


Lecture 5:


Infrastructure-as-a-Service




As of today we are in the cloud



Provisioning: How does software get into the boxes?




Hardware Virtualization

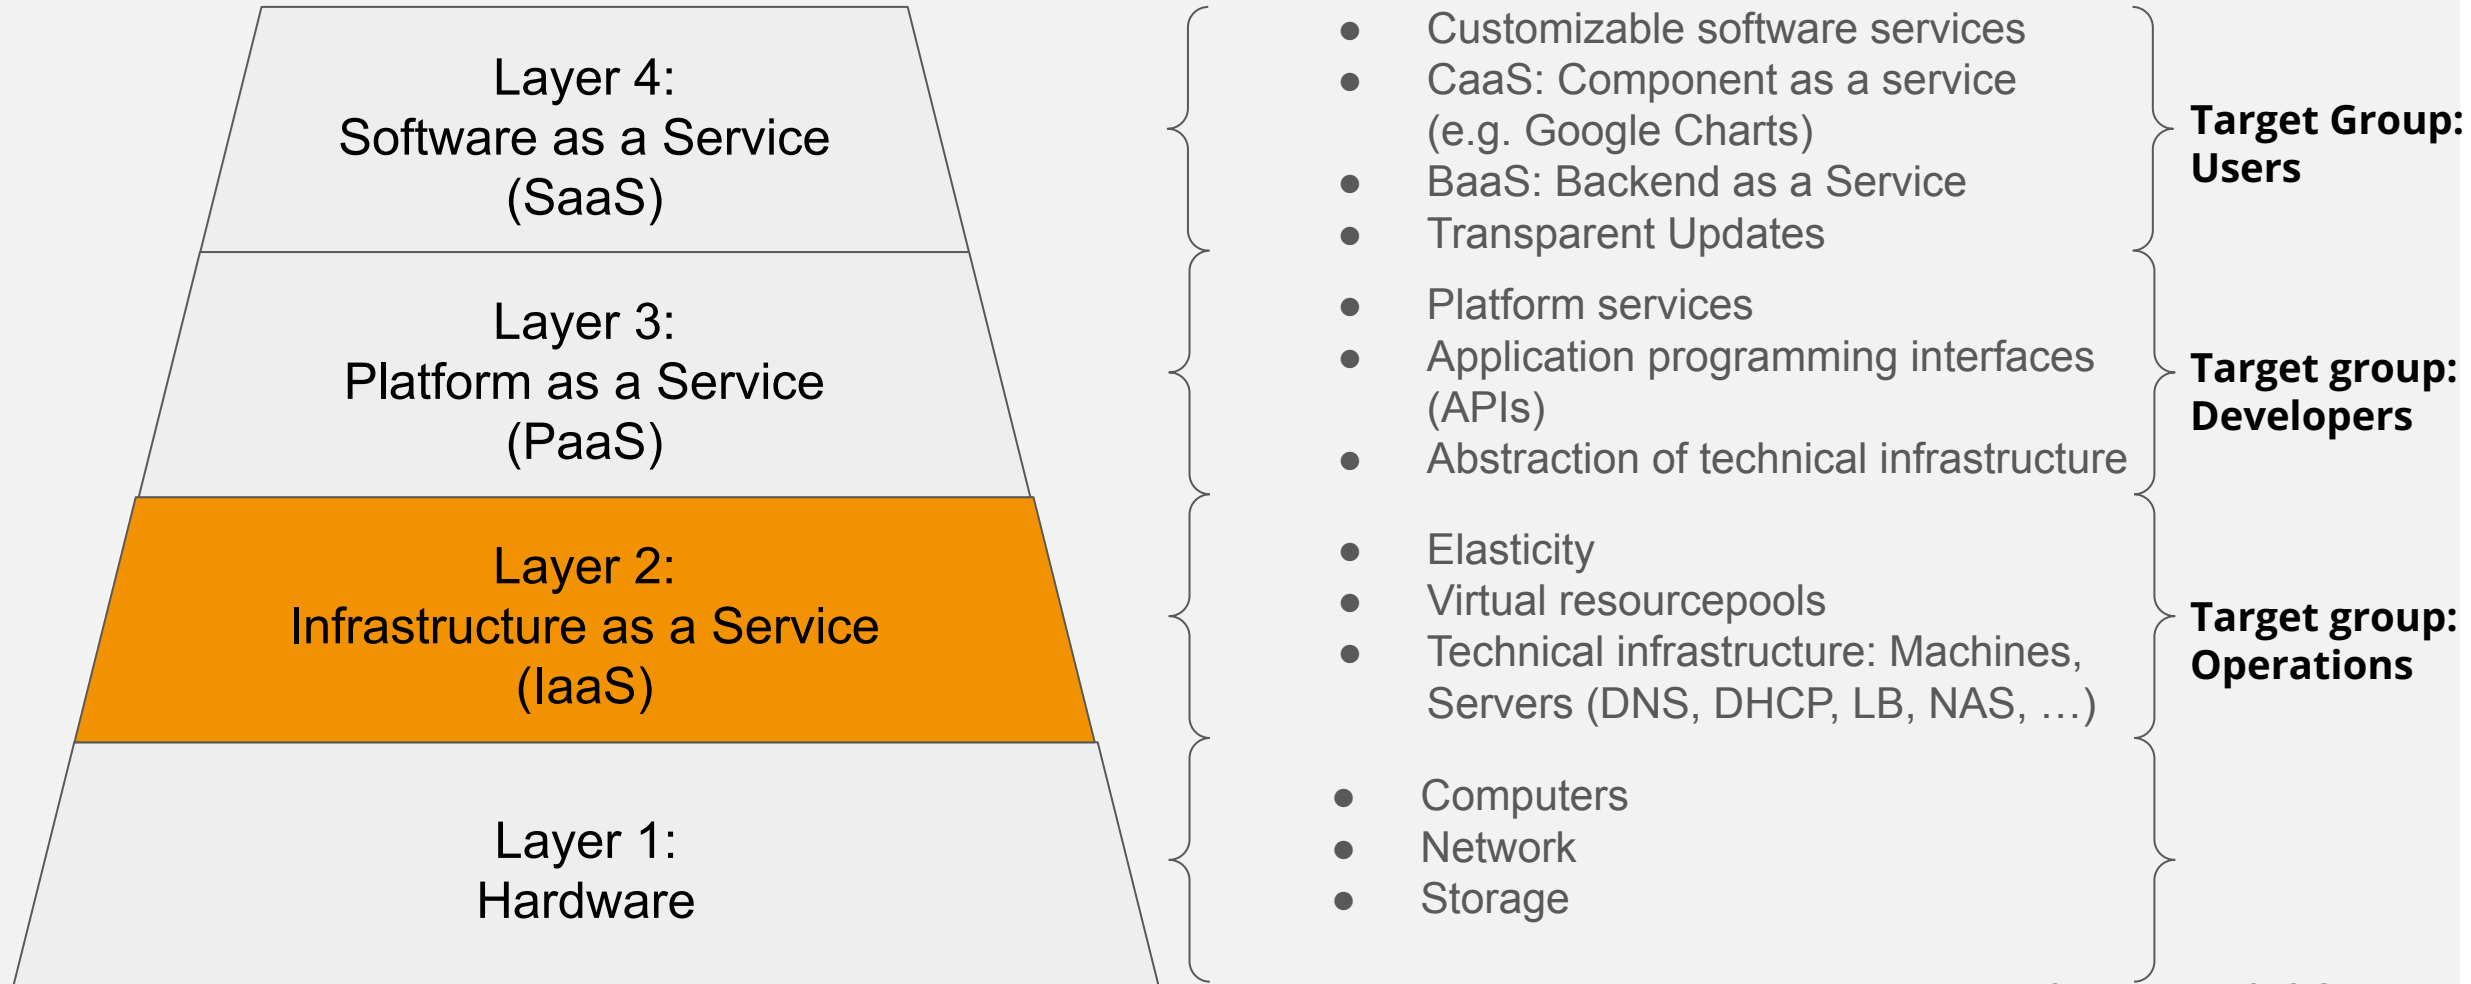


OS Virtualization

In this lecture: How to get infrastructure/hardware?



The layered model of cloud computing: From metal to application.



https://www.youtube.com/watch?v=M988_fsOSWo



QA|WARE

Introduction: Infrastructure-as-a-Service

Time2System during the last century: > 1 Jahr.



<http://de.wikipedia.org/wiki/Gro%C3%9Frechner>


Time2System in the Cloud-Era: In real time.

The **Slashdot effect**, also known as **slashdotting**, occurs when a popular [website](#) links to a smaller website, causing a massive increase in traffic. This [overloads](#) the smaller site, causing it to slow down or even temporarily become unavailable. Typically, less robust sites are unable to cope with the huge increase in traffic and become unavailable – common causes are lack of sufficient [data bandwidth](#), [servers](#) that fail to cope with the high number of requests, and traffic [quotas](#). Sites that are maintained on [shared hosting](#) services often fail when confronted with the Slashdot effect. This has the same effect as a [denial-of-service attack](#), albeit accidentally. The name stems from the huge influx of [web traffic](#) which would result from the technology news site [Slashdot](#) linking to websites. The term **flash crowd** is a more generic term.^[1]

The original circumstances have changed, as flash crowds from *Slashdot* were reported in 2005 to be diminishing due to competition from [similar sites](#),^[2] and the general adoption of elastically scalable cloud hosting platforms.



Classical operations are expensive in times of dynamic demand



Source: Amazon Web Services

Definition IaaS

IaaS refers to a business model that, contrary to the traditional purchasing of computing infrastructure, provides for renting and releasing it as needed.

Properties of an IaaS-Cloud:

- **Resource-Pools:** Availability of seemingly unlimited resources, that process requests in a distributed manner
- **Elasticity:** Dynamic allocation of additional resources based on demand
- **Pay-as-you-go Modell:** Only pay for what you actually use

Resource-Types in an IaaS-Cloud:


- **Compute:** Compute-Nodes with CPU & RAM
- **Storage:** Storage-capacity via mountable filesystems, block storage or database services.
- **Network:** Network and network-services like DNS, DHCP, VPN, CDN and Load Balancer.

Infrastructure-services in an IaaS-Cloud:


- **Monitoring**
- **Resource-Management**

Scalability: Effects

- **Daily and seasonal effects:** midday peak, prime-time peak, weekend peak, Christmas, Valentine's Day, Mother's Day, etc. (predictable load peaks)




- Continuous growth




Source: Amazon Web Services

- **Special effects:** z.B. Slashdot-Effekt (unpredictable load peaks)



- **Temporary Platforms:** Projects, Tests, Batch...



Kinds of elasticity

Demand elasticity: Allocated resources increase/decrease with demand.



- **Pseudo-elasticity:** Quick setup. Short cancellation period.
- **Real-time elasticity:** Allocation and release of resources within seconds. Automated process with manual triggers or on schedule.
- **Self-adaptive elasticity:** Automatic allocation and release of resources in real-time based on rules and metrics.

Supply elasticity: Allocated resources increase/decrease with supply.

- This is typical behavior of a grid: all available machines are allocated. Variants are also available where one can bid for free resources.

Income elasticity: Allocated resources increase/decrease with income or budget.

There are many cloud providers.




The IaaS market 2020

Figure 1. Magic Quadrant for Cloud Infrastructure and Platform Services



Amazon Leads \$100 Billion Cloud Market

Worldwide market share of leading cloud infrastructure service providers in Q2 2020*



* includes platform as a service (PaaS) and infrastructure as a service (IaaS) as well as hosted private cloud services

** 12 months ended June 30, 2020

Source: Synergy Research Group



statista


The IaaS market 2024

Figure 1: Magic Quadrant for Strategic Cloud Platform Services



Amazon Maintains Dominant Lead in the Cloud Market

Worldwide market share of leading cloud infrastructure service providers in Q3 2024*



* Includes platform as a service (PaaS) and infrastructure as a service (IaaS) as well as hosted private cloud services
Source: Synergy Research Group



Criteria when selecting an IaaS cloud

- Supported cloud variants (Private Cloud, Public Cloud, Hybrid Cloud, etc.)
- Reliability / Availability
- Security and data protection
- Predictable and stable performance
- Pricing model: Fixed and flexible costs
- Scalability: Limits, automation, and response times
- Lock-in of data and applications: Open APIs
- Liability
- Support

Service Level Agreement

Service Level Objective

An SLO is a specific measurable target that defines a key aspect of the service’s performance. It is essentially the goal or objective that the service provider aims to achieve and can include metrics like uptime, response time, or error rates.

Service Level Agreement

An SLA is a formal contract between a service provider and a customer that outlines the expected level of service, including the SLOs, and specifies the responsibilities of both parties.

Availability-classes:

Availability %	Downtime per Year	Downtime per Month	Downtime per Week
99.9% (three nines)	8.76 hours	43.2 minutes	10.1 minutes
99.95%	4.38 hours	21.56 minutes	5.04 minutes
99.99% (four nines)	52.6 minutes	4.32 minutes	1.01 minutes
99.999% (five nines)	5.26 minutes	25.9 seconds	6.05 seconds
99.9999% (six nines)	31.5 seconds	2.59 seconds	.0605 seconds

Example: Amazon S3 (Storage)

Service Commitment

AWS will use commercially reasonable efforts to make Amazon S3 available with a Monthly Uptime Percentage (defined below) of at least 99.9% during any monthly billing cycle (the “Service Commitment”). In the event Amazon S3 does not meet the Service Commitment, you will be eligible to receive a Service Credit as described below.

Monthly Uptime Percentage	Service Credit Percentage
Equal to or greater than 99% but less than 99.9%	10%
less than 99%	25%

Security Aspects of an IaaS-Cloud.

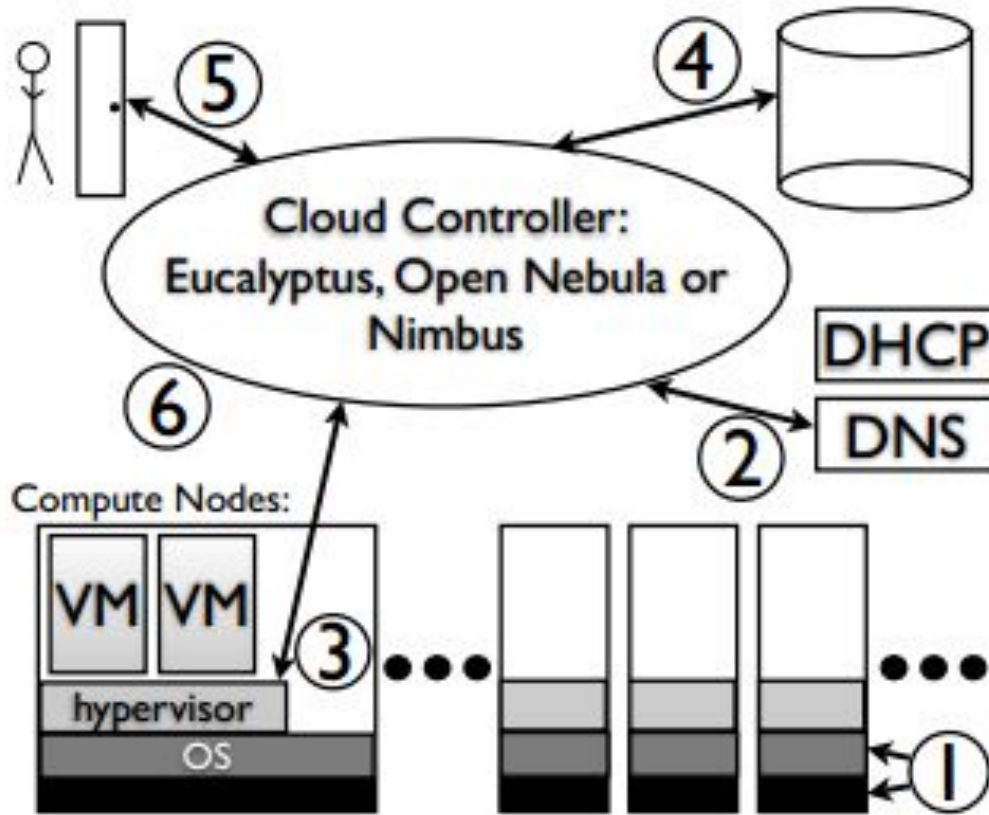
- Confidentiality of data and data communication: data encryption, VPNs
- Traceability of data: compliance with national laws (e.g., EU data protection regulation, US Patriot Act) through geographical data storage
- Firewalls and strong authentication methods
- Backup of VMs, storage, and databases
- Certifications: ISO 27001, TÜV IT



QAWARE

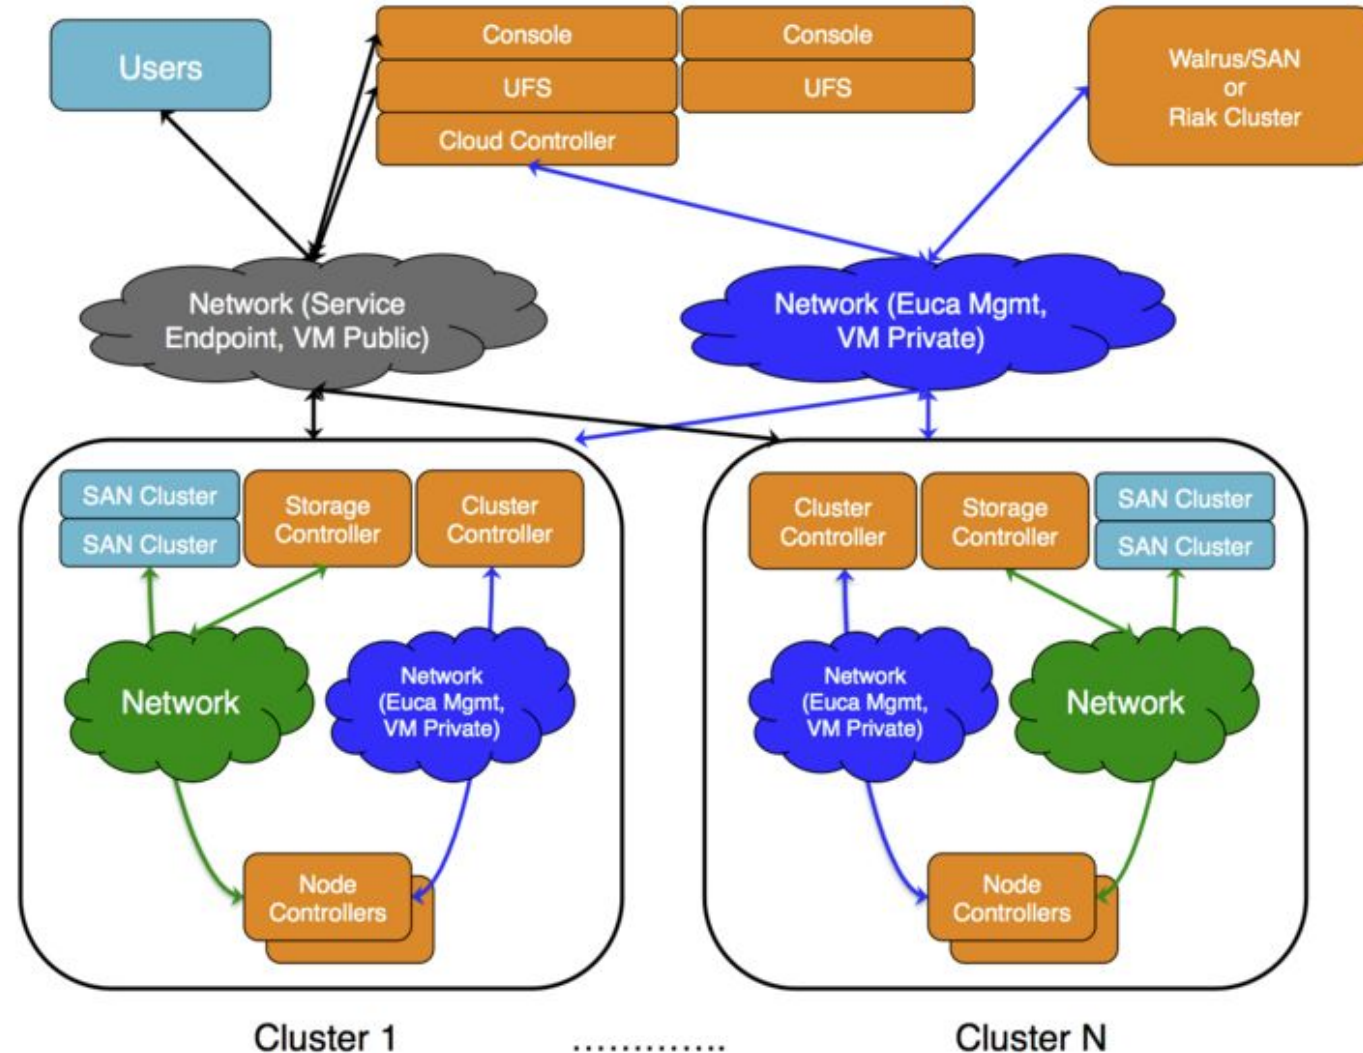
Architecture of an IaaS-Cloud

An IaaS reference architecture.



1. Hardware and OS
2. Virtual network and network services
3. Virtualization
4. Storage and Image management
5. Management interface for admins and users
6. Cloud Controller for tenant specific management of Cloud-Ressourcen

Internal architecture of an IaaS-Cloud based on Eucalyptus.



Internal architecture of an IaaS-Cloud based on Eucalyptus.



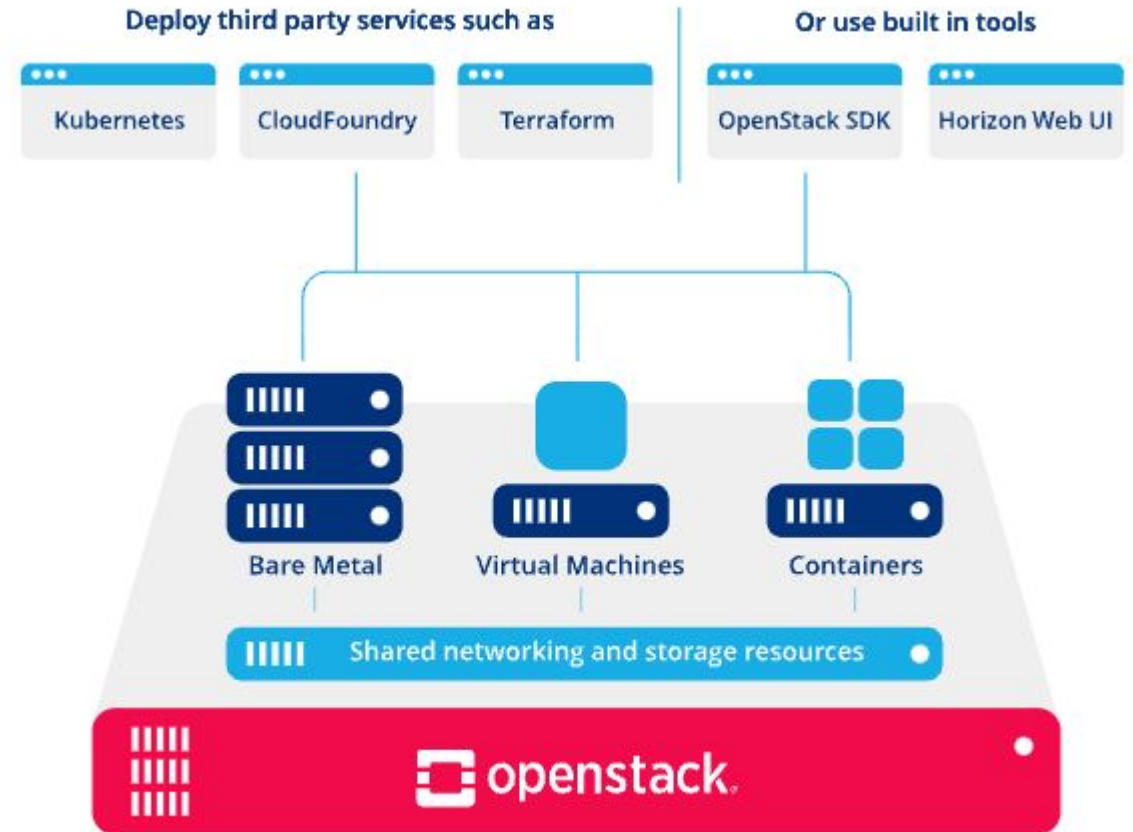


QAWARE

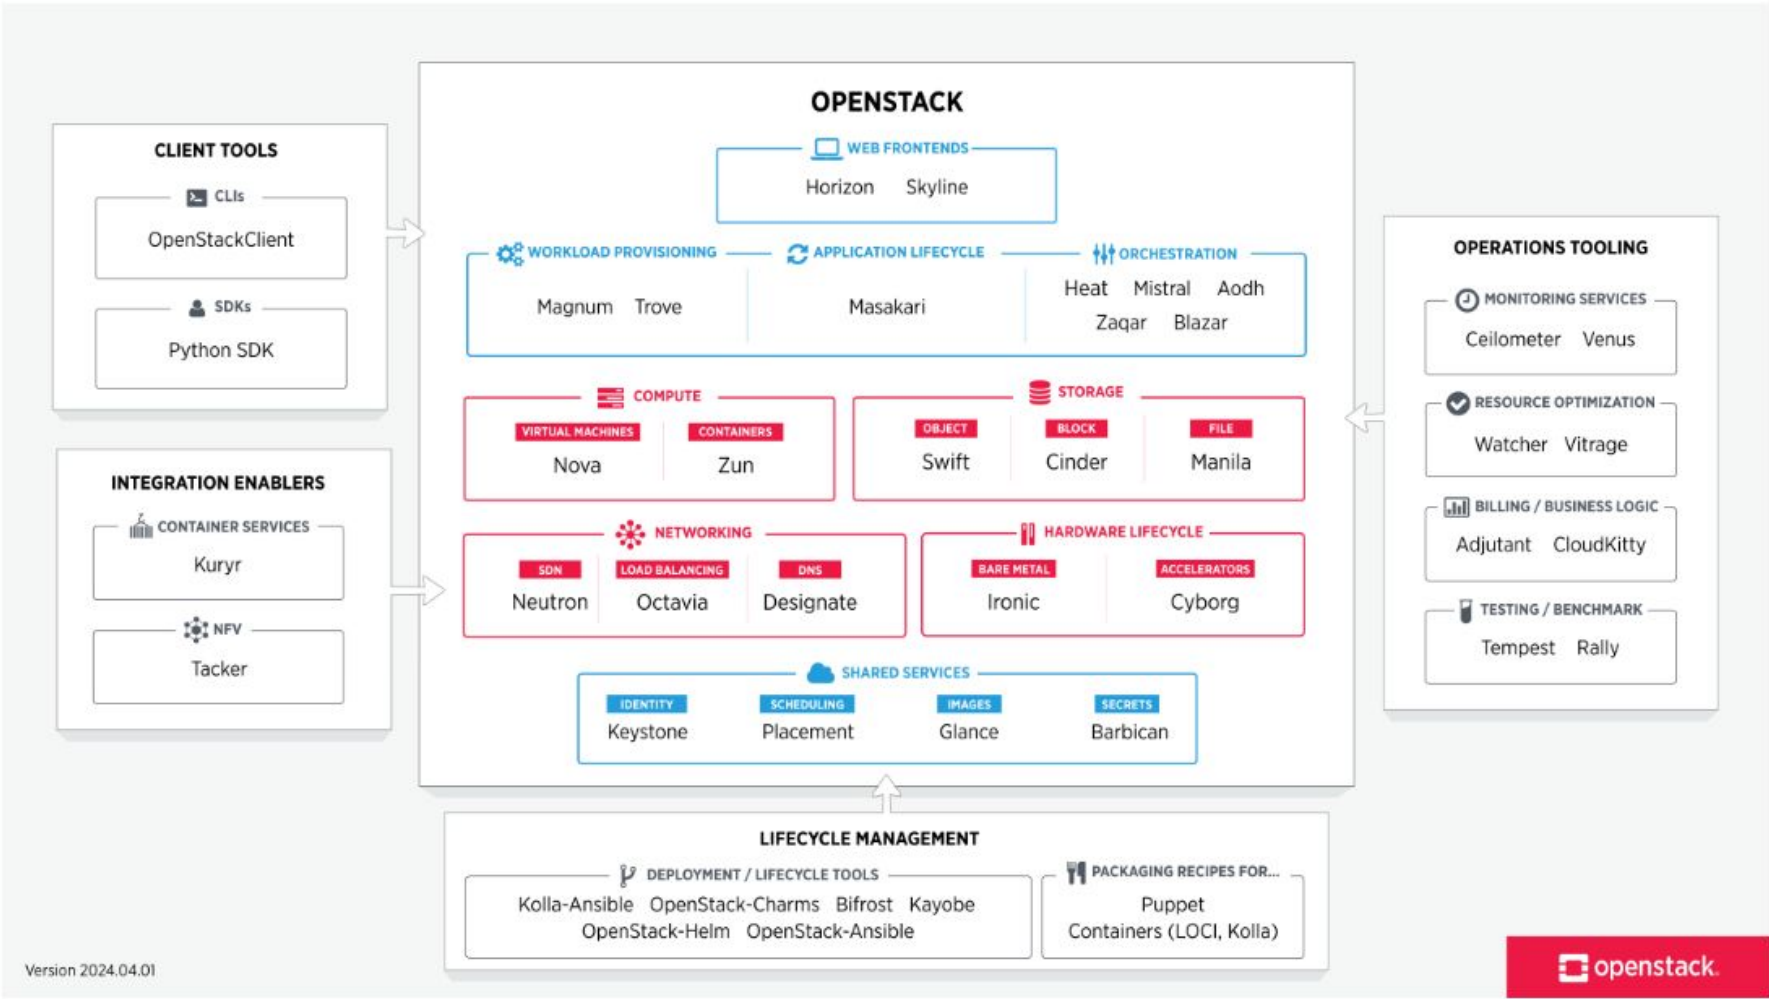
IaaS with OpenStack

OpenStack: the de-facto standard for Open-Source Private IaaS Clouds.

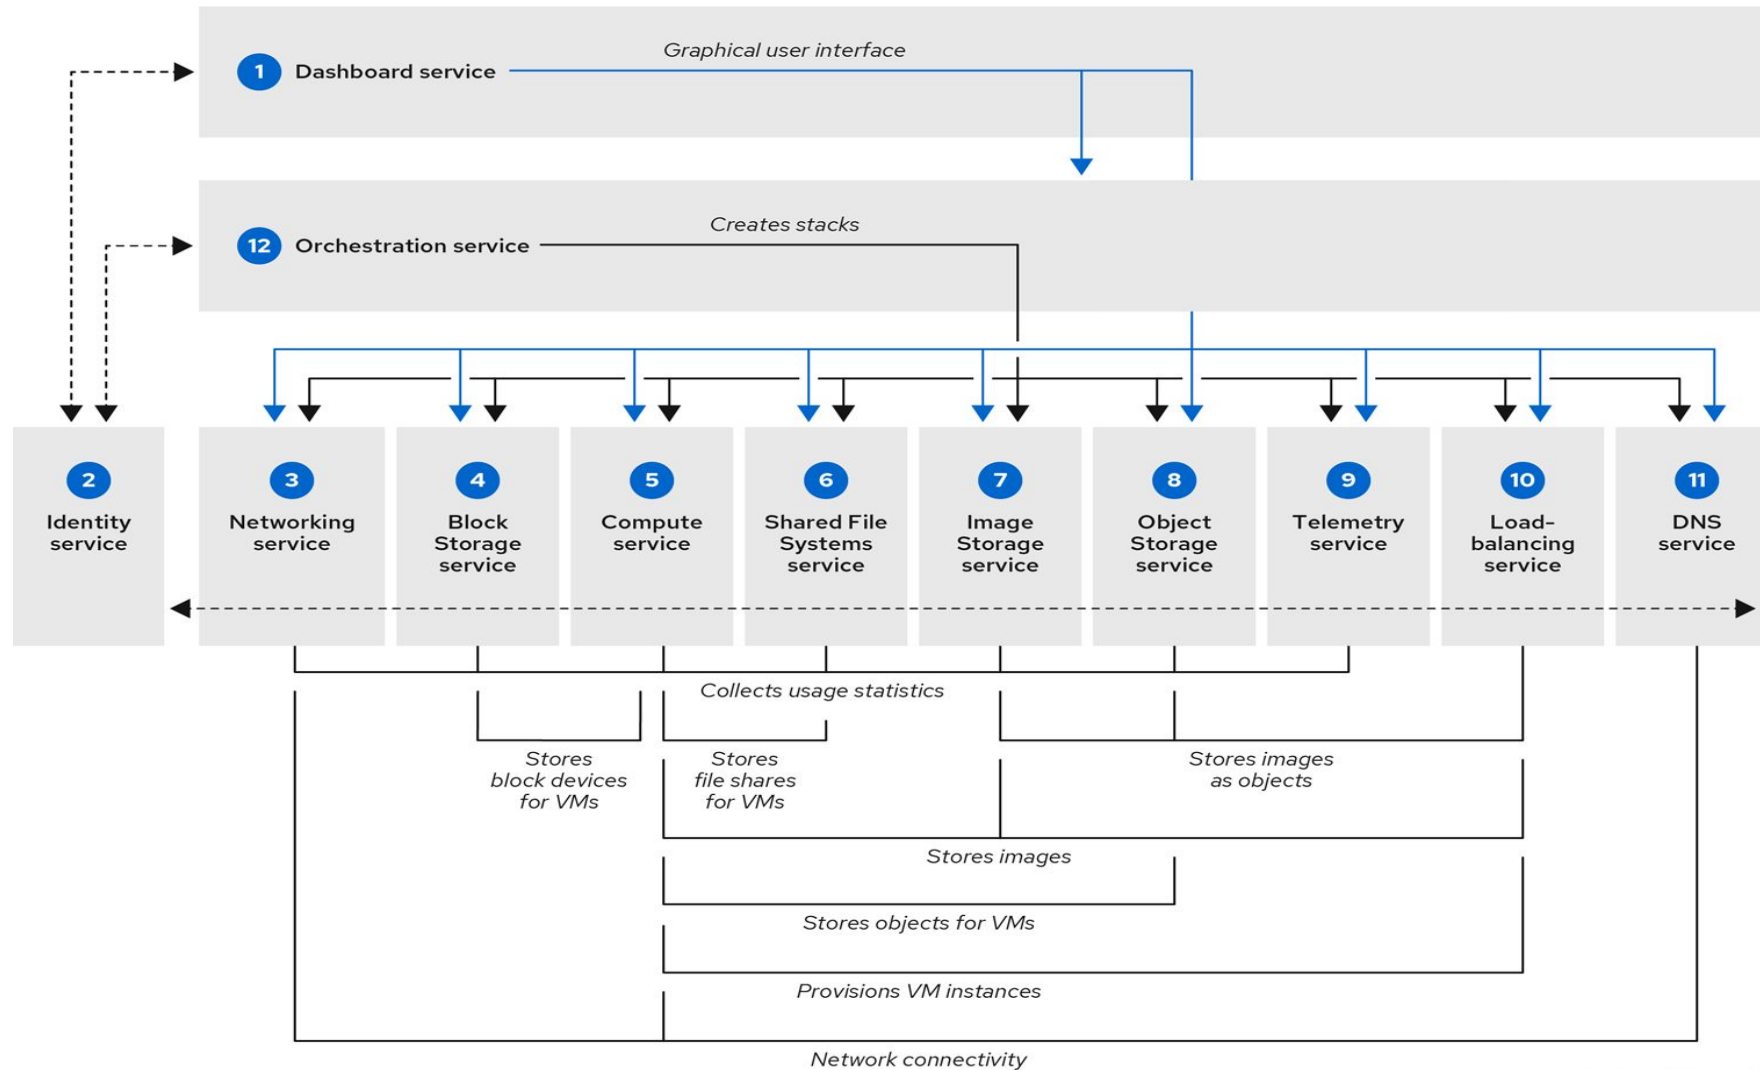
- The open-source project was significantly initiated by RackSpace and NASA.
- The first full release was in October 2010.
- Licensed under the Apache License.
- A wide range of traditional IT players (SAP, IBM, vmWare, HP, Oracle, Cisco) are part of the OpenStack community.
- Designed more as a framework than a finished system for IaaS clouds.



The OpenStack components.



Core-Components working together in OpenStack.



188_OpenStack_1221



QA|WARE

IaaS with Amazon EC2

Besides Amazon EC2 IaaS Cloud AWS offers many more IaaS-components, PaaS, Serverless-, und SaaS-services.

Product listing: <https://aws.amazon.com/de/products/>

Compute EC2 Lightsail ↗ ECR ECS EKS Lambda Batch Elastic Beanstalk Serverless Application Repository	Robotics AWS RoboMaker	Analytics Athena EMR CloudSearch Elasticsearch Service Kinesis QuickSight ↗ Data Pipeline AWS Glue AWS Lake Formation MSK	Business Applications Alexa for Business Amazon Chime ↗ WorkMail
	Customer Enablement AWS IQ ↗ Support Managed Services		End User Computing WorkSpaces AppStream 2.0 WorkDocs WorkLink
Storage S3 EFS FSx S3 Glacier Storage Gateway AWS Backup	Blockchain Amazon Managed Blockchain		
Database RDS DynamoDB ElastiCache Neptune Amazon Redshift Amazon QLDB Amazon DocumentDB	Satellite Ground Station	Security, Identity, & Compliance IAM Resource Access Manager Cognito Secrets Manager GuardDuty Inspector Amazon Macie ↗ AWS Single Sign-On Certificate Manager Key Management Service CloudHSM Directory Service WAF & Shield Artifact Security Hub	Internet Of Things IoT Core Amazon FreeRTOS IoT 1-Click IoT Analytics IoT Device Defender IoT Device Management IoT Events IoT Greengrass IoT SiteWise IoT Things Graph
Migration & Transfer AWS Migration Hub Application Discovery Service Database Migration Service Server Migration Service AWS Transfer for SFTP Snowball DataSync	Management & Governance AWS Organizations CloudWatch AWS Auto Scaling CloudFormation CloudTrail Config OpsWorks Service Catalog Systems Manager Trusted Advisor Control Tower AWS License Manager AWS Well-Architected Tool Personal Health Dashboard ↗ AWS Chatbot	Mobile AWS Amplify Mobile Hub AWS AppSync Device Farm	Game Development Amazon GameLift
Networking & Content Delivery VPC CloudFront Route 53 API Gateway Direct Connect AWS App Mesh AWS Cloud Map Global Accelerator ↗	Media Services Elastic Transcoder Kinesis Video Streams MediaConnect MediaConvert MediaLive MediaPackage MediaStore MediaTailor Elemental Appliances & Software	AR & VR Amazon Sumerian	
Developer Tools CodeStar CodeCommit CodeBuild CodeDeploy CodePipeline Cloud9 X-Ray	Machine Learning Amazon SageMaker Amazon Comprehend AWS DeepLens Amazon Lex Machine Learning Amazon Polly Rekognition Amazon Transcribe Amazon Translate Amazon Personalize Amazon Forecast Amazon Textract AWS DeepRacer	Application Integration Step Functions Amazon EventBridge Amazon MQ Simple Notification Service Simple Queue Service SWF	
		AWS Cost Management AWS Cost Explorer AWS Budgets AWS Marketplace Subscriptions	
		Customer Engagement Amazon Connect Pinpoint Simple Email Service	

Global distribution of AWS

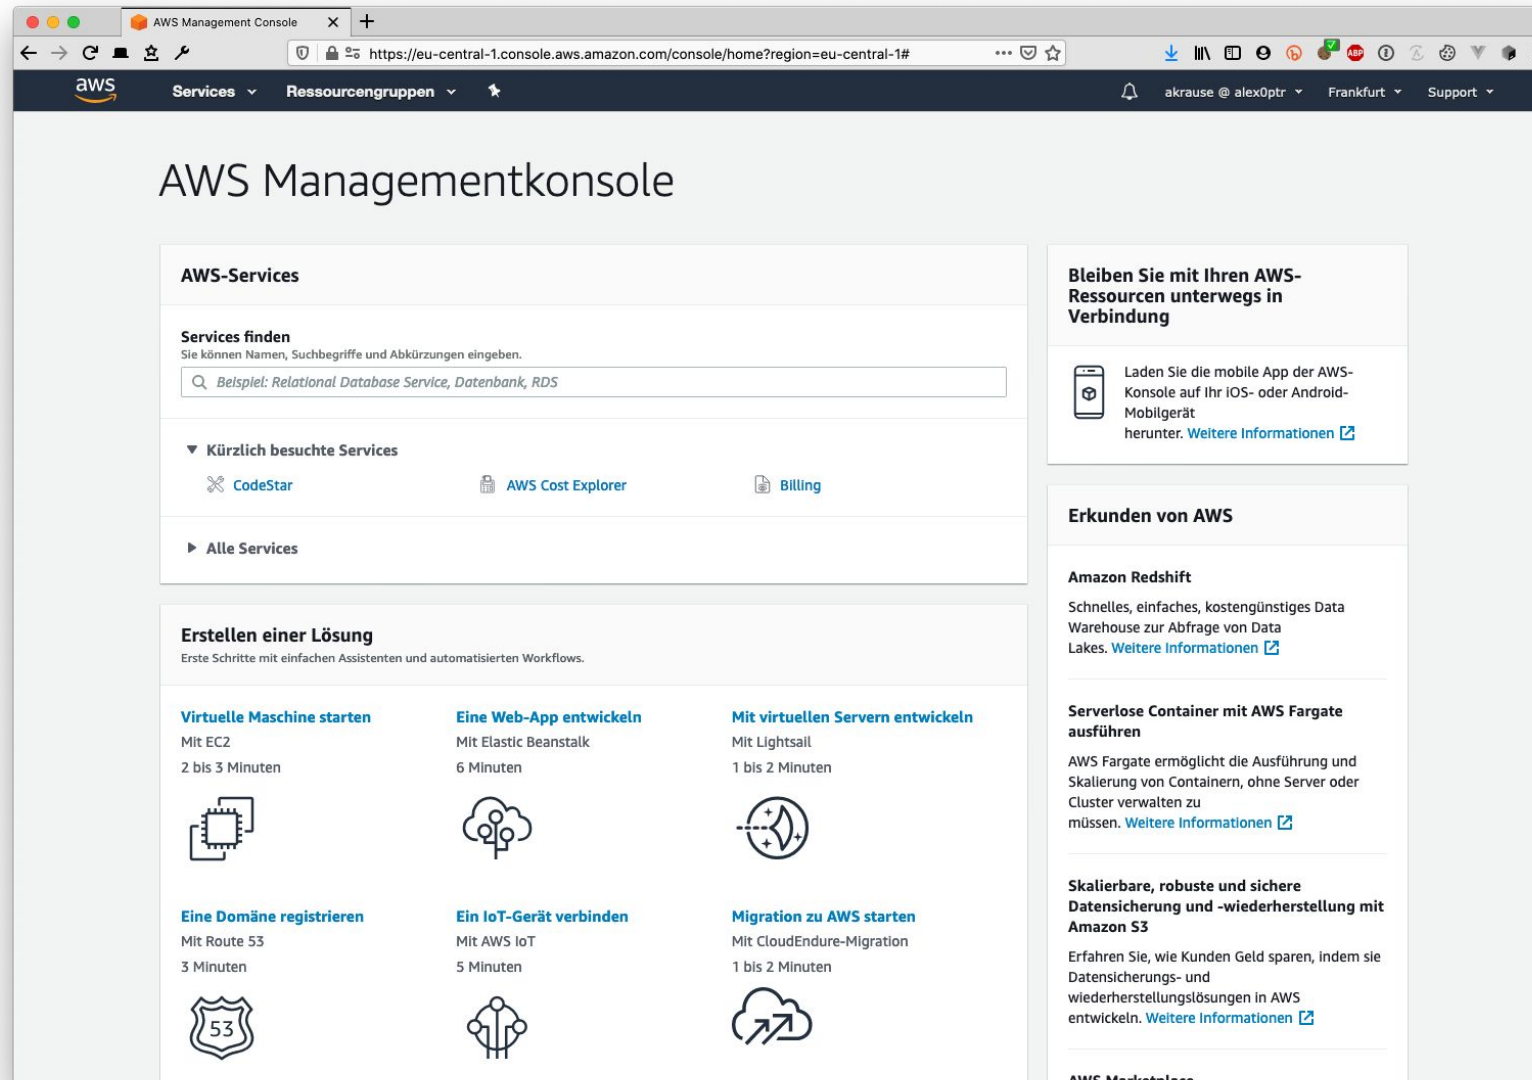
- 24 Regions
- split into 77 Availability Zones
- 216 Points of Presence

Example:

- eu-central-1 (Frankfurt)
- since 2014
- 3 Availability Zones



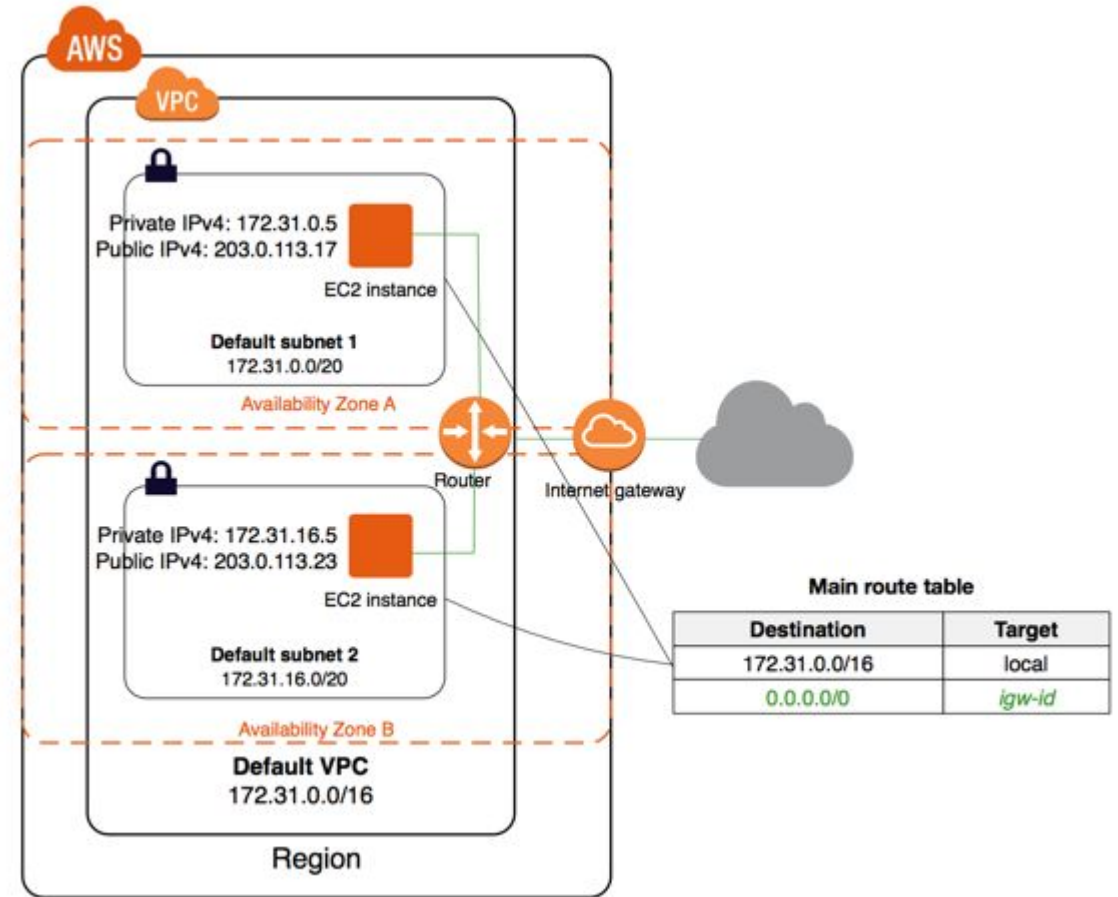
The AWS Management console allows to control all services



Amazon VPC (Virtual Private Cloud)

- Software Defined Network is an enforced prerequisite for any IaaS architecture on AWS
- Allows to provision:
 - VPCs and subnets
 - Network interfaces
 - Security Groups
 - Routing tables
 - Internet Gateways
 - NAT
 - DHCP
 - DNS
 - Elastic IP addresses
 - ...

Preconfigured Standard Network:




Amazon EC2 (Elastic Compute Cloud)

Amazon also offers an IaaS cloud as part of AWS (Amazon Web Services).

History:

- Started within Amazon in 2001
- Public beta launched on August 25, 2006
- By mid-2007, more bandwidth was consumed in the cloud by third parties than by Amazon's websites
- Production-ready as of October 23, 2008
- Approximately \$12 billion investment in infrastructure from 2005 to 2012
- 2015: 1.5 to 2 million servers in 10 global data centers.


On-demand, reserved, and spot instances in various sizes and with different hardware: [AWS EC2 Instance Types](#), as well as various storage and network services.



<http://aws.typepad.com/aws/2008/05/lots-of-bits.html>

Architecture of Amazon EC2.

- AWS Management Console
- Webservice-API
- EBS (Elastic Block Store)
- S3 (Simple Storage Service)



- VPC (Virtual Private Cloud)
- Route 53
- Elastic Load Balancer
- CloudFront CDN

- EC2-Nodes with Xen- or HVM-Virtualisierung
- Monitoring via CloudWatch
- AutoScaling based on CloudWatch-Metrics

EC2 Metadata Service

- Offered by the hypervisor it is available on every EC2-Instance under <http://169.254.169.254/latest/meta-data>

Allows an instance to query data:

- about its environment
- about itself
- user-defined metadata

Enables the implementation of advanced security and automation mechanisms:

- Instance profile via tokens to call AWS services, including automatic token exchange
- Short-lived management of SSH keys (EC2 Instance Connect)
- Tagging instance metrics in CloudWatch, etc.
- Provisioning via Cloud Init

Demonstration (5 Min): <https://www.youtube.com/watch?v=tPQsl8n6er0>

cloud-init



cloud-init

“Cloud-init is the defacto multi-distribution package that handles early initialization of a cloud instance.”

- since 2008
- Init System for the Cloud
- Initially only AWS and Ubuntu
- Now the de-facto standard across all cloud environments
- Configuration via user data from the metadata service

cloud-init - configuration options

Shell / Bash Skript:

```
#!/bin/sh
echo "Hello cloud-init!"
```

Templated Shell / Bash Skript:

```
## template: jinja
#!/bin/bash
{% if v1.region == 'us-east-2' -%}
echo 'Installing custom proxies for {{ v1.region }}'
sudo apt-get install my-xtra-fast-stack
{%- endif %}
```

Cloud-Config:

```
#cloud-config
packages:
  - cowsay
users:
  - default
  - name: app
    groups: docker
write_files:
  - content: nVc+Xj7rPhMqb...
    encoding: b64
    owner: app:app
    path: /home/app/application.yml
    permissions: '0655'
```

EC2 AutoScaling



EC2 AutoScaling enables organizing instances into **groups** as a logical unit. For example, the group can be used to control the number of instances, allowing new ones to be created automatically if some fail.



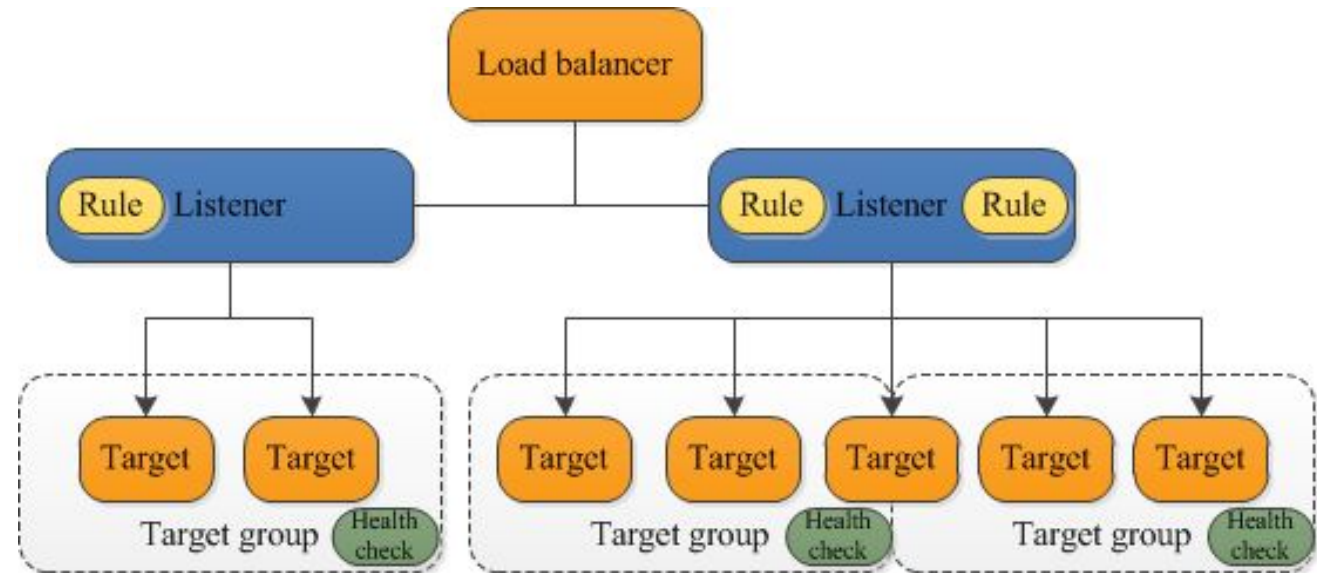
EC2 AutoScaling groups use **launch templates** to create new instances or to replace them in a rolling manner. Launch templates define instance parameters such as AMI ID, instance type, security group, or block device mappings.



Scaling options allow automatic scaling of instances based on conditions, such as CPU load, scheduled timing, or predictive scaling.

Elastic Load Balancing

- Accepts public traffic and distributes it across instances.
- Monitors the functionality of instances/applications (health check) and only routes requests/connections to "healthy" targets.
- Different variants:
 - Application Load Balancer - Layer 7
 - Network Load Balancer - Layer 4
 - Classic Load Balancer - Legacy
- Supports TLS, integration with AutoScaling, etc.

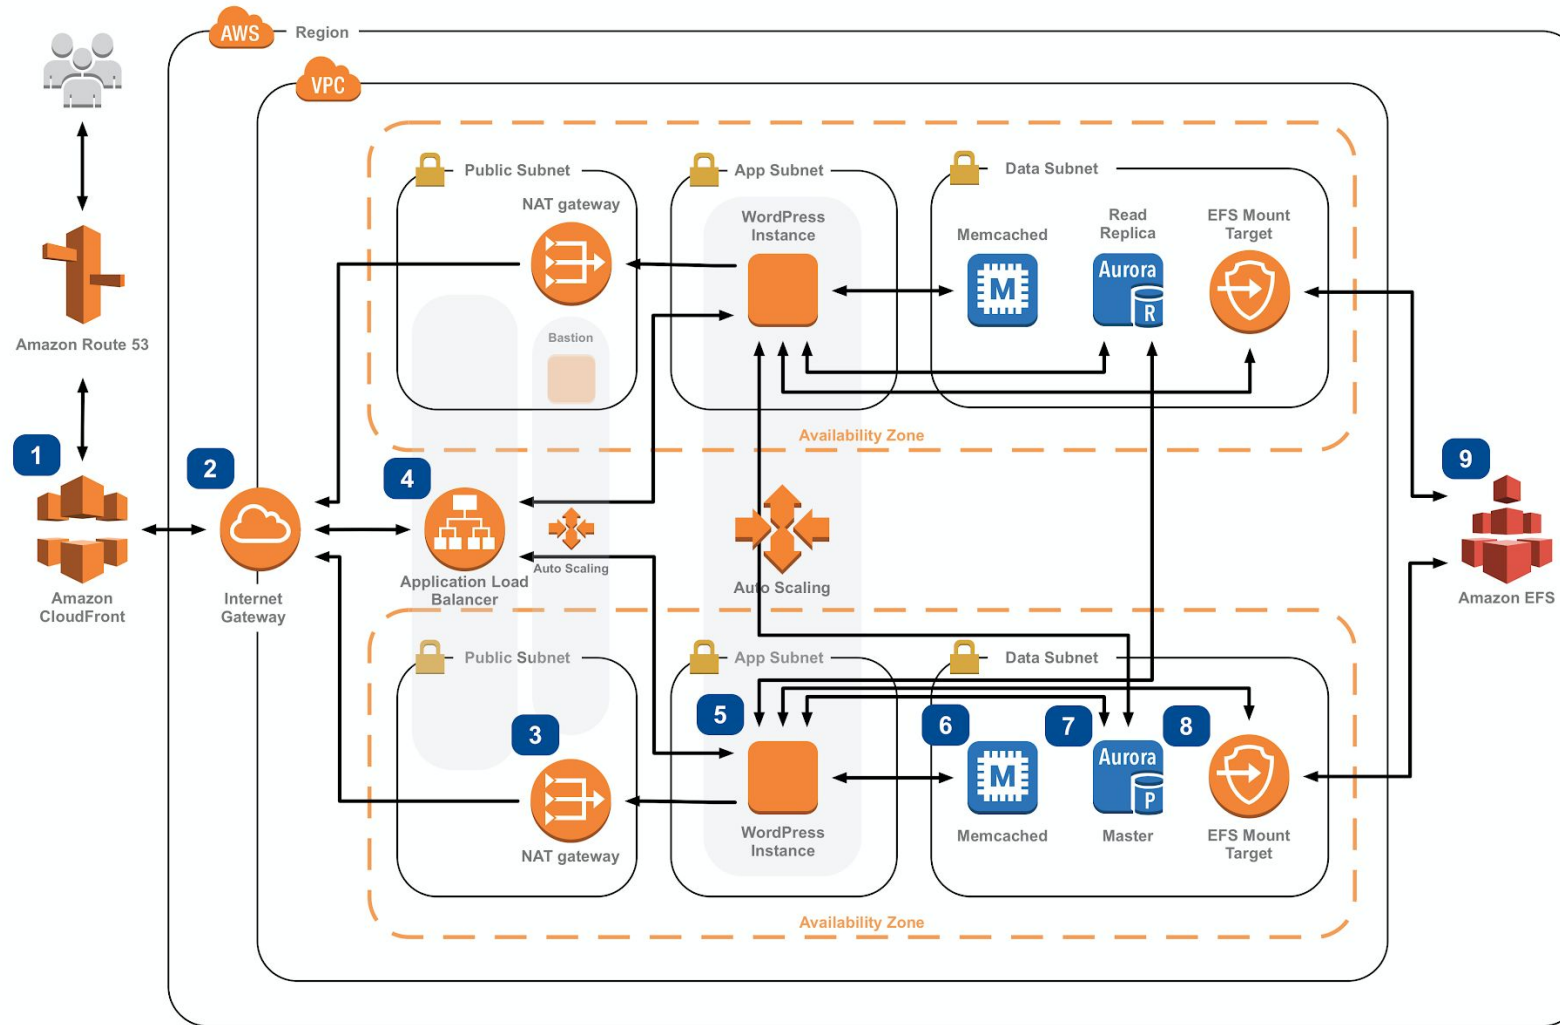


Example: Application Load Balancer

WordPress Hosting

How to run WordPress on AWS

WordPress is one of the world's most popular web publishing platforms, being used to publish 27% of all websites, from personal blogs to some of the biggest news sites. This reference architecture simplifies the complexity of deploying a scalable and highly available WordPress site on AWS.



- 1 Static and dynamic content is delivered by **Amazon CloudFront**.
- 2 An **Internet gateway** allows communication between instances in your VPC and the Internet.
- 3 **NAT gateways** in each public subnet enable Amazon EC2 instances in private subnets (App & Data) to access the Internet.
- 4 Use an **Application Load Balancer** to distribute web traffic across an Auto Scaling Group of Amazon EC2 instances in multiple AZs.
- 5 Run your WordPress site using an **Auto Scaling group of Amazon EC2 instances**. Install the latest versions of WordPress, Apache web server, PHP 7, and OPcache and build an Amazon Machine Image that will be used by the Auto Scaling group launch configuration to launch new instances in the Auto Scaling group.
- 6 If database access patterns are read-heavy, consider using a WordPress plugin that takes advantage of a caching layer like **Amazon ElastiCache (Memcached)** in front of the database layer to cache frequently accessed data.
- 7 Simplify your database administration by running your database layer in **Amazon RDS** using either Aurora or MySQL.
- 8 Amazon EC2 instances access shared WordPress data in an Amazon EFS file system using **Mount Targets** in each AZ in your VPC.
- 9 Use **Amazon EFS**, a simple, highly available, and scalable network file system so WordPress instances have access to your shared, unstructured WordPress data, like php files, config, themes, plugins, etc.



Security aspects of AWS

Certified according to ISO 27001 / C5 (recommendation from BSI) and many other standards: [AWS Compliance Programs](#)

European data centers and offices are subject to EU data protection regulations.
Amazon is also subject to the US Patriot Act and the CLOUD Act.

AWS offers services and products to meet security and compliance requirements:

- **Identity and Access Management:** IAM, Single Sign-On, Cognito...
- **Detection:** GuardDuty, Config, CloudTrail...
- **Infrastructure protection:** Shield, Web Application Firewall, Firewall Manager
- **Data protection:** KMS, CloudHSM, Macie...
- **Incident response:** Detective, CloudEndure Disaster Recovery
- **Compliance:** Artifact



QA|WARE

Infrastructure as Code

Infrastructure as Code

Provisioning and managing entire data centers—not just individual virtual machines


Distinction from configuration management (e.g., Ansible):

- Explicit creation and destruction of the infrastructure of a (virtual) data center
- Immutable infrastructure, instead of continuously modifying existing resources
- Typically declarative rather than imperative
- First introduced for the cloud in 2010 with AWS CloudFormation

Advantages:

- Versioning of the data center, enabling easy staging and rollbacks
- Accelerated delivery of infrastructure changes
- Consistency across environments
- Provides security and auditability of infrastructure in code
- Reusable and modular
- Enables collaboration through code management

Infrastructure as Code with Terraform



HashiCorp


Terraform

“Write, Plan, and Create
Infrastructure as Code”

- Since 2014, open source by Hashicorp
- Supports around 40 cloud providers
- Also integrates with database systems, monitoring, and infrastructure software such as Kubernetes
- Wide selection of plugins and reusable modules
- Declarative configuration language
- Commercial extensions available

Terraform Basics

- **Write:** Describe the target state using a domain-specific language, HCL (HashiCorp Configuration Language)
- **Plan** (terraform plan): Determine the current state. Plan necessary changes (ordered and parallelized according to dependencies, minimizing disruptions as much as possible)
- **Apply** (terraform apply): Idempotent creation of the target state. The state (.tfstate file) is usually stored in remote storage (e.g., S3, HTTP, etc.)



Video Pause at 06m:30s, then up to 10m18s:
<https://www.youtube.com/watch?v=h970ZBgKINg>

Core entities of a Terraform-configuration

Resource: managed infrastructure component

```
resource "aws_instance" "web" {  
  ami          = "ami-408c7f28"  
  instance_type = "t1.micro"  
}
```

Resources have arguments and attributes. Further they define dependencies between each other, forming a directed graph.

Provider: integration to the infrastructure/software provider e.g. AWS

Alicloud	Archive	AWS
Azure	Bitbucket	CenturyLinkCloud
Chef	Circonus	Cloudflare
CloudScale.ch	CloudStack	Cobbler
Consul	Datadog	DigitalOcean
DNS	DNSMadeEasy	DNSimple
<pre>provider "aws" { access_key = "\${var.aws_access_key}" secret_key = "\${var.aws_secret_key}" region = "us-east-1" }</pre>		
Nomad	NS1	Null
1&1	OpenStack	OpenTelekomCloud
OpsGenie	Oracle Public Cloud	Oracle Cloud Platform
OVH	Packet	PagerDuty
Palo Alto Networks	PostgreSQL	PowerDNS
ProfitBricks	RabbitMQ	Rancher
Random	Rundeck	Scaleway
SoftLayer	StatusCake	Spotinst
Template	Terraform	Terraform Enterprise
TLS	Triton	UltraDNS
Vault	VMware vCloud Director	VMware NSX-T

Provisioner: execute custom actions in the graph locally or remotely. Use only when **absolutely necessary**, in cases the provider does not fulfill your needs.

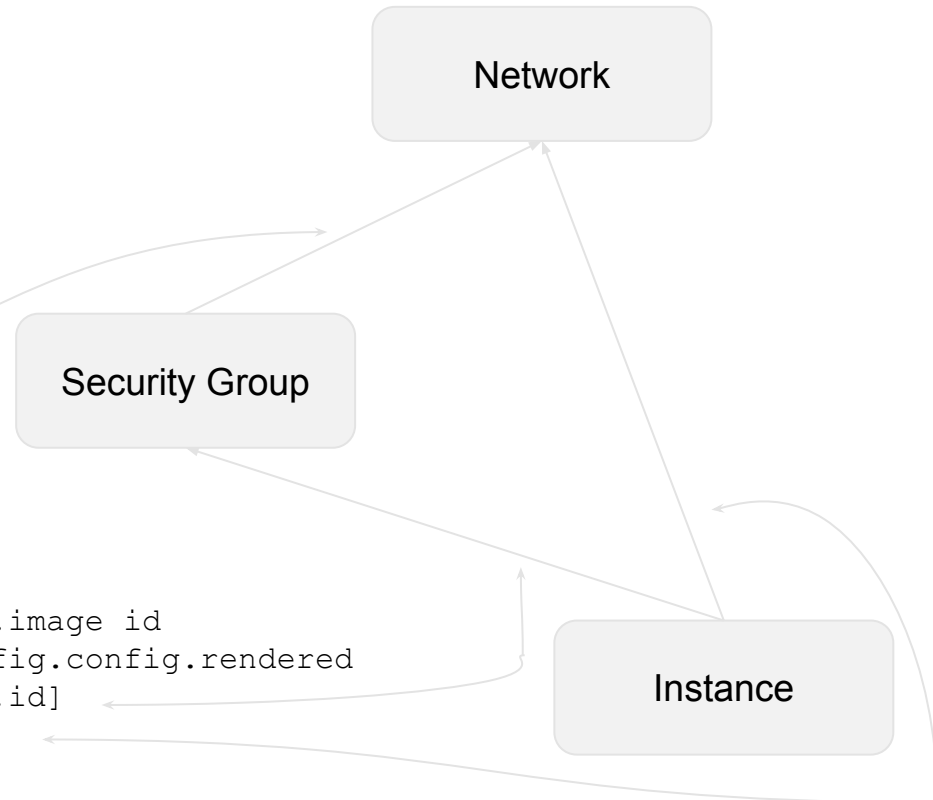
```
resource "aws_instance" "web" {  
  # ...  
  
  provisioner "local-exec" {  
    command = "echo ${self.private_ip} > file.txt"  
  }  
}
```


Example Hashicorp Configuration Language

```
module vpc {  
  source = "terraform-aws-modules/vpc/aws"  
  version = "2.18.0"  
  name    = local.env  
  # <shortened>  
}
```

```
resource aws security group bastion {  
  name          = "${local.env}-bastion"  
  description    = "For Bastion Hosts"  
  vpc id        = module.vpc.vpc_id  
  # <shortened>  
}
```


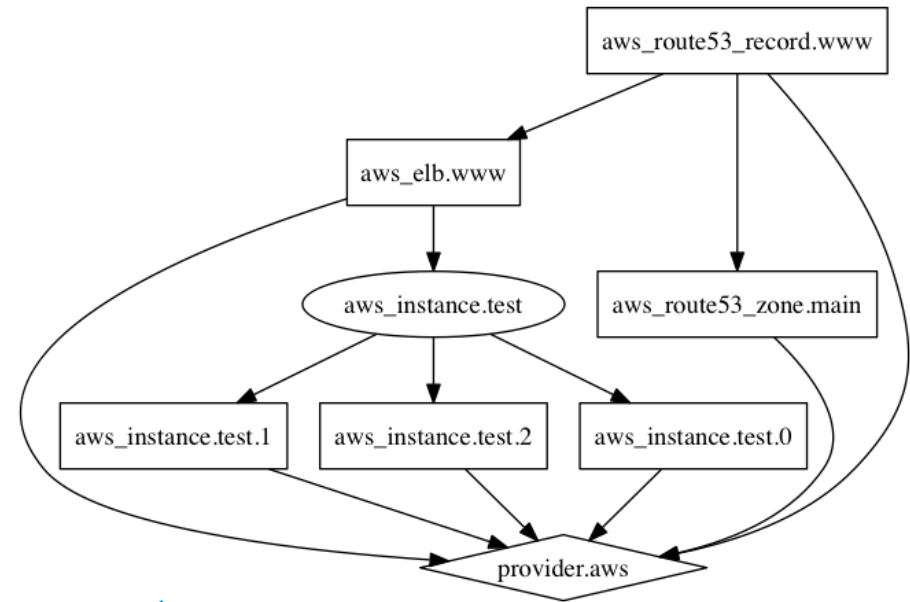
```
resource aws instance bastion {  
  instance type    = "t3.nano"  
  ami              = data.aws.ami.amazon linux 2.image id  
  user data base64 = data.template.cloudinit.config.config.rendered  
  vpc security group ids = [aws.security.group.bastion.id]  
  subnet id        = module.vpc.public_subnets[0]  
  # <shortened>  
}
```



Workflow

```
resource aws instance bastion {  
  instance type      = "t3.nano"  
  ami                = data.aws_ami.amazon_linux_2.image_id  
  user_data_base64   = data.template_cloudinit_config.config.rendered  
  vpc_security_group_ids = [aws_security_group.bastion.id]  
  subnet_id          = module.vpc.public_subnets[0]  
  # <shortened>  
}
```

```
terraform/  
├─ main.tf  
└─ terraform.tfvars
```



Terraform Deployment Levels

Level 3: Application
Deployment-packages, Data, Cron-Jobs, ...



Application Provisioning
Terraform uses Provisioner or Provider

Level 2: Software-Infrastructure
Server, virtual Machines, libraries, ...



Server Provisioning
Terraform uses Provisioner or Provider

Level 1: System-Software
Virtualization, OS, ...



Bootstrapping
Terraform uses Provider



Bare Metal Provisioning
Terraform uses Provider

Example: Provider

```
provider aws {  
  version = "2.56.0"  
  region  = "eu-central-1"  
}
```


Example: Data

```
data aws ami amazon linux_2 {  
    most_recent = true  
  
    owners = ["amazon"]  
  
    filter {  
        name      = "name"  
        values    = ["amzn2-ami-hvm*"]  
    }  
  
    filter {  
        name      = "root-device-type"  
        values    = ["ebs"]  
    }  
  
    filter {  
        name      = "architecture"  
        values    = ["x86_64"]  
    }  
}
```

Example: Resources

```
resource aws instance bastion {  
    instance type      = "t3.nano"  
    ami                = data.aws_ami.amazon_linux_2.image_id  
    user_data_base64   = data.template_cloudinit_config.config.rendered  
    vpc_security_group_ids = [aws_security_group.bastion.id]  
    subnet_id          = module.vpc.public_subnets[0]  
    tags = merge(local.standard_tags, {  
        "Name" = "${local.env}-bastion"  
    })  
}
```

<https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/instance>


Example: Output

```
output "bastion" {  
  description = "Information about the bastion host."  
  value = {  
    instance      = aws_instance.bastion.id  
    dns_name      = aws_route53_record.bastion.name  
    key_push_policy = aws_iam_policy.allow_bastion_ssh_access.arn  
  }  
}
```

Collaboration, State and Workspaces

- Terraform is stateful and usually stores this state in the cloud; local storage is possible but not recommended. This allows, for example, the removal of resources in the code to also logically remove the cloud resource. In this case, the reference in the state represents the difference from the target state.
- Remote state locking prevents the application of changes in parallel.
- The state may contain sensitive secrets, such as passwords and certificates.
- Multiple states (workspaces) for the same configuration enable:
 - easy staging, with one workspace per environment
 - completely independent development

```
terraform {  
  backend s3 {  
    bucket      = "prj-aws-dev-tfstate-storage"  
    key         = "prj-main"  
    dynamodb_table = "prj-tfstate-lock"  
    region      = "eu-central-1"  
  }  
}
```



Video ab 10m18s, bis 13m:15s:
<https://www.youtube.com/watch?v=h970ZBgKING>

Modules: Reusable abstractions

- akin to a library in the programming language of your choice
- contains complete, “final” Terraform declaration
- contains multiple resource definitions
- “Inputs” are the parameters of the module resource
- “Output” are the attributes of the module resource
- versioned
- published via <https://registry.terraform.io> or a Git repository



Video ab 13m53s, bis Ende:

<https://www.youtube.com/watch?v=h970ZBgKINg>

Example: Module

```
module vpc {
  source = "terraform-aws-modules/vpc/aws"
  version = "2.18.0"
  name = local.env
  cidr = "10.0.0.0/16"
  azs = data.aws_availability_zones.azs.names
  private_subnets = [ "10.0.0.0/19", "10.0.32.0/19", "10.0.64.0/19" ]
  public_subnets = [ "10.0.96.0/21", "10.0.104.0/21", "10.0.112.0/21" ]
  enable_nat_gateway = true
  single_nat_gateway = ! local.config.ha_nat_gateways
  one_nat_gateway_per_az = local.config.ha_nat_gateways
  enable_dns_hostnames = true
  enable_dns_support = true
  tags = merge(local.standard_tags, map( "kubernetes.io/cluster/${local.env}", "shared" ))

  enable_s3_endpoint = true
  enable_ecr_dkr_endpoint = true
  ecr_dkr_endpoint_private_dns_enabled = true
  ecr_dkr_endpoint_security_group_ids = [aws_security_group.vpc_endpoints.id]

  private_subnet_tags = {
    "kubernetes.io/role/internal-elb" = "1"
    "kubernetes.io/role/elb" = "1"
  }
}
```

Example: Directory structure

```
terraform/  
├── base/  
│   ├── vpc.tf  
│   ├── network.tf  
│   ├── variables.tf  
│   └── terraform.tfvars  
├── qa/  
│   ├── ec2.tf  
│   ├── cloudwatch.tf  
│   ├── route53.tf  
│   ├── variables.tf  
│   └── terraform.tfvars  
└── prod/  
    ├── ec2.tf  
    ├── cloudwatch.tf  
    ├── route53.tf  
    ├── variables.tf  
    └── terraform.tfvars
```

For getting your hands dirty

<https://github.com/brikis98/terraform-up-and-running-code>