



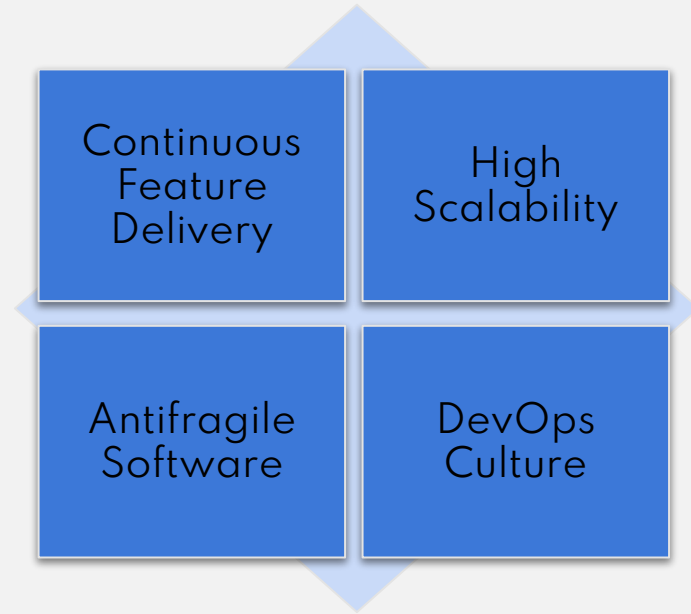
# DevOps & Continuous Delivery

Lukas Buchner  
lukas.buchner@qaware.de

# Drivers of Cloud-native applications



QA|WARE



# What is DevOps?



QA|WARE

**DevOps** is a methodology in the software development and IT industry. Used as a set of practices and tools, DevOps integrates and automates the work of **software development** (*Dev*) and **IT operations** (*Ops*) as a means for improving and shortening the **systems development life cycle**.<sup>[1]</sup> DevOps is complementary to **agile software development**; several DevOps aspects came from the *agile* approach.



... the term is used in multiple contexts. At its most successful, DevOps is a combination of specific practices, culture change, and tools.<sup>[8]</sup>

# What is DevOps?



QA|WARE

**DevOps** is the improved integration of **development** and **operations** through greater **collaboration** and **automation**, aiming to **deploy** changes to production more quickly and keep the **MTTR** (Mean Time to Recovery) low. **DevOps** is therefore a **culture**.



MVP + Feature Stream


Per Feature:

- Minimal manual post-commit effort until PROD
- Diagnosability of a feature's success
- Ability to disable/roll back the feature

# DevOps connects DEvelopment, OPerations and Quality Assurance



QA|WARE



# DevOps Topologies

Not everything labeled as DevOps truly embodies DevOps.

Numerous anti-patterns and types can be found at:  
<https://web.devopstopologies.com/>



QA|WARE

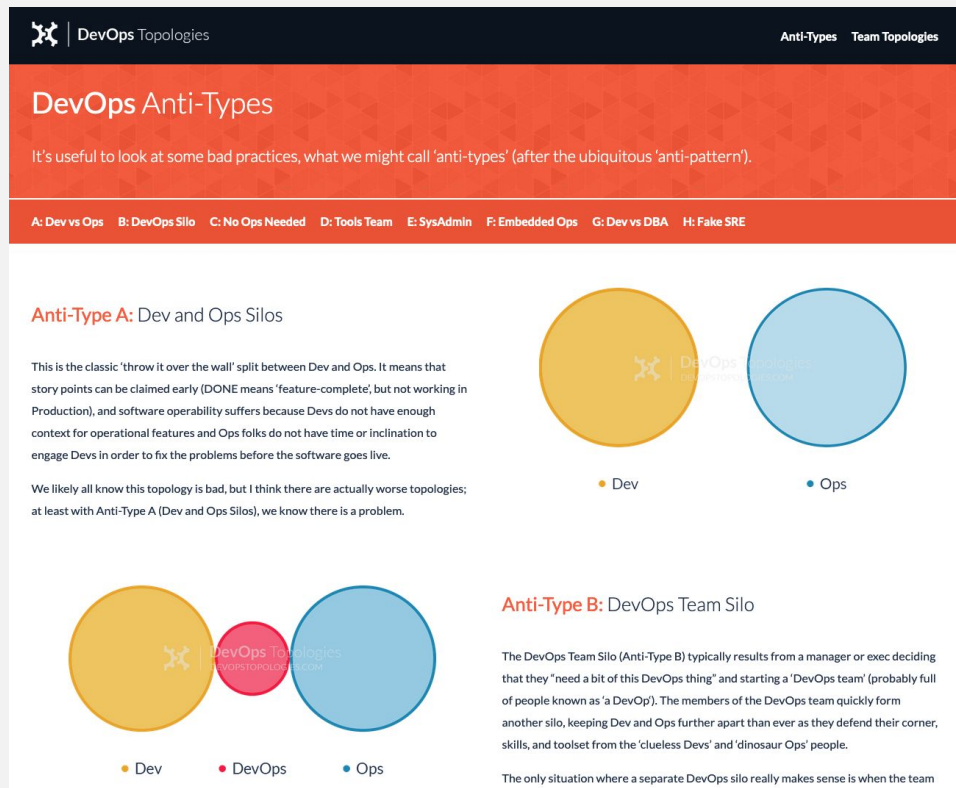


Image based on work at [devopstopologies.com](https://web.devopstopologies.com/) - licensed under [CC BY-SA](https://creativecommons.org/licenses/by-sa/4.0/).



# Pipelines



# Continuous X



QA|WARE

## Continuous Integration (CI)

- All changes are immediately integrated into the current development branch and tested.
- This ensures continuous testing of compatibility between changes.

## Continuous Delivery (CD)

- The code *can* be deployed at any time.
- However, it does not have to be deployed immediately.
- This means the code must (ideally) always build, be tested, and debugged when necessary.

## Continuous Deployment

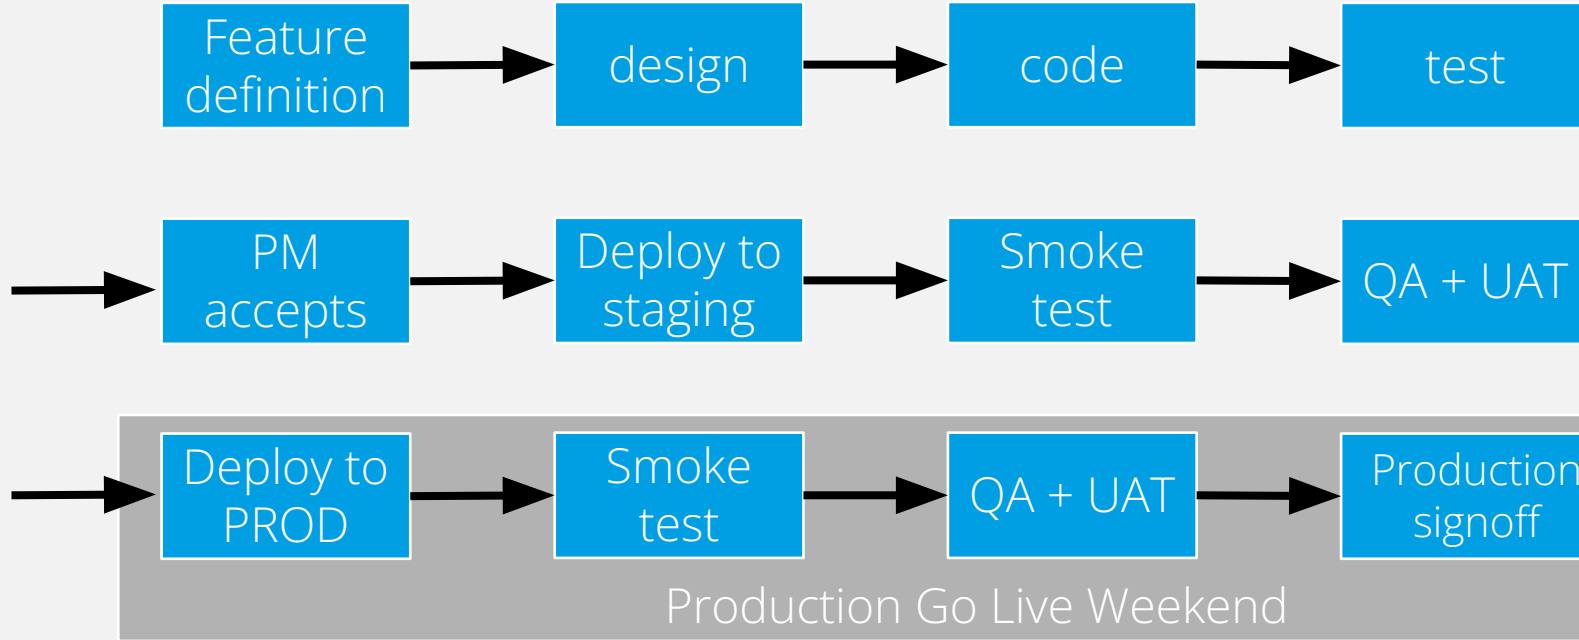
- Every stable change is deployed to production.
- Part of the quality testing takes place directly in production.
- The ability to handle errors must be in place (e.g., Canary Release, see later).



# Example: Pipeline without Continuous Delivery



QA|WARE

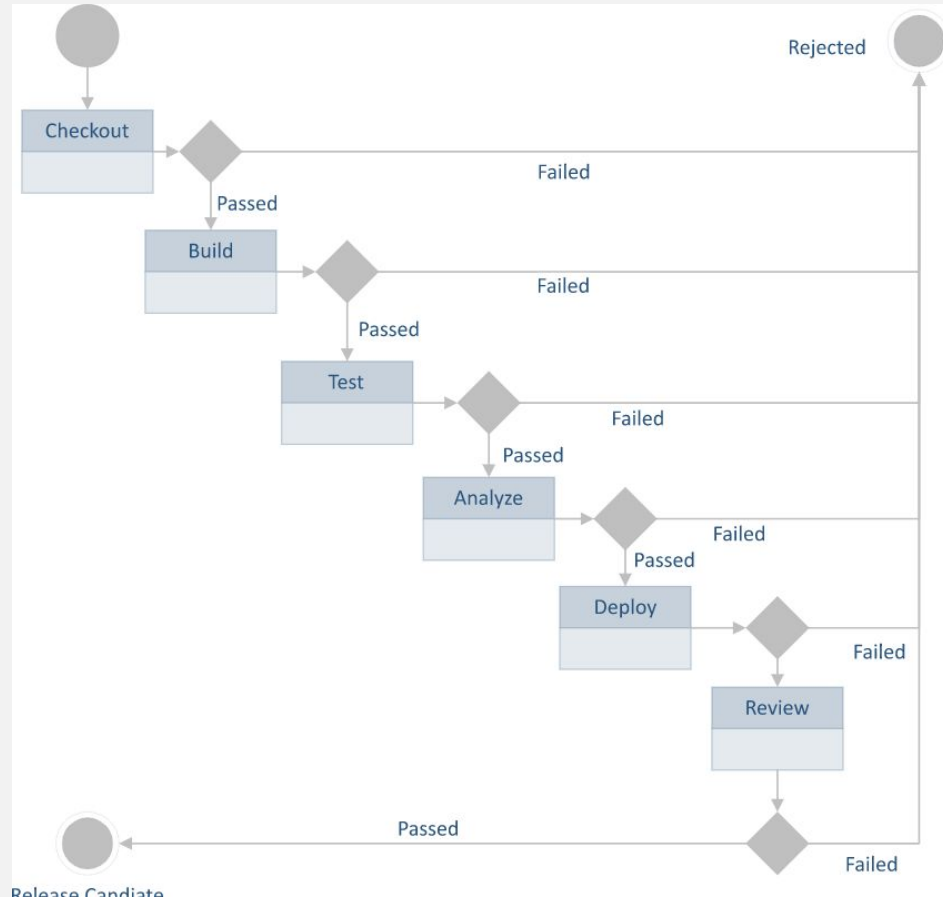


<https://www.scaledagileframework.com/continuous-delivery-pipeline/>

# Example: Continuous-Delivery-Pipeline




QA|WARE



# Shift left philosophy- give feedback to developers as early as possible



QA|WARE



„Continuous delivery process  
diagramm“  
by Grégoire Détrez  
Creative Commons 4.0

# Example of a test pyramid that aids early feedback on errors



QA|WARE

## Unit Tests:

Classic unit tests (e.g., JUnit, Mockito).

## Service Tests:

Tests for a single microservice, including REST controllers and client calls (e.g., JUnit, Spring MVC Tests, Wiremock).

Mocks for other microservices are necessary (e.g., using Wiremock).

## Integration Tests:

Tests the integration of multiple services and their interactions (e.g., JUnit, Spring MVC Tests).

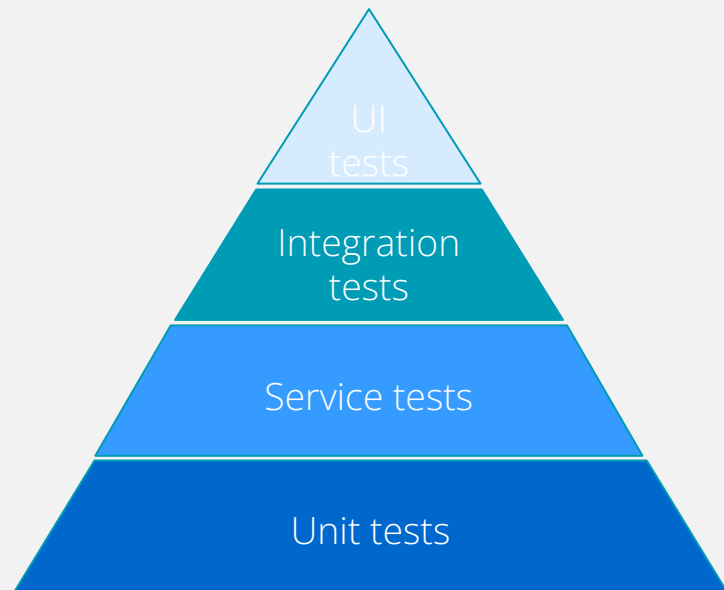
## Performance Tests:

Checks for significant performance changes (e.g., Gatling).

## UI Tests:

Tests UI functionality and its interaction with the backend (e.g., Selenium, Protractor).

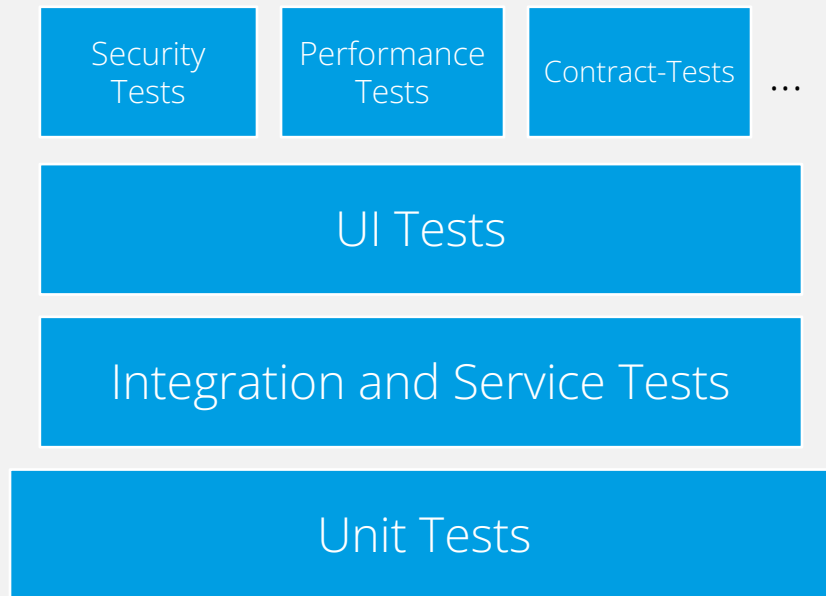
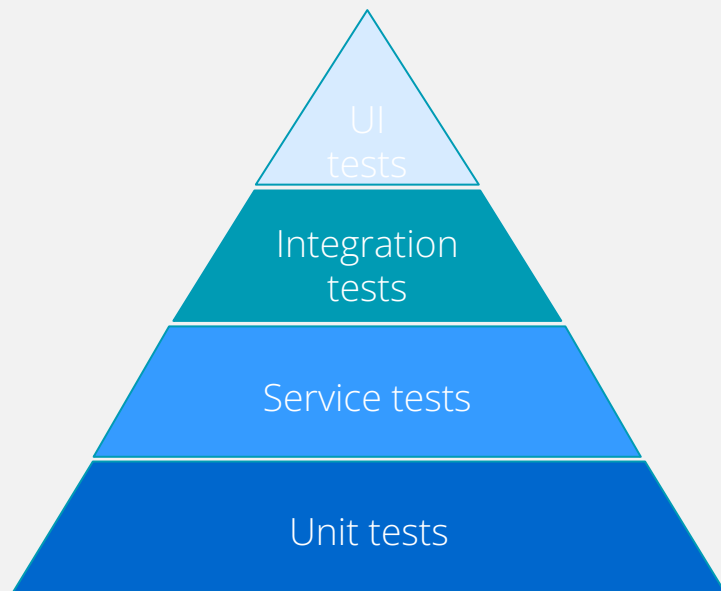
All tests should be executed as often as possible—ideally with every commit!



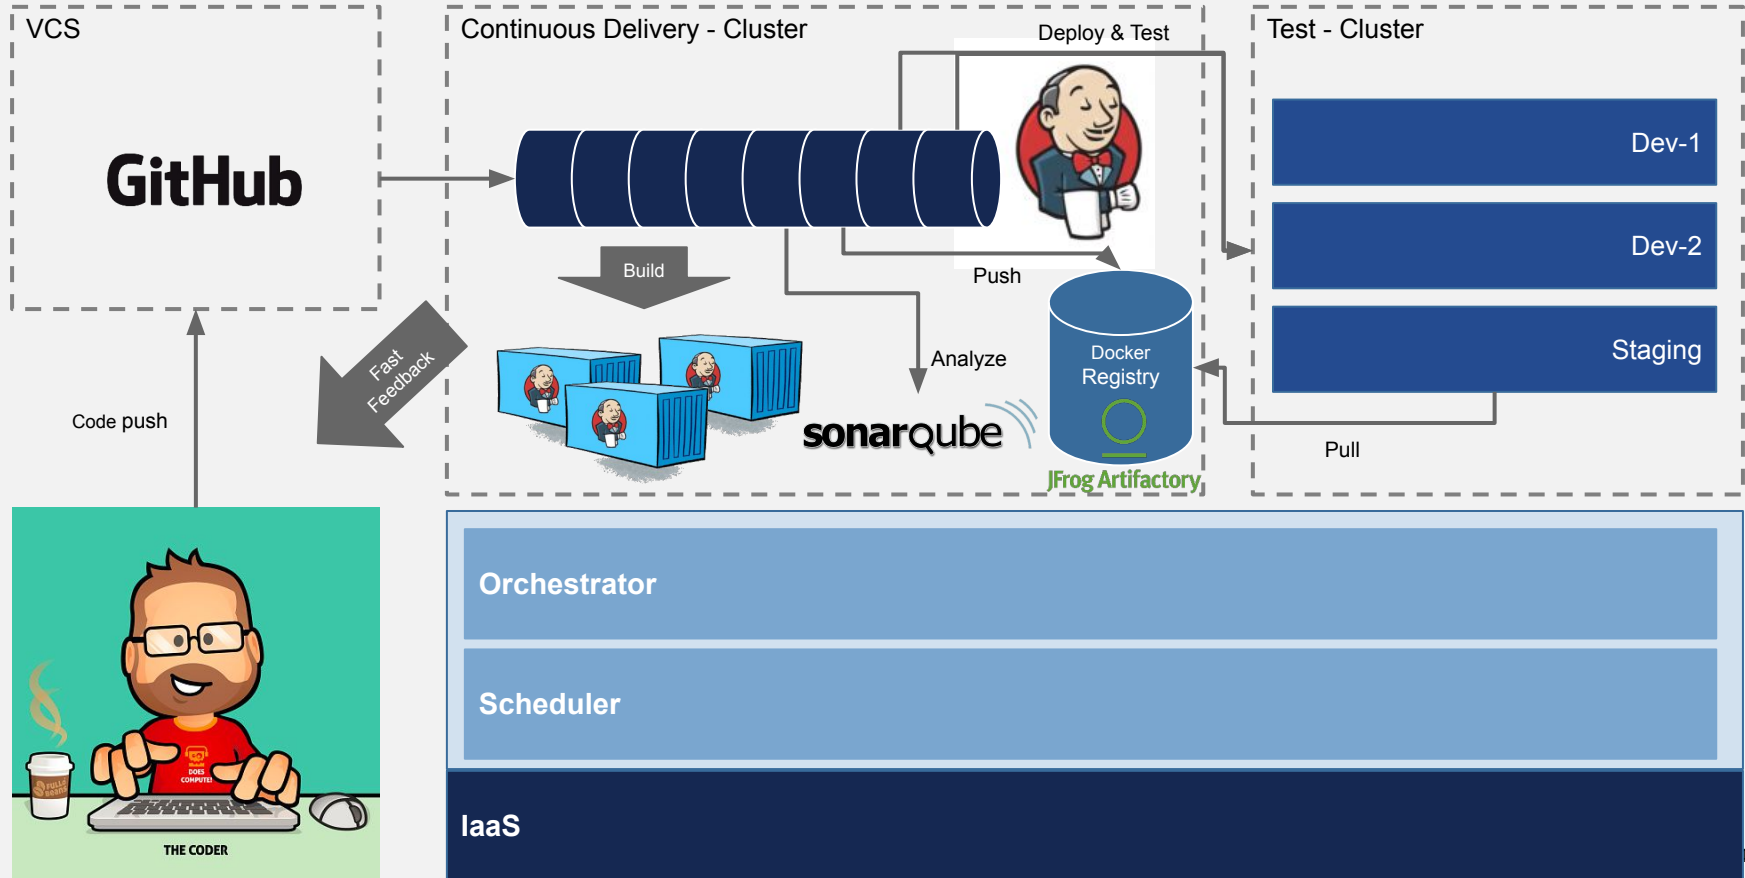
**All layers should run in an automated fashion.  
The pyramid may also come in different shapes:**



QA|WARE



# Example architecture of continuous delivery pipeline

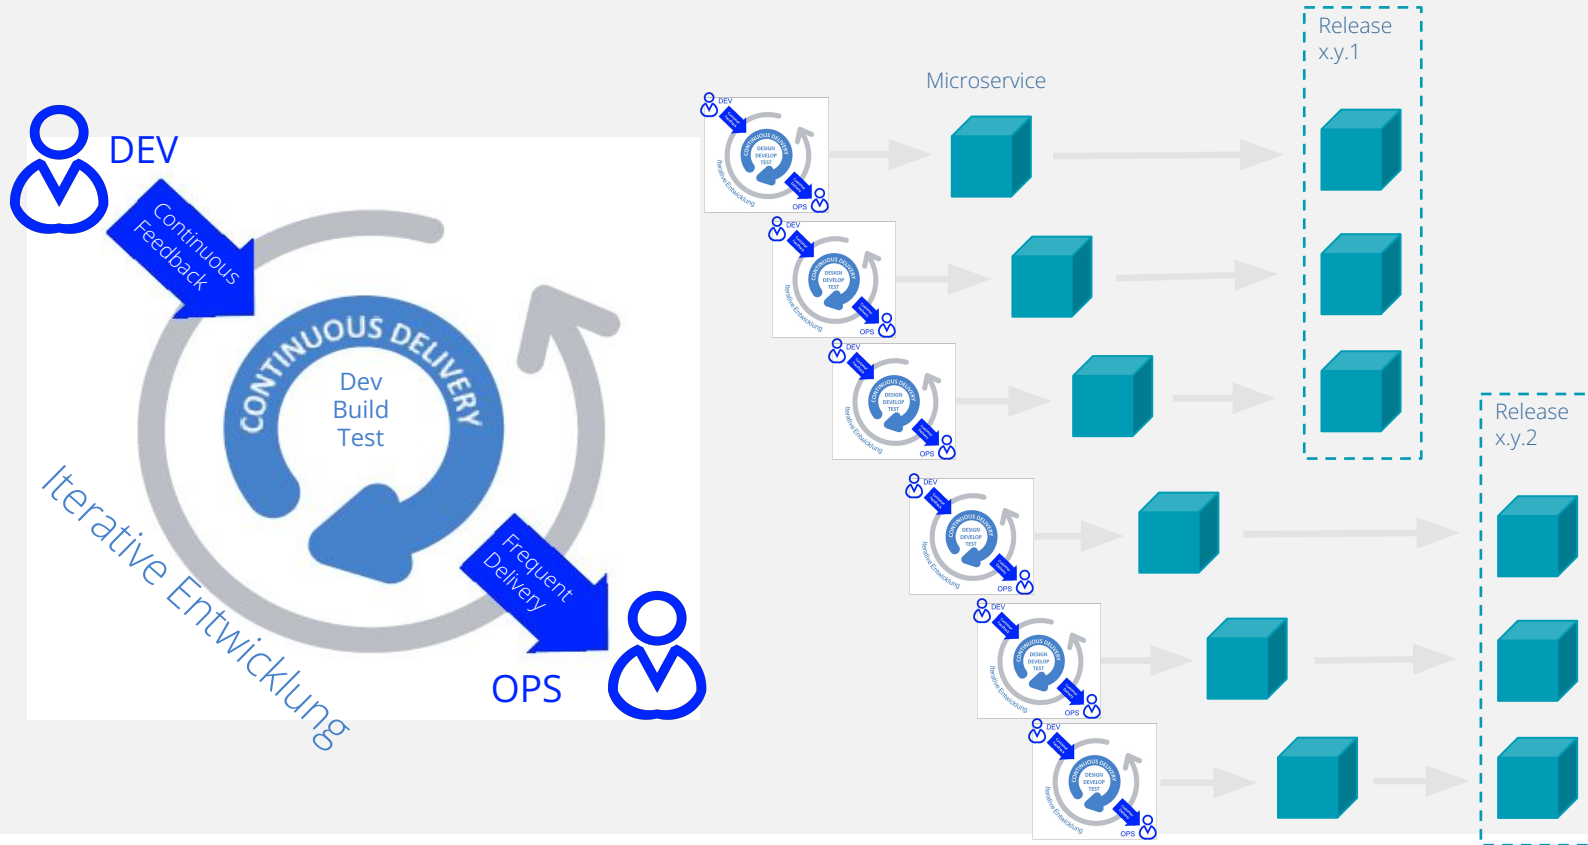




# Continuous Delivery in Microservice-Architectures:



QA|WARE



# Continuous Delivery: Building Blocks





Everything as C<>de



# Everything that brings the CD environment to life comes from the VCS



QA|WARE

## Build-as-Code

- Maven, Gradle, etc.
- Describes how the application is built.

## Test-as-Code

- Unit, component, integration, API, UI, and performance tests.
- Describes how the project is tested.

## Infrastructure-as-Code

- Docker, Terraform, Vagrant, Ansible, Kubernetes deployments.
- Describes how runtime environments are set up.


## Pipeline-as-Code

- Build pipeline via Jenkinsfile, Gitlab CI etc..
- Build wrapper: Describes all steps up to a runnable installation.



ARCHITECTURAL DRAWINGS

# Self-Service und Blueprints



# Blueprints and Templates aid the Setup



QA|WARE

Blueprints & Templates:

The build pipelines of (micro)services within a project are often very similar.

Blueprints and templates provide a framework and implicitly establish conventions.

Examples:

- Using the Jenkins Pipeline Multibranch Plugin automatically creates a dedicated pipeline for each branch.
- Consistent integration with platform components (e.g., SonarQube, Artifactory).

```
stages {  
  
    stage('Send Build started Notification') {  
        steps {  
            slackSend (color: '#FFFF00', message: "STARTED: Job '${env.JOB_NAME}' [${env.BUILD_URL}]")  
        }  
    }  
  
    stage('Build project') {  
        steps {  
            sh './gradlew clean build --info --no-daemon'  
        }  
    }  
  
    stage('Unit Test Reporting') {  
        steps {  
            junit allowEmptyResults: true, testResults: '**/build/test-results/*.xml'  
        }  
    }  
}
```





QA|WARE

# GitOps

# The Idea



QA|WARE

## 1 Declarative

A system managed by GitOps must have its desired state expressed declaratively.

## 2 Versioned and Immutable

Desired state is stored in a way that enforces immutability, versioning and retains a complete version history.

## 3 Pulled automatically

Software agents automatically pull the desired state declarations from the source.

## 4 Continuously reconciled

Software agents continuously observe actual system state and attempt to apply the desired state.

# Advantages gained out of the box



QA|WARE

- Ideally allows restoring any system state in history, e.g., for easy rollbacks.
- Enforces pipelines.
- Provides transparency for changes; in the case of Git, it also shows who made changes.
- Ensures that the system state does not deviate from the desired state.

# GitOps in the K8s ecosystem



QA|WARE



ArgoCD




Flux

# Continuous Integration



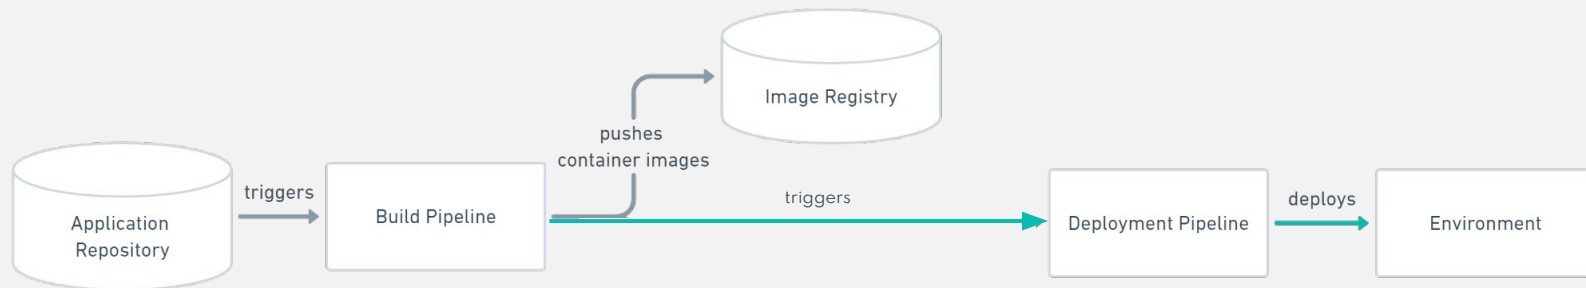
QA|WARE



# Continuous Integration with automated deployments



QA|WARE



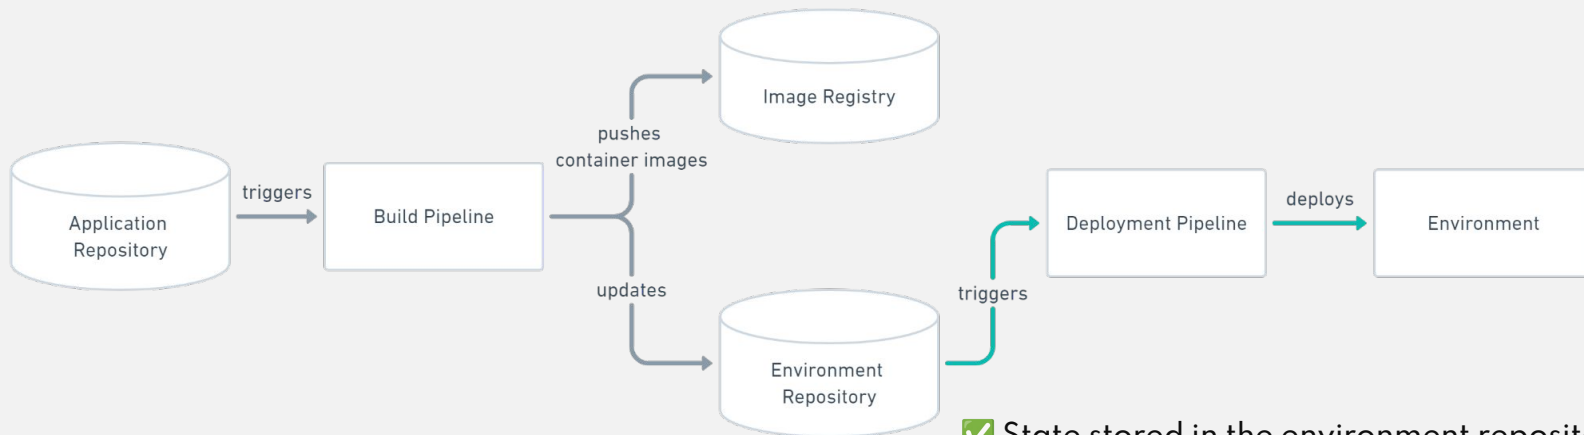
- ✓ Automated deployments
- ✗ Deployment pipeline requires extensive permissions
- ✗ Execution is often not idempotent
- ✗ Cluster state is not fixed



# Push-based deployments with GitOps



QA|WARE

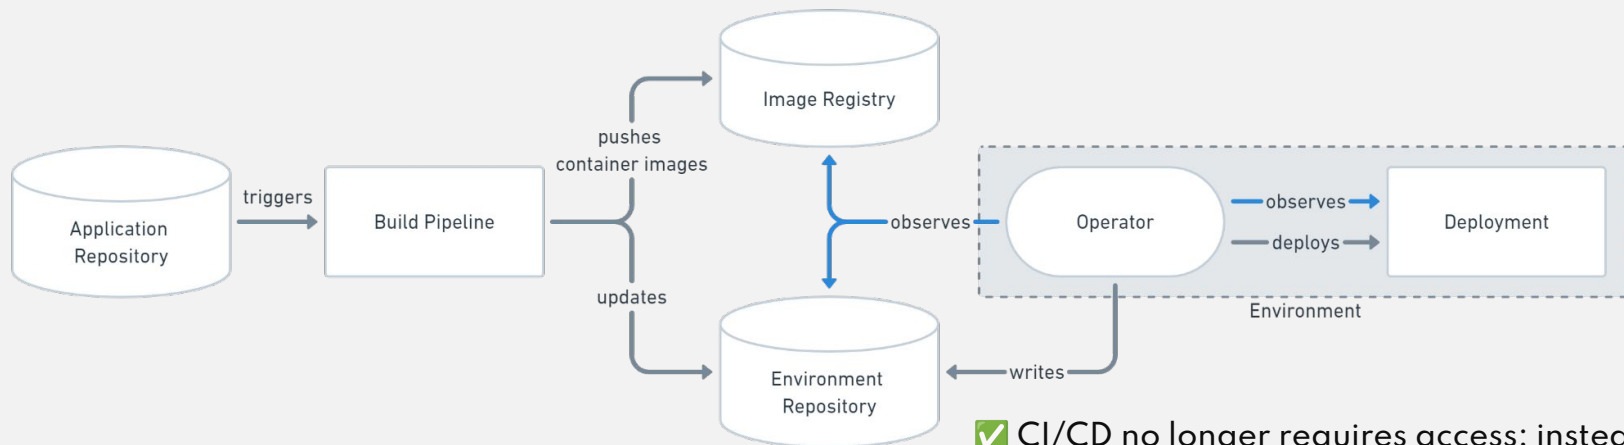


- ✓ State stored in the environment repository and can be restored
- ✗ Deployment pipeline requires extensive permissions
- ✗ Deviations in the cluster outside of a deployment are not detected or corrected
- ✗ Deleting resources is difficult

# Pull-based deployments mit GitOps



QA|WARE



✓ CI/CD no longer requires access; instead, the operator accesses the environment repo and registry.

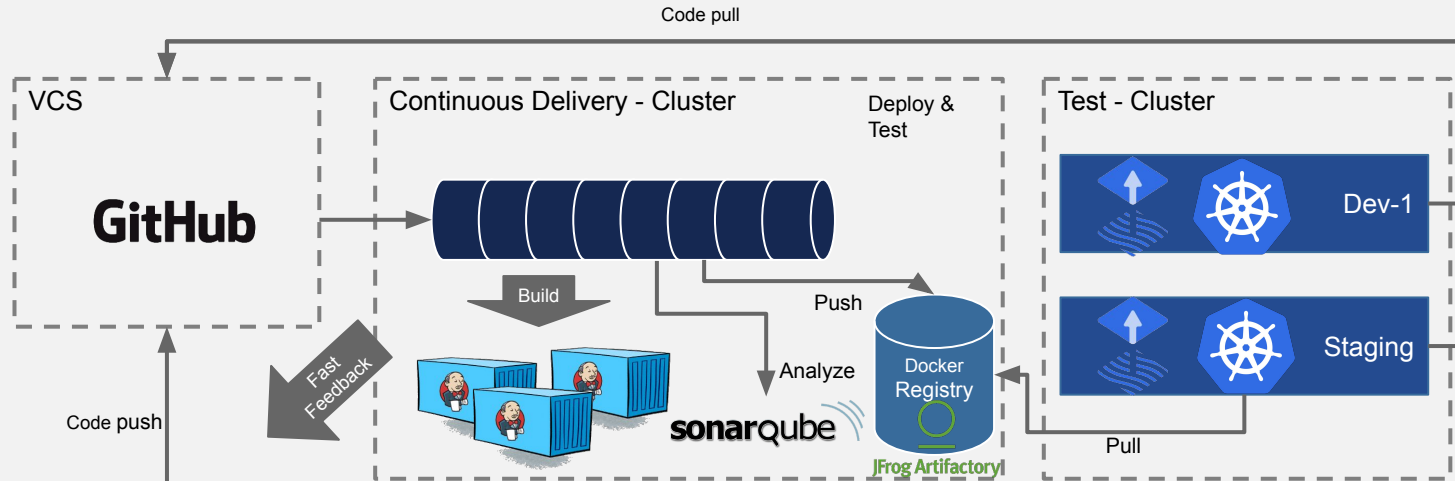
✓ The operator attempts to establish the desired cluster state, including resource deletion.

✗ More complex

# GitOps with Flux



QA|WARE



# Differences in the architecture



QA|WARE

CI/CD pipeline no longer requires cluster access

=> Improved security, as no high-privilege credentials are needed for the pipeline.

Clusters now actively pull source repositories that contain the desired state.

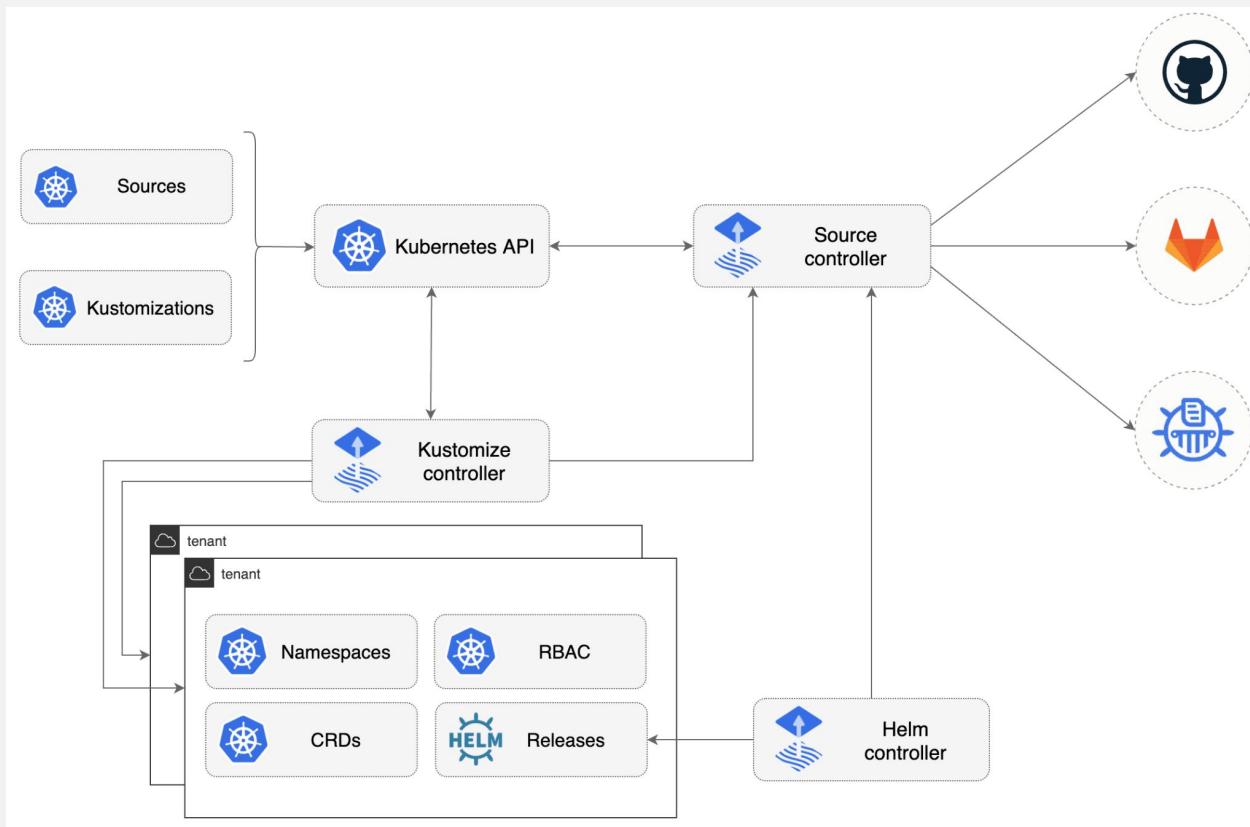
Clusters reconcile their state, meaning the current state is compared with the desired state, and the cluster attempts to reach the desired state.

=> Clusters automatically converge toward the desired state.

# Flux in detail



Q|WARE



# Flux in detail



Q|WARE

## Source Controller

- Connects to various sources, e.g., Git repositories, OCI repositories, etc.
- Provides the downloaded packages to other controllers.

## Helm Controller

- Offers Helm integration for Flux.
- In combination with the Source Controller, Helm charts can be automatically downloaded, installed, and updated.

## Kustomize Controller

- Uses the Git repositories managed by the Source Controller to execute Kustomize and apply the configured resources to the cluster.
- Tracks all resources, enabling automated garbage collection of unused resources.

## Notification Controller

- Provides insights into Flux events, such as successes or errors.
- Includes integrations with chat systems.



# Configure a source



QAWARE

```
apiVersion: source.toolkit.fluxcd.io/v1
kind: GitRepository
metadata:
  name: podinfo
  namespace: default
spec:
  interval: 5m0s
  url: https://github.com/stefanprodan/podinfo
  ref:
    branch: master
```

- Configures a GitRepository as a source in the Source Controller.
- Clones the master branch from the URL and makes it available to other controllers as a tar.gz file.
- Polls the URL every 5 minutes to check for changes.

# Configure a kustomization



QAWARE

```
apiVersion: kustomize.toolkit.fluxcd.io/v1
kind: Kustomization
metadata:
  name: podinfo
  namespace: default
spec:
  interval: 10m
  targetNamespace: default
  sourceRef:
    kind: GitRepository
    name: podinfo
  path: "./kustomize"
  prune: true
  timeout: 1m
```

Creates a Kustomization in the Kustomize Controller that references the GitRepository from the Source Controller via **sourceRef**.

The Kustomize Controller retrieves the source code from the Source Controller and uses Kustomize (a Kubernetes tool) to apply everything under the specified **path**.



QA|WARE

# Flux in action!

# Deployment Strategies



# Rolling Upgrades



QA|WARE

## **Strategy:**

Update one instance at a time, before proceeding with the next one.

## **Advantages:**

- Easily understandable
- Easily achievable with Kubernetes
- Does not require fancy monitoring infrastructure

## **Disadvantages:**

- Requires special care regarding database migrations as the old and new version running at the same time. If not done right, the application is flaky for users during deployment.

# Blue/Green Deployments



QA|WARE

## Strategy:

For the time of the deployment the old and new software versions are running in side by side. Once ensured, that the new version meets all requirements, the load balancer switches over to the new version.

## Advantages:

- Easily understandable
- Rather easy to achieve with Kubernetes
- Does not require fancy monitoring infrastructure

## Disadvantages:

- Requires special care regarding database migrations as the old and new version running at the same time. If not done right, the application is flaky for users during deployment.

# Canary Deployments



QA|WARE

## Strategy:

For the time of the deployment the old and new software versions are running in side by side. Gradually a part of the traffic (e.g. 5%) is routed to the new version. If there are no errors in monitoring, more traffic is shifted. The deployment is done, when 100% of traffic hits the new software version.

## Advantages:

- Theoretically ideal, as bad software versions are never rolled out fully. Bad software impacts a small amount of users only

## Disadvantages:

- Requires special tooling (e.g. monitoring infrastructure)
- Hard to really get right. What happens if there's no continuous traffic incoming during deployment?
- Requires special care regarding database migrations as the old and new version running at the same time. If not done right, the application is flaky for users during deployment.