



# Platform-as-a-Service (PaaS) & Internal Developer Platforms

Lukas Buchner

lukas.buchner@qaware.de






QA|WARE

# PaaS-Cloud Basics

# The layered model of cloud computing: From metal to application.



QA|WARE



- Customizable software services
- CaaS: Component as a service (e.g. Google Charts)
- BaaS: Backend as a Service
- Transparent Updates
- Platform services
- Application programming interfaces (APIs)
- Abstraction of technical infrastructure
- Elasticity
- Virtual resourcepools
- Technical infrastructure: Machines, Servers (DNS, DHCP, LB, NAS, ...)
- Computers
- Network
- Storage

**Target Group:  
Users**

**Target group:  
Developers**

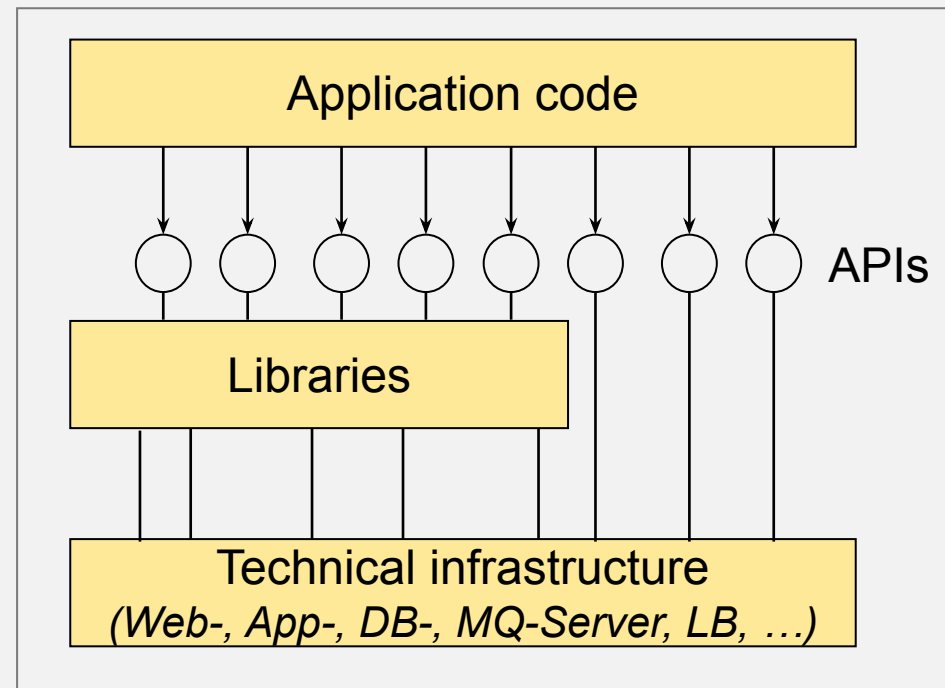
**Target group:  
Operations**

[https://www.youtube.com/watch?v=M988\\_fsOSWo](https://www.youtube.com/watch?v=M988_fsOSWo)

# The Problem: Stovepipe Architecture.

## Integration is time consuming.

The System: connected with great manual efforts



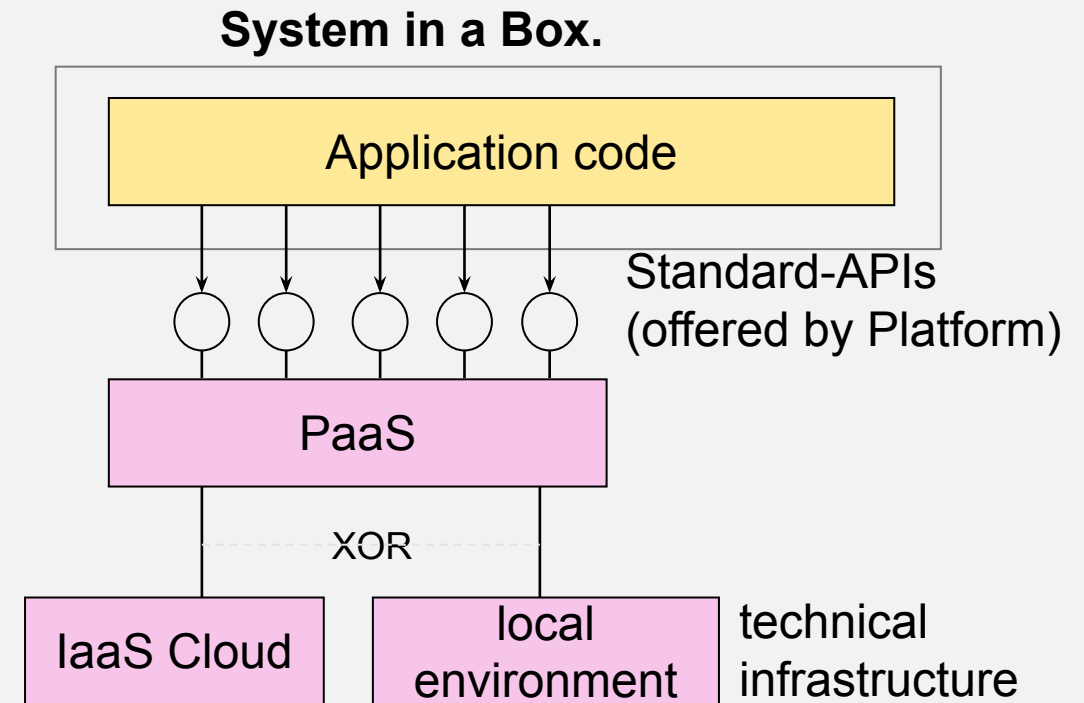
# Idea: Platform-as-a-Service offers an ad-hoc Development- and Operations Platform.



The application is deployed either as an application package or as source code. No image with technical infrastructure is required.

The application only interacts with programming or access interfaces of its runtime environment.

“Engine and operating system should not matter...”  
The application is automatically scaled.

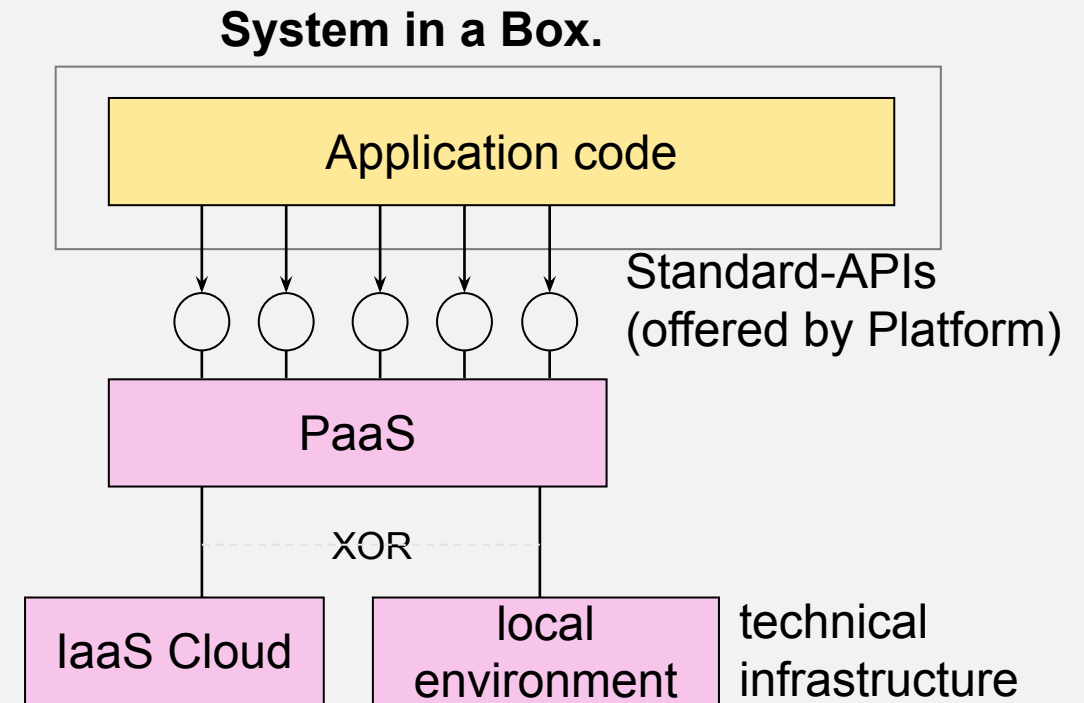


# Idea: Platform-as-a-Service offers an ad-hoc Development- and Operations Platform.



Development tools (especially plugins for IDEs and build systems, as well as a local testing environment) are available:  
“deploy to cloud.”

The platform provides an interface for administration and monitoring of applications.





# PaaS: Definitions

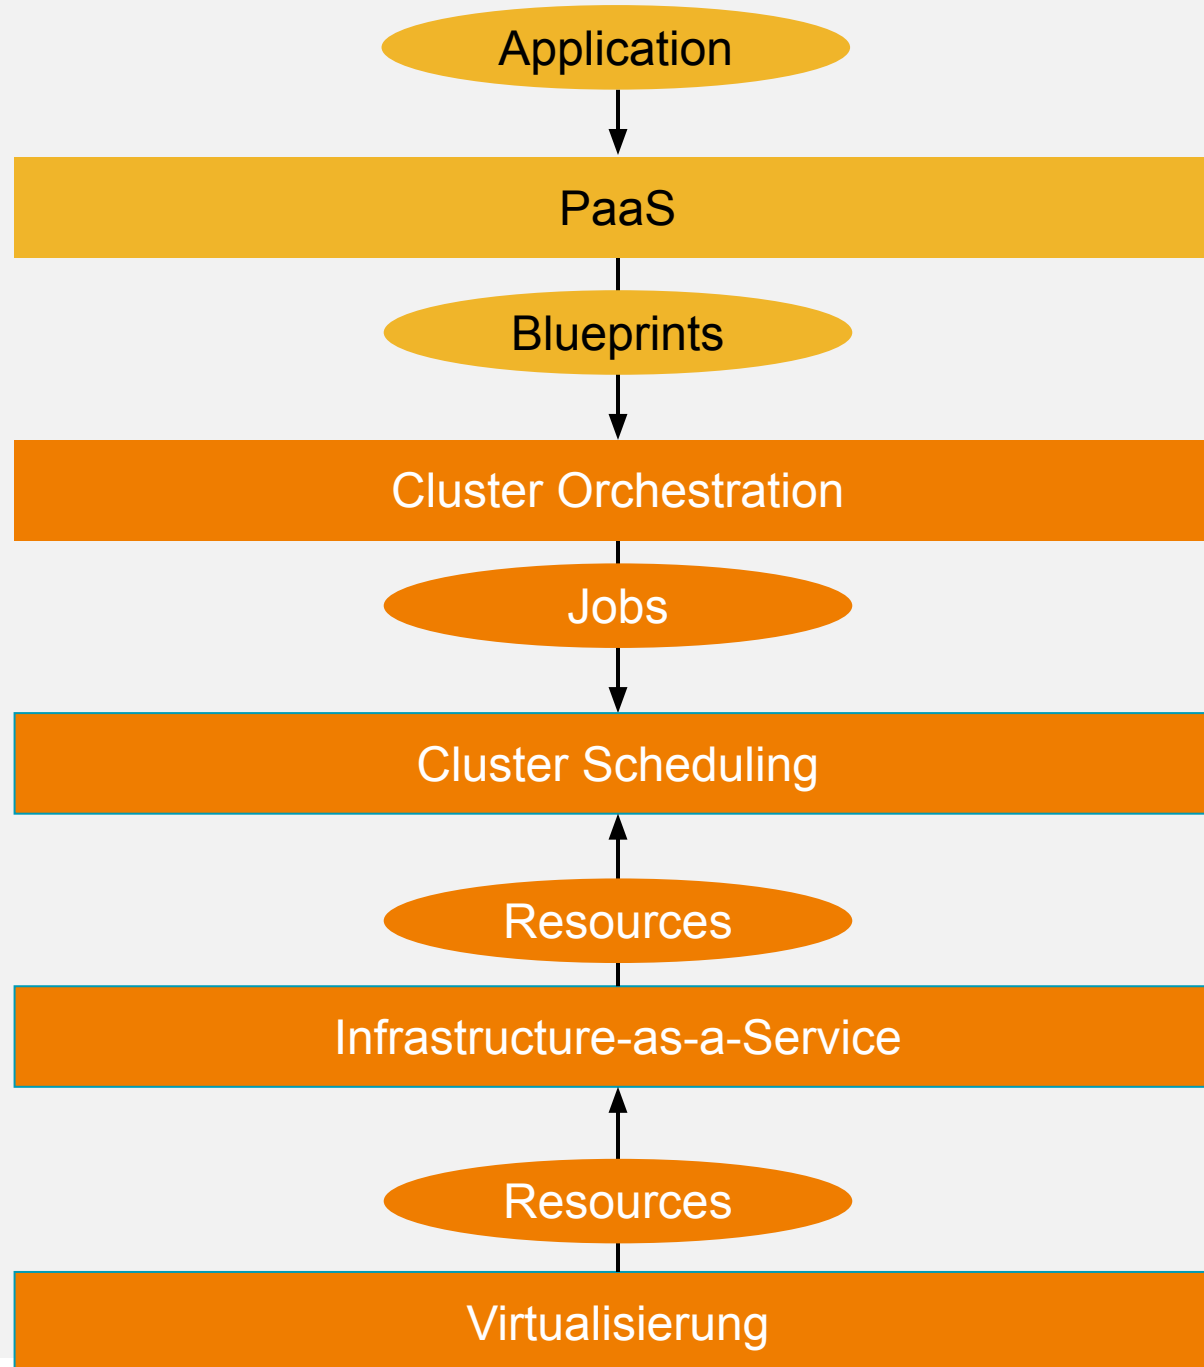


QA|WARE

NIST: The capability provided to the consumer is to **deploy onto the cloud infrastructure** consumer-created or acquired applications created **using programming languages, libraries, services, and tools supported by the provider**. The **consumer does not manage or control the underlying cloud infrastructure** including network, servers, operating systems, or storage, but **has control over the deployed applications** and possibly configuration settings for the application-hosting environment.

Forrester: A **complete application platform** for multitenant cloud environments that **includes development tools, runtime, and administration** and management tools and services. PaaS **combines an application platform with managed cloud infrastructure** services.

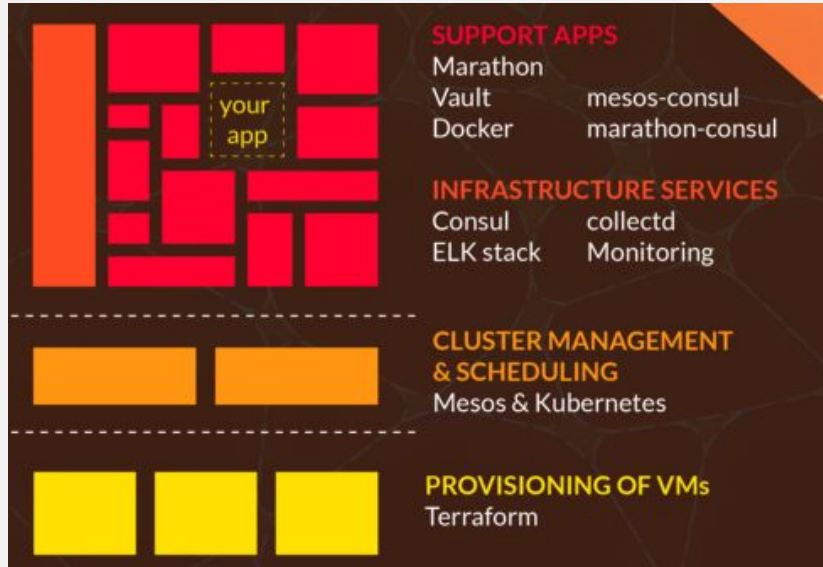
# The Big Picture





# Building-Blocks of PaaS-Solutions:

## Do you see similarities?



Quelle: <https://mantl.io>

Quelle: <https://github.com/yelp/paasta>

*Note:* PaaS is an opinionated platform that uses a few un-opinionated tools. It requires a non-trivial amount of infrastructure to be in place before it works completely:

- [Docker](#) for code delivery and containment
- [Mesos](#) for code execution and scheduling (runs Docker containers)
- [Marathon](#) for managing long-running services
- [Chronos](#) for running things on a timer (nightly batches)
- [SmartStack](#) for service registration and discovery
- [Sensu](#) for monitoring/alerting
- [Jenkins](#) (optionally) for continuous deployment



QA|WARE



Apollo is built on top of the following components:

- [Packer](#) for automating the build of the base images
- [Terraform](#) for provisioning the infrastructure
- [Apache Mesos](#) for cluster management, scheduling and resource isolation
- [Consul](#) for service discovery, DNS
- [Docker](#) for application container runtimes
- [Weave](#) for networking of docker containers
- [HAProxy](#) for application container load balancing

Quelle: <https://github.com/Capgemini/Apollo>

Cloud-ready Software Architecture

Cluster Orchestration

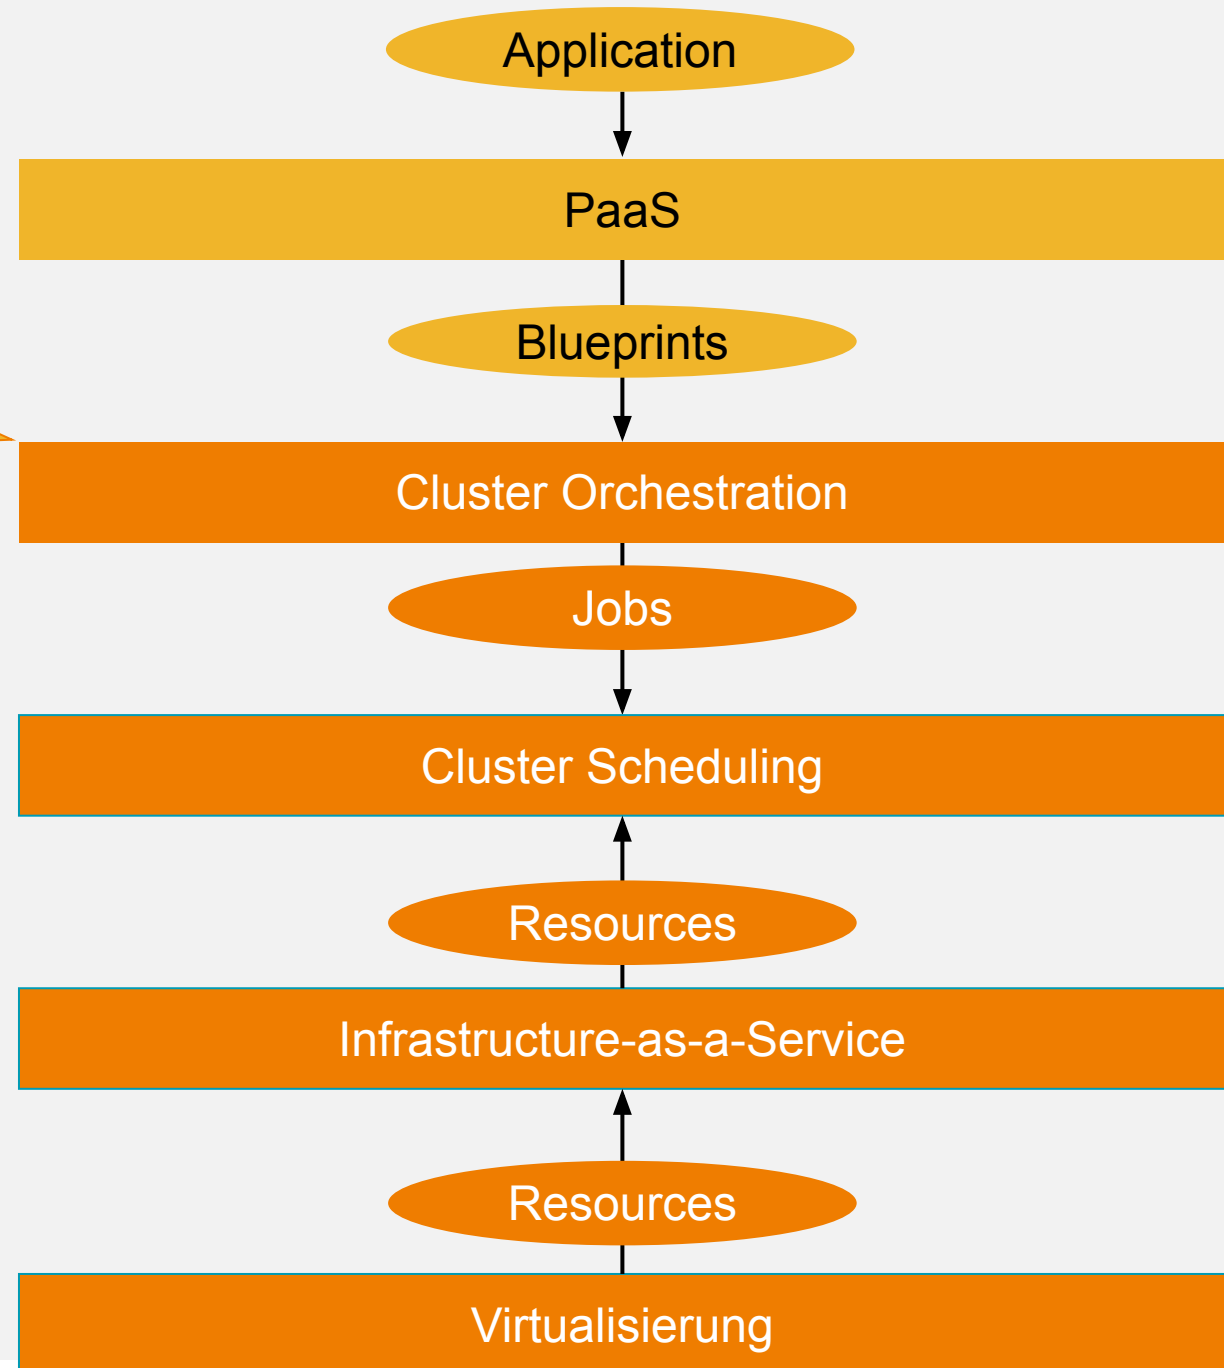
Cluster Scheduling

# The Big Picture




This is already 80% of a PaaS.  
What's still missing:

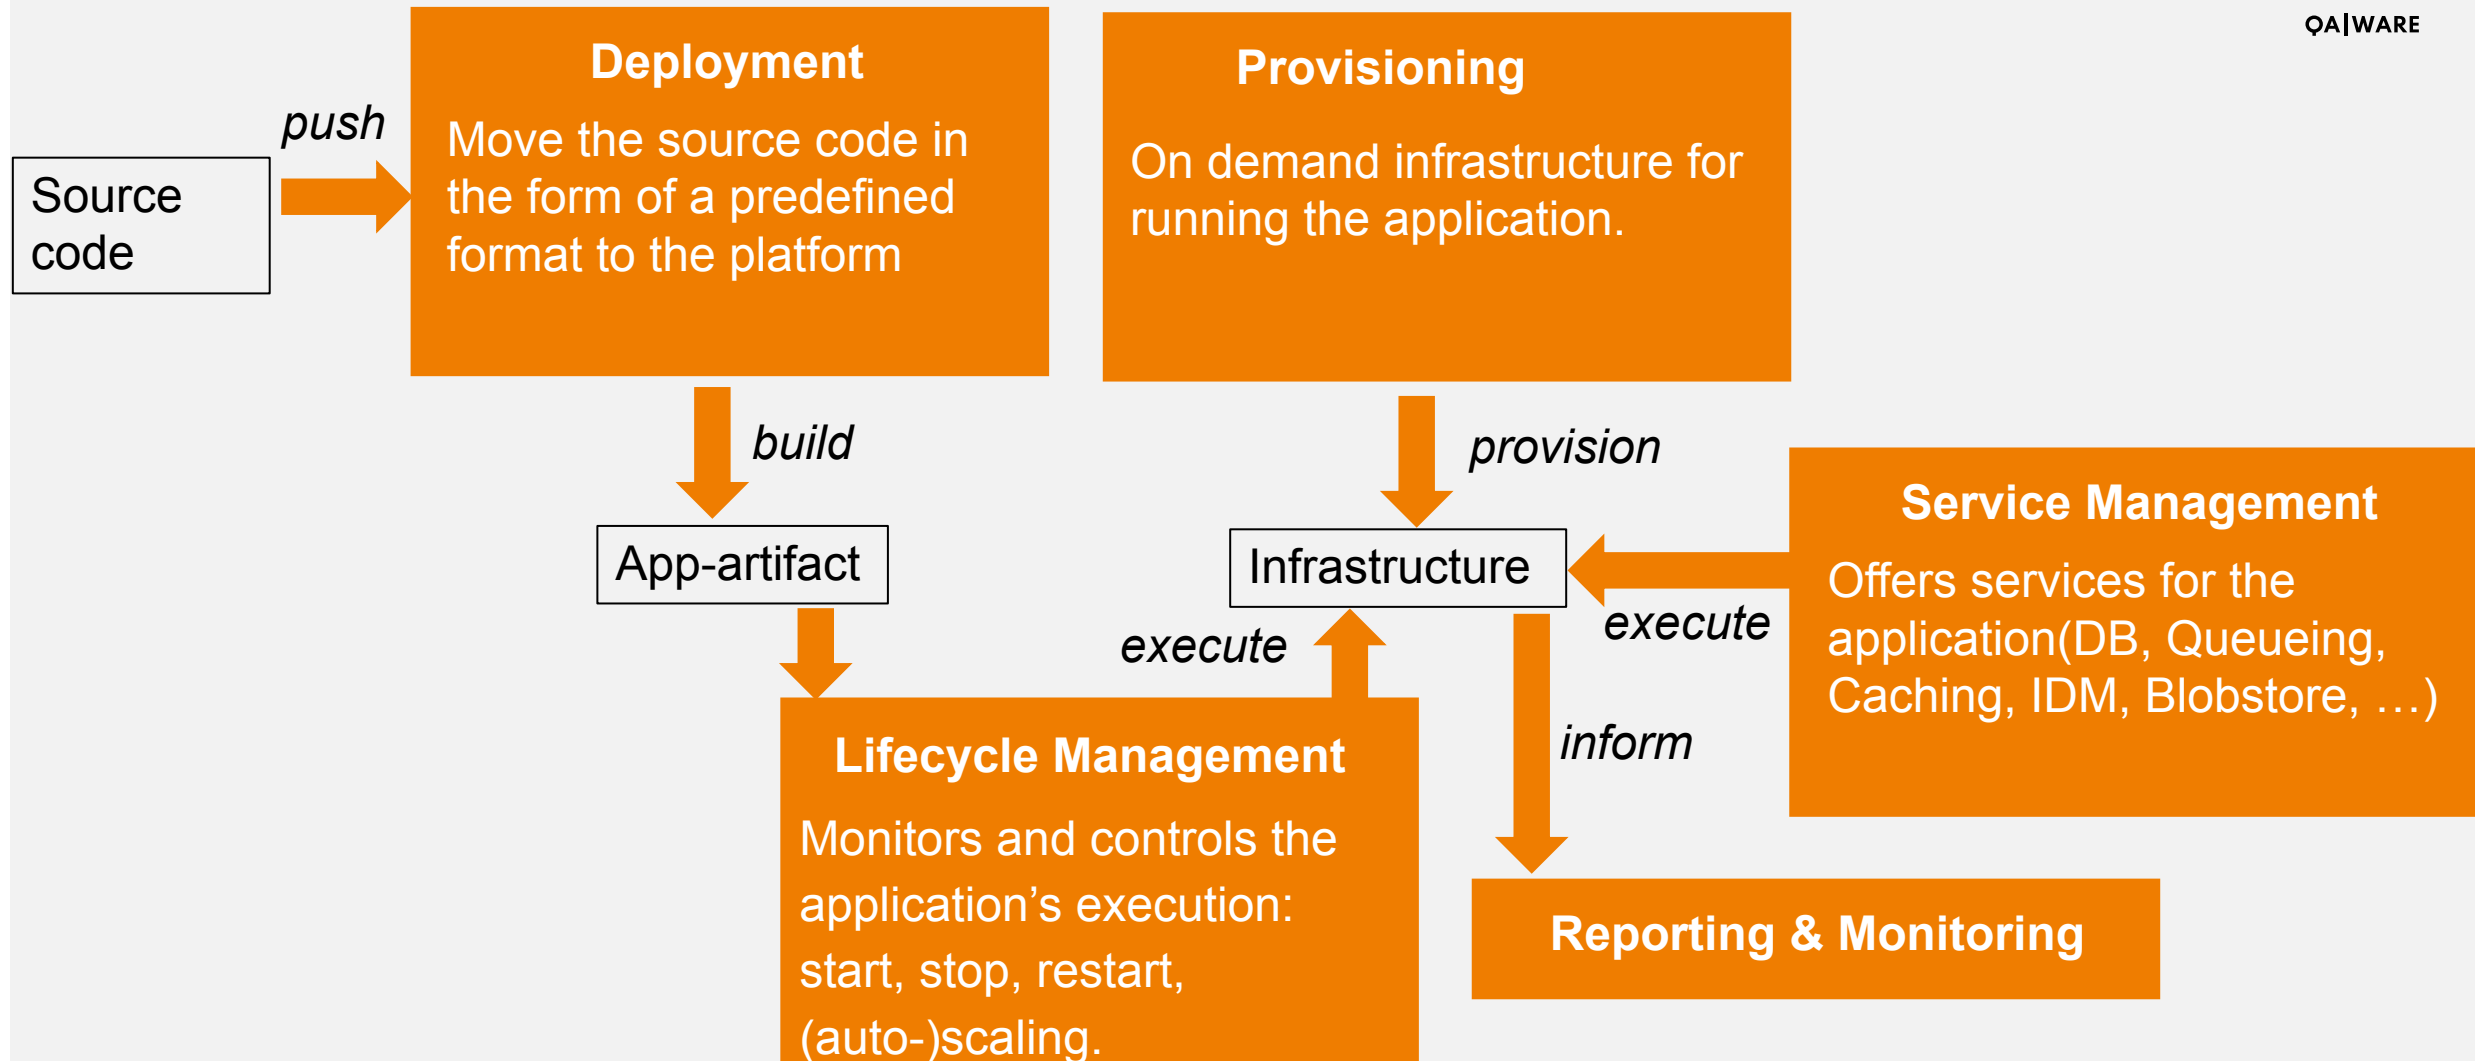
- Reuse of infrastructure/APIs
- Convenience services for developers



# IaaS vs. PaaS



# The functional building blocks of a PaaS Cloud.



← = Dataflow




# Example: Cloud Foundry



QA|WARE



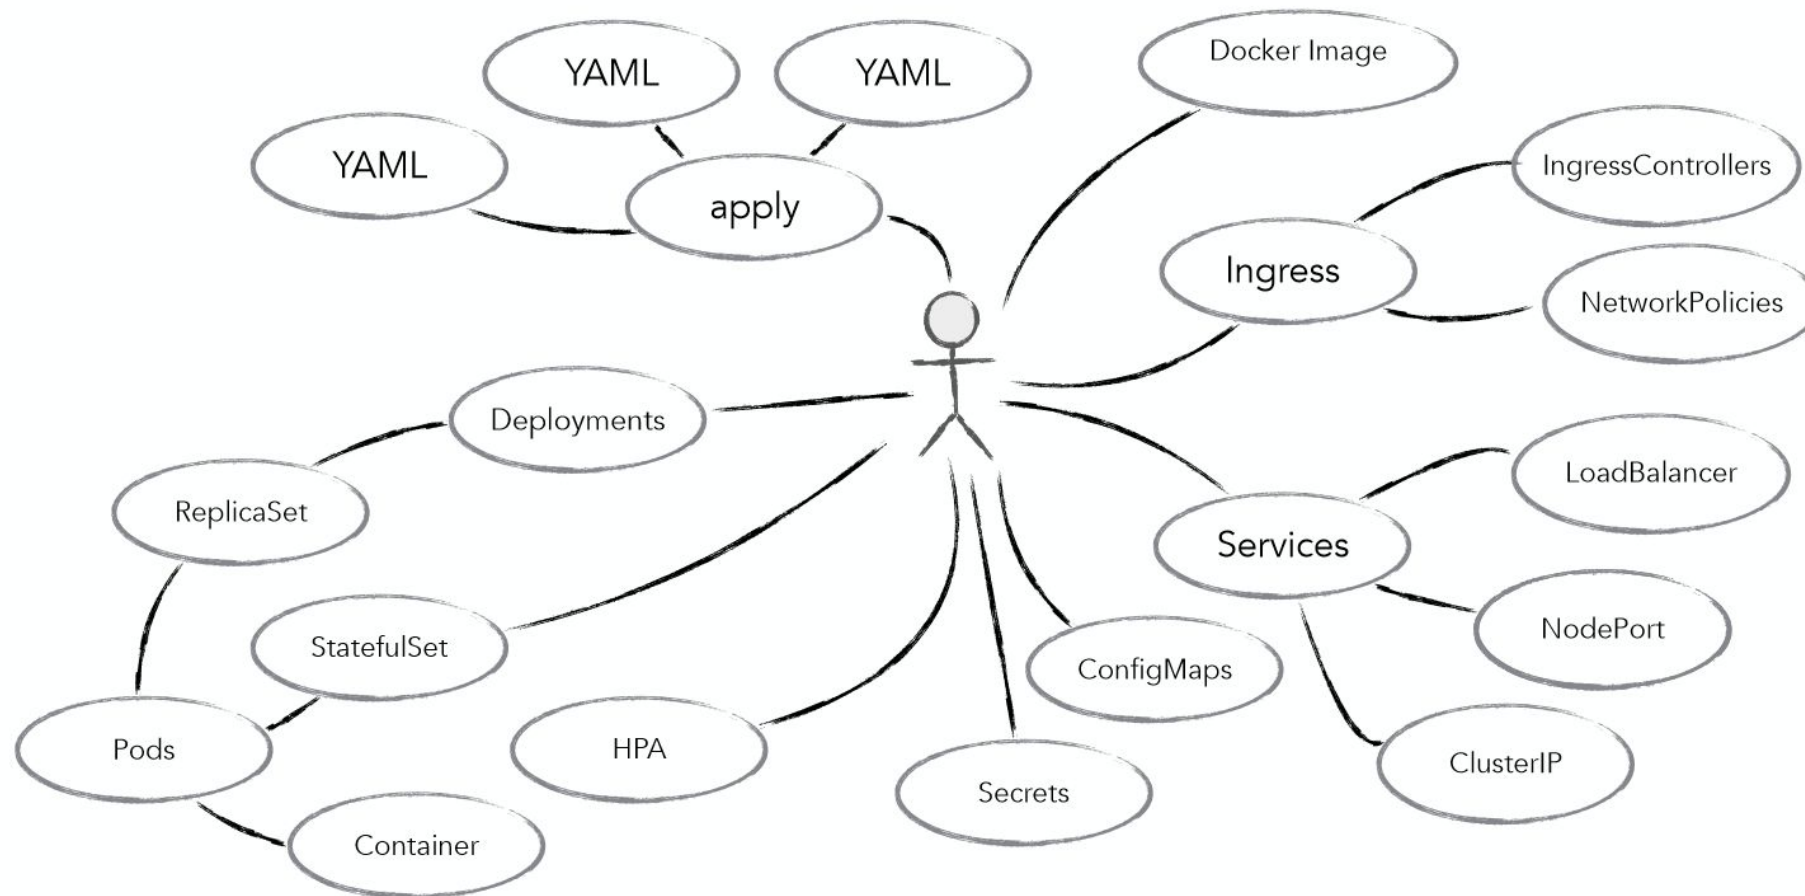
## Minimal Concepts



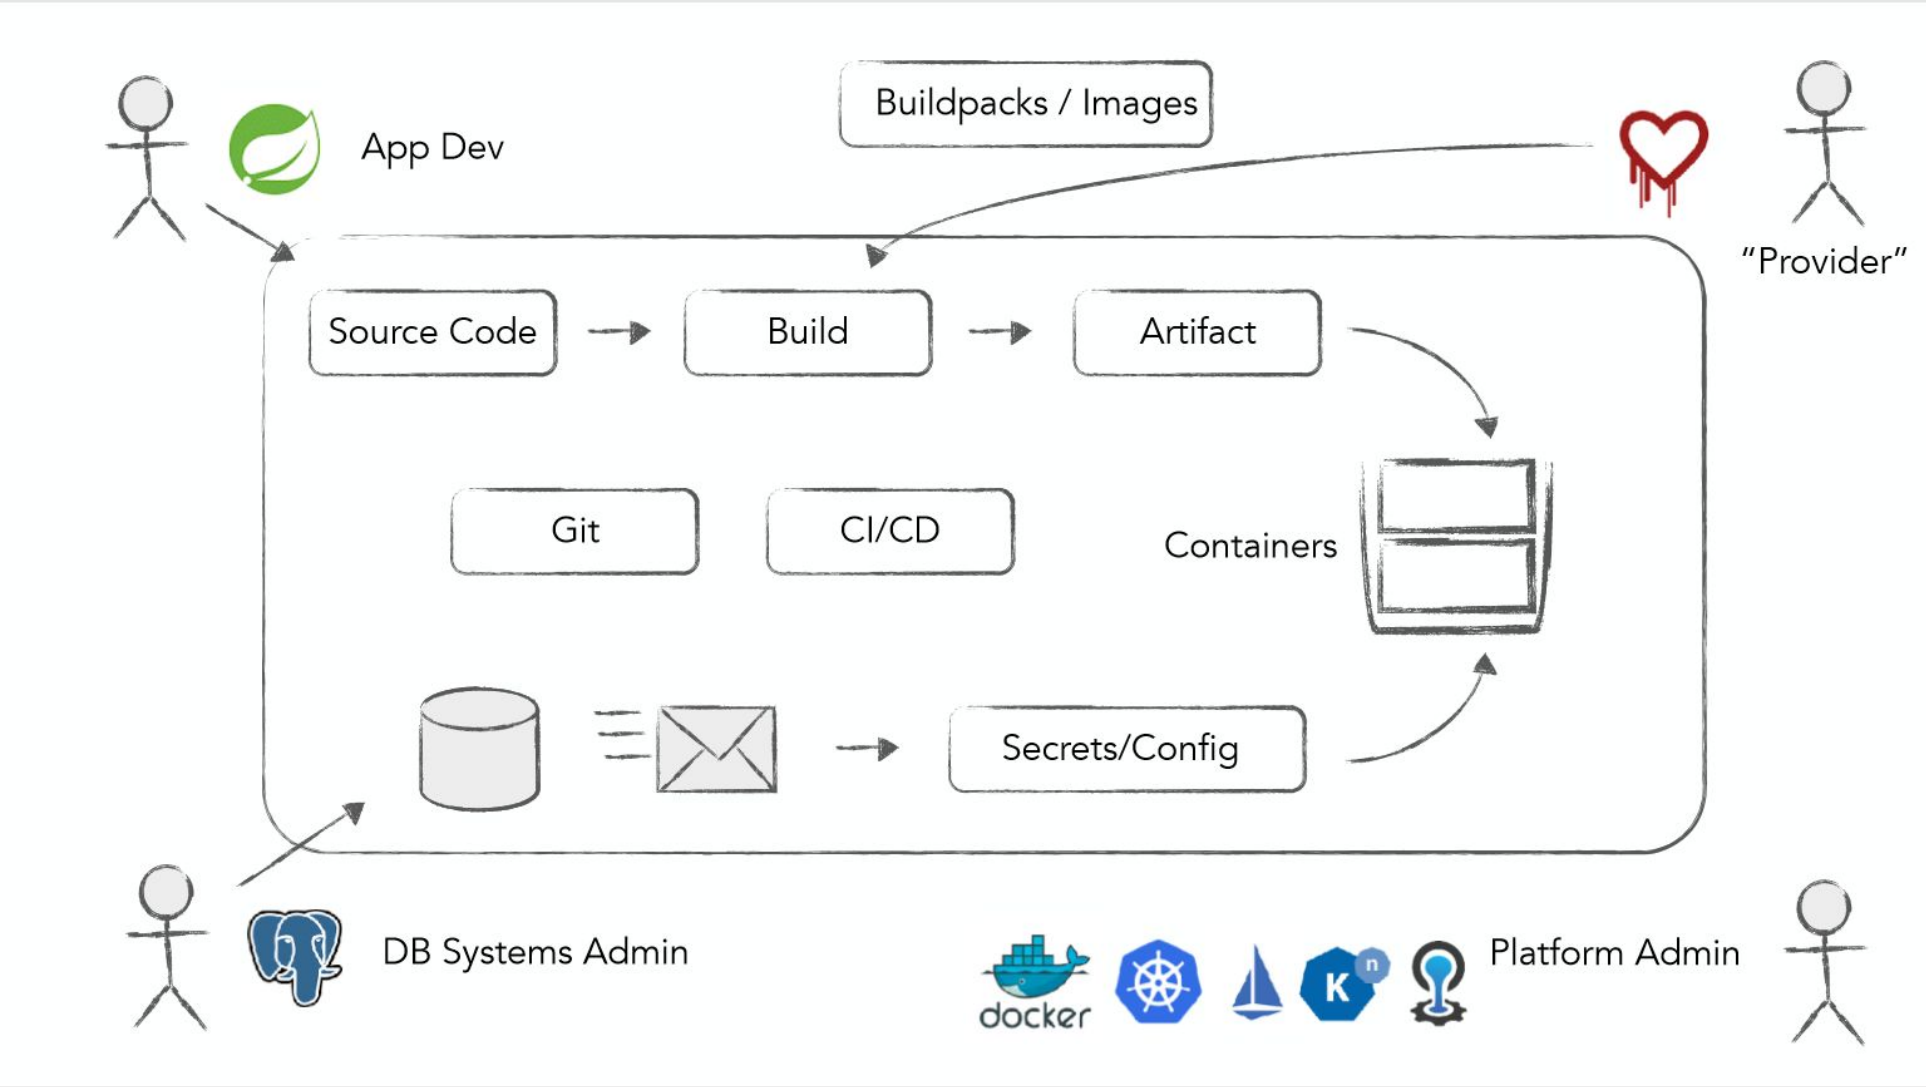
# In comparison: Kubernetes



## Minimal Concepts

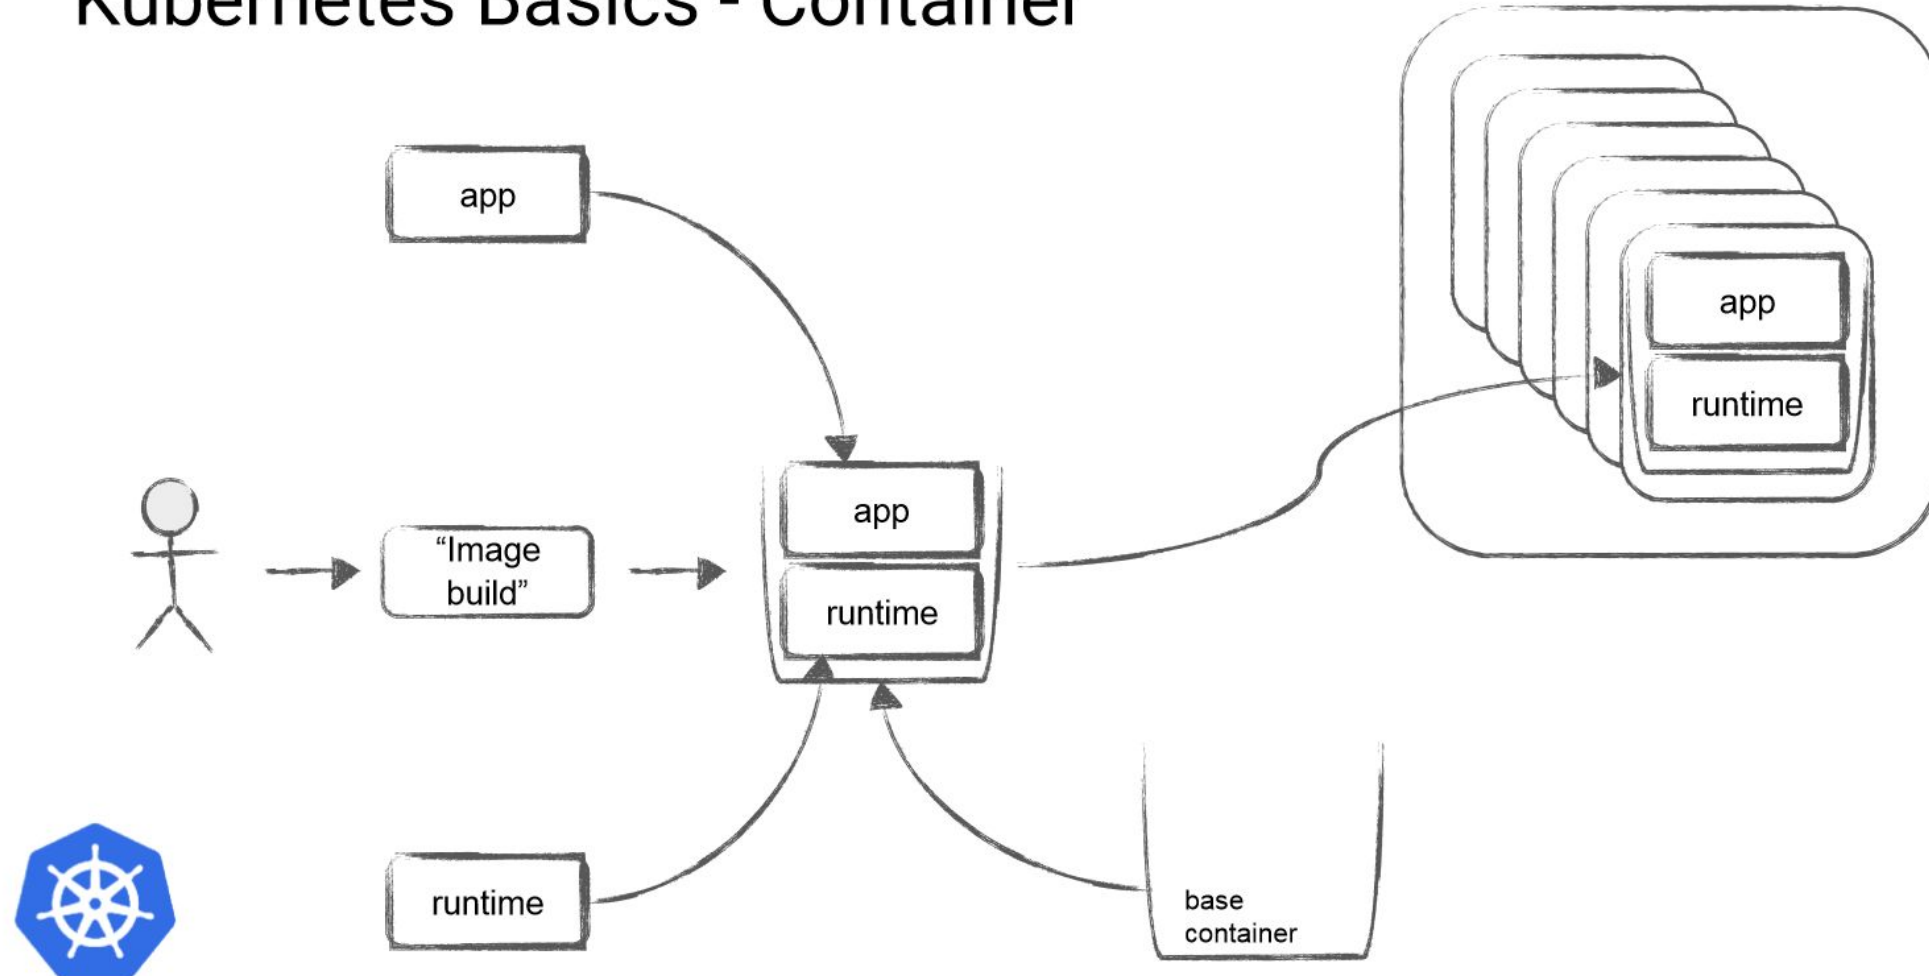


# Example: Cloud Foundry - Workflow



# In comparison: Kubernetes – Workflow (1/2)

## Kubernetes Basics - Container






## In comparison: Kubernetes – Workflow (2/2)

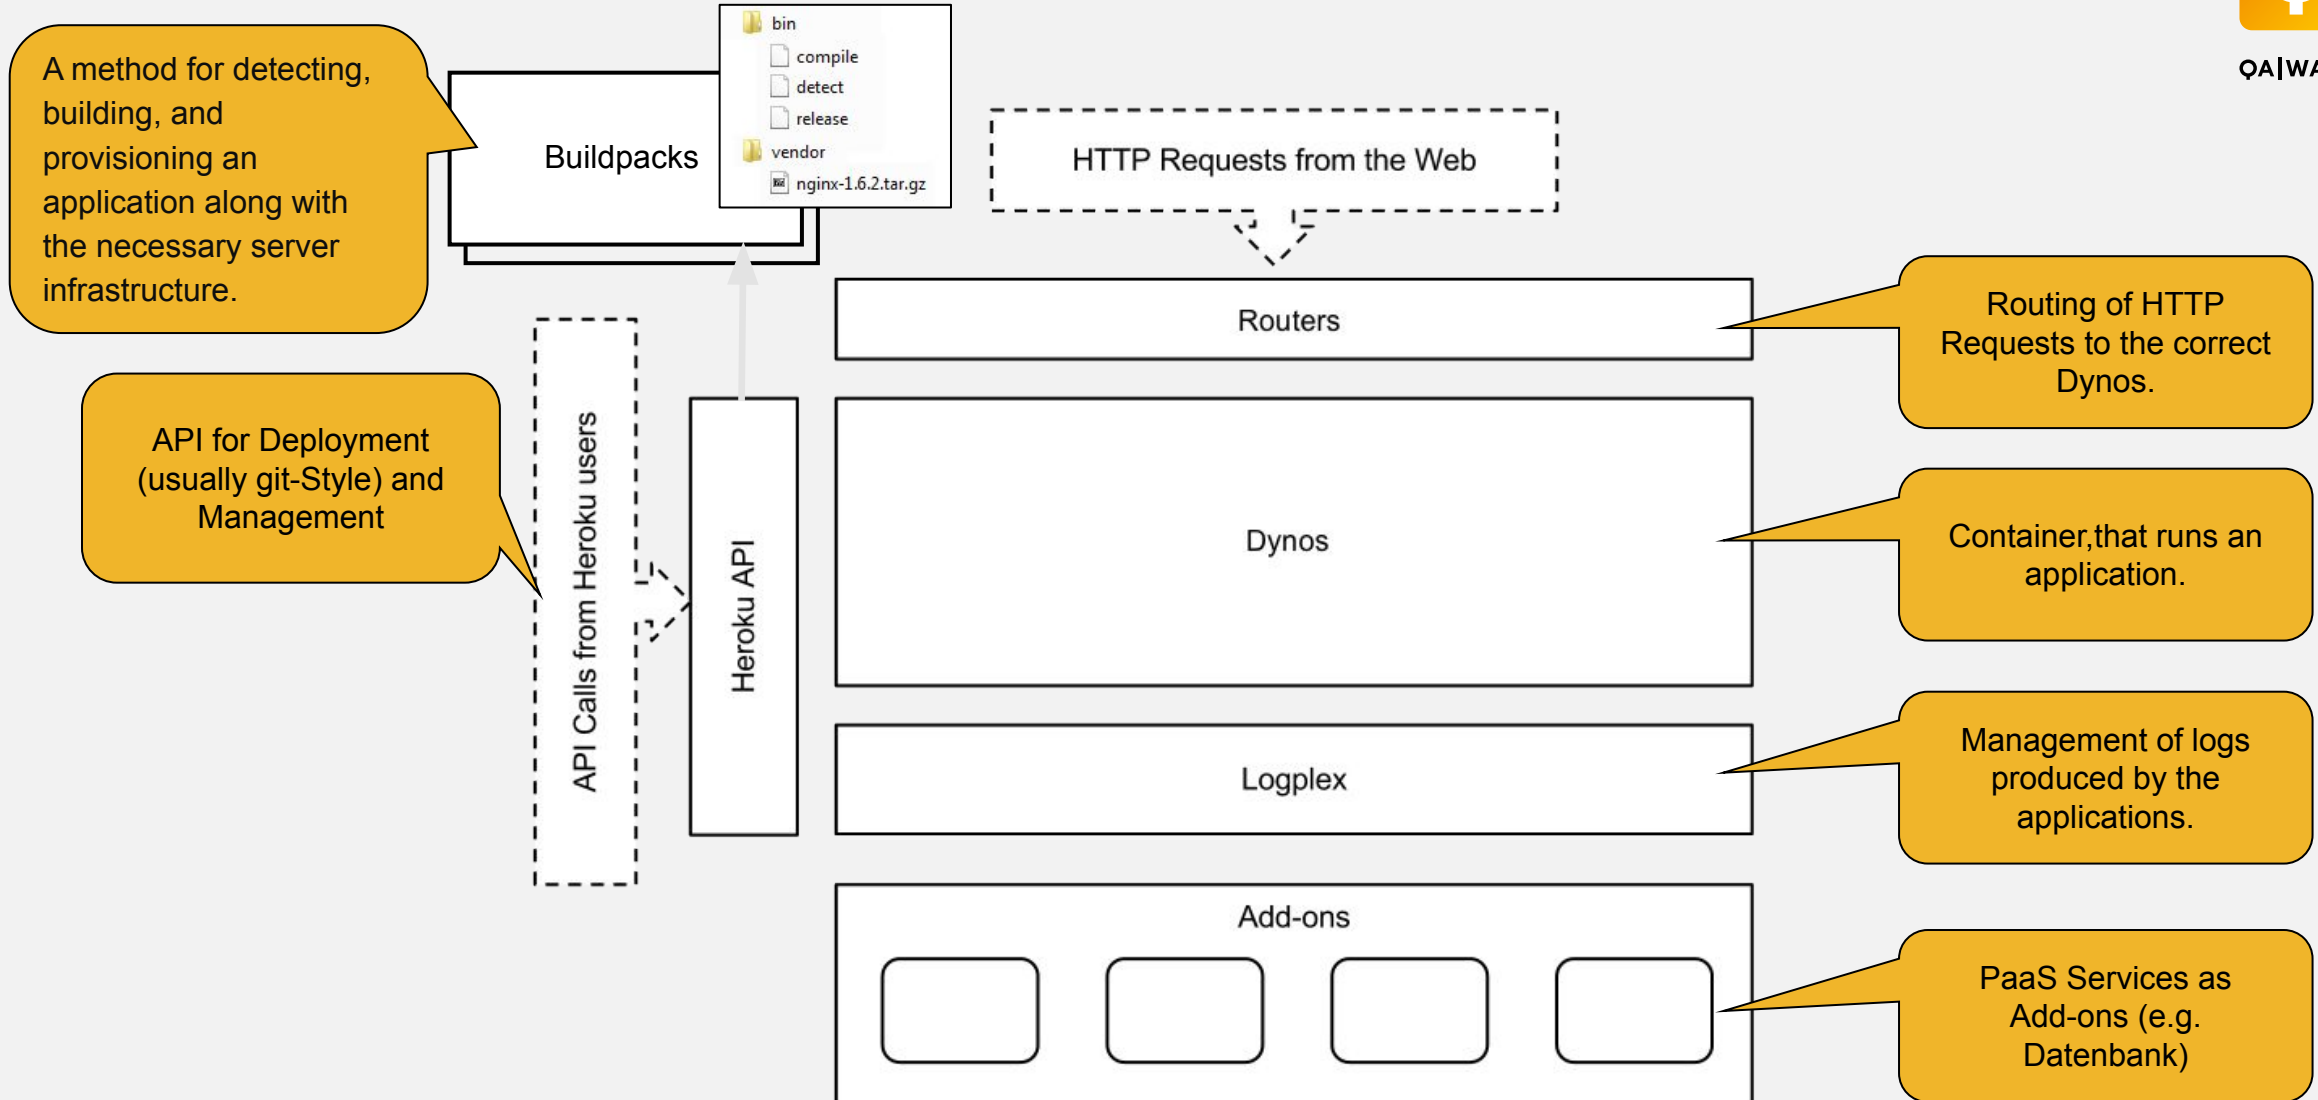


QA|WARE

### Kubernetes Basics - Orchestration



# High-Level Architecture of a PaaS, looking at Heroku.





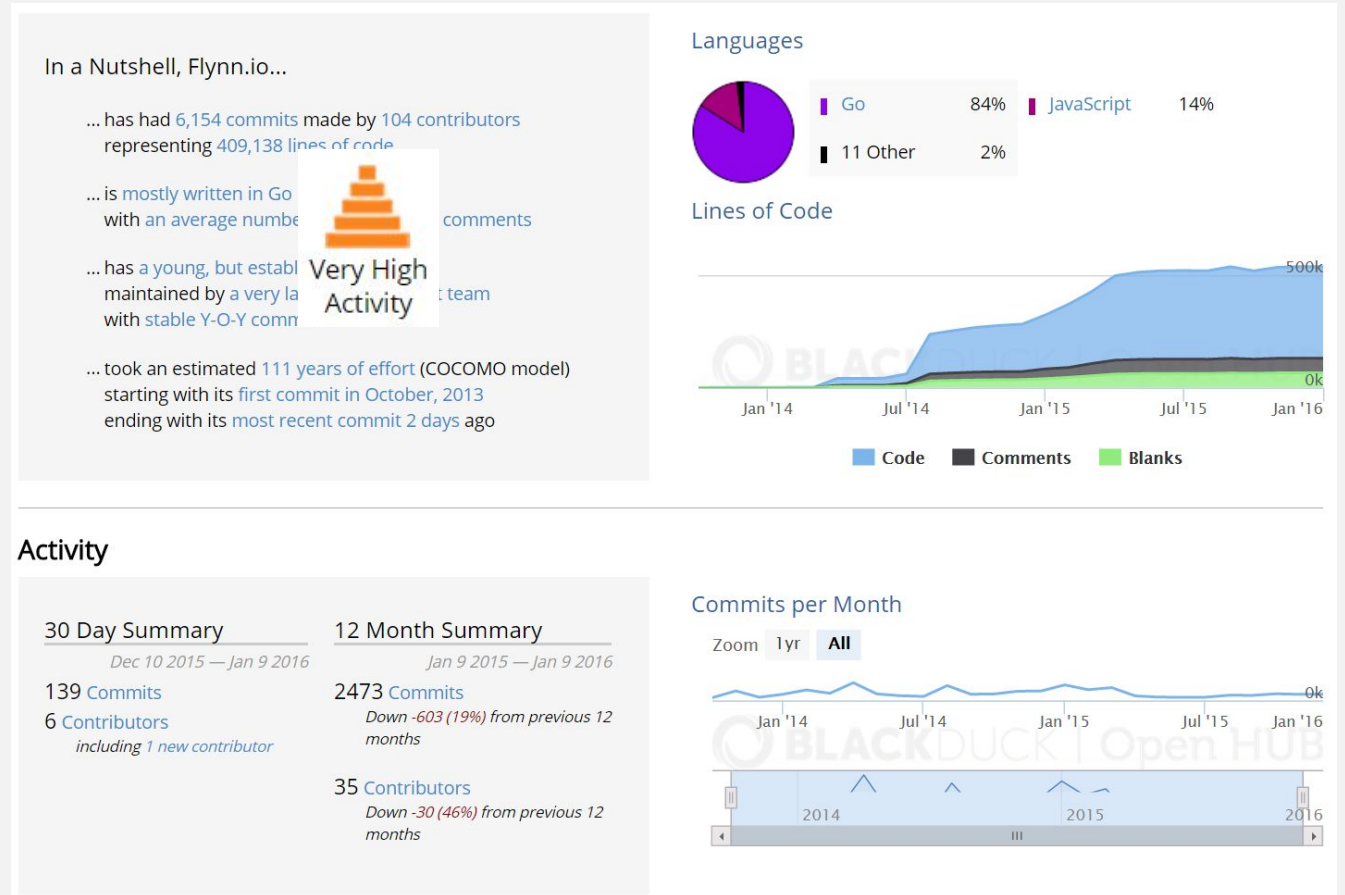
QA|WARE

# Private PaaS Clouds



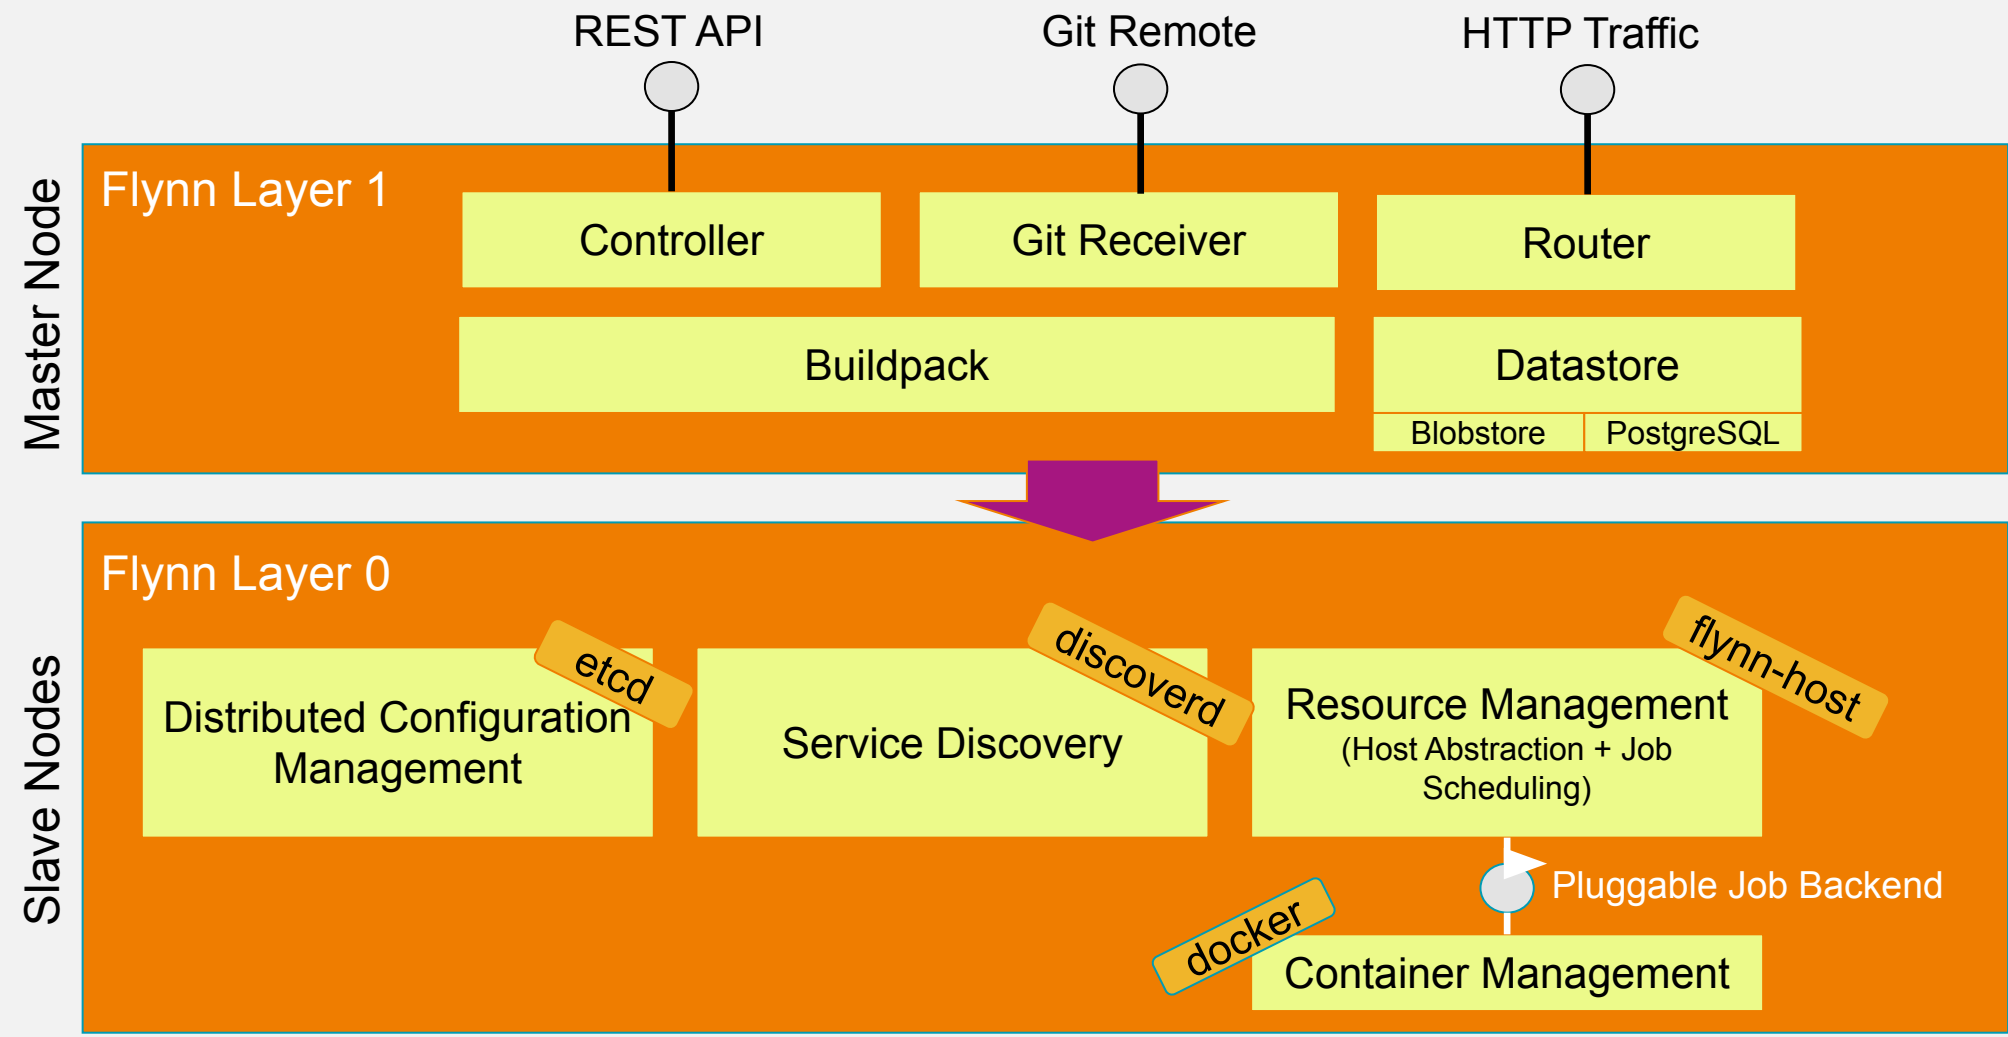
# Beispiel: Flynn

- Private PaaS auf Basis Docker
- Open-Source-Projekt unter einer BSD Lizenz






# Die Architektur von Flynn



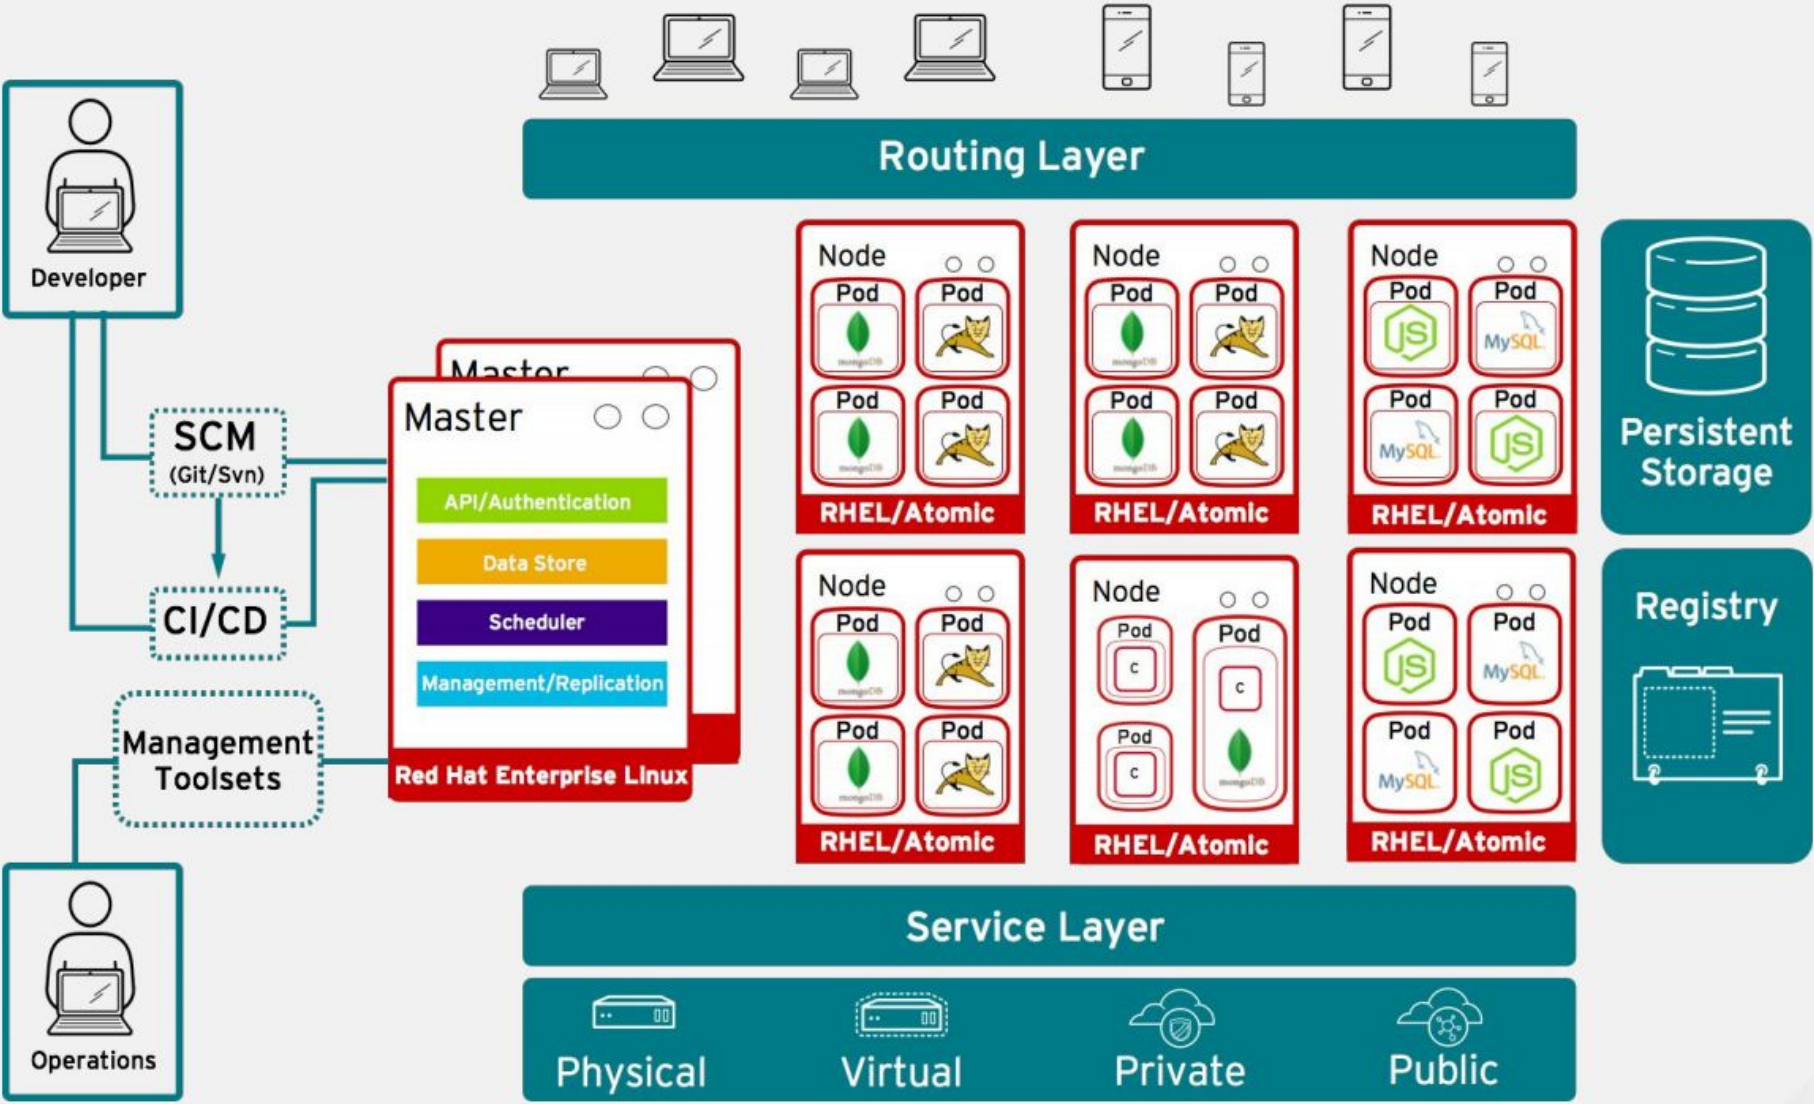
# Alternative Private PaaS Clouds




- OpenShift (<https://www.openshift.com>, PaaS created by Red Hat, based on Kubernetes)
- CloudFoundry (<http://www.cloudfoundry.org>, PaaS based on Kubernetes)
- PaaS-TA (<https://github.com/yelp/paasta>). Open-Source private PaaS based on Kubernetes (originally Mesos and Marathon)

# Example: OpenShift.

<https://www.redhat.com/cms/managed-files/OpenShift.pdf>

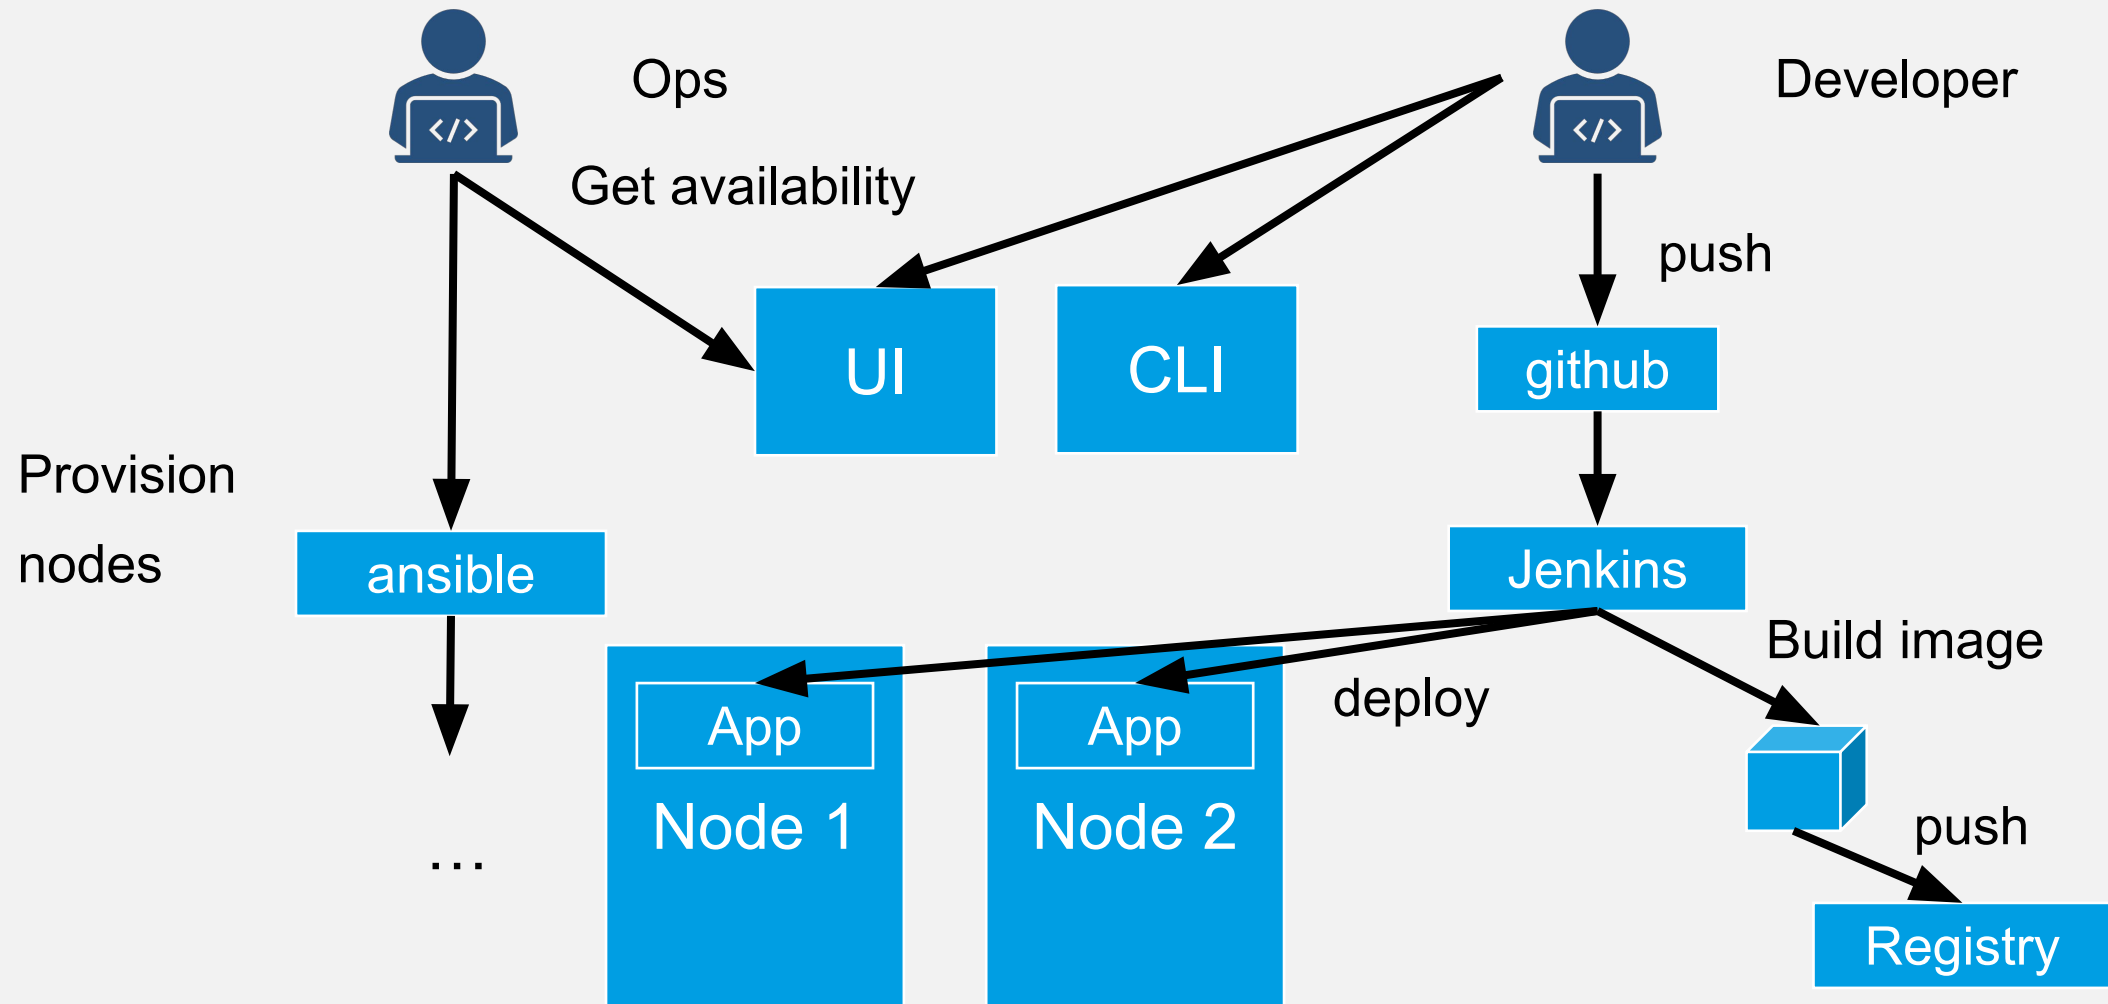


# OpenShift (Whiteboard)





# OpenShift (Whiteboard)





QA|WARE

# Public PaaS Clouds

# Ein PaaS-Vergleich über die angebotenen APIs und Services.

	GAE-J	AWS
<b>Datenspeicher</b>	App Engine Datastore (Key/Value mit JDO und JPA API) Cloud Storage (Objekte) Blobstore (Dateien), Cloud SQL (relational)	DynamoDB (Key/Value), S3 (Objekte und Dateien), RDS (relational)
<b>Messaging</b>	Mail (mit javax.mail API), XMPP, Channel (Push-API)	SES (E-Mails), SNS (Notifications), SQS (Message Queuing)
<b>Engine</b>	Servlet Engine, Capabilities, LogService	Elastic Beanstalk (Servlet Engine)
<b>Integration</b>	URLFetch, App Identity, OAuth	
<b>Parallele Verarbeitung</b>	Task Queue	Elastic MapReduce
<b>Volltextsuche</b>	Search, Prospective Search	CloudSearch
<b>Cache</b>	Memcache mit JCache-API	ElastiCache
<b>User-Authentifizierung</b>	Google Accounts, OpenID	IAM
<b>SaaS-APIs</b>	Google Data API, Images, Conversion	SWF (Workflows)
<b>Mandantenfähigkeit</b>	Multitenancy (Namespaces API)	

# The Google App Engine




QA|WARE

The Google App Engine (GAE) is Google's PaaS offering. Applications run within Google's infrastructure.

Running applications is free within certain quotas. Beyond that, costs are incurred based on factors such as service calls, storage volume, and actual CPU seconds used.

Supported languages:

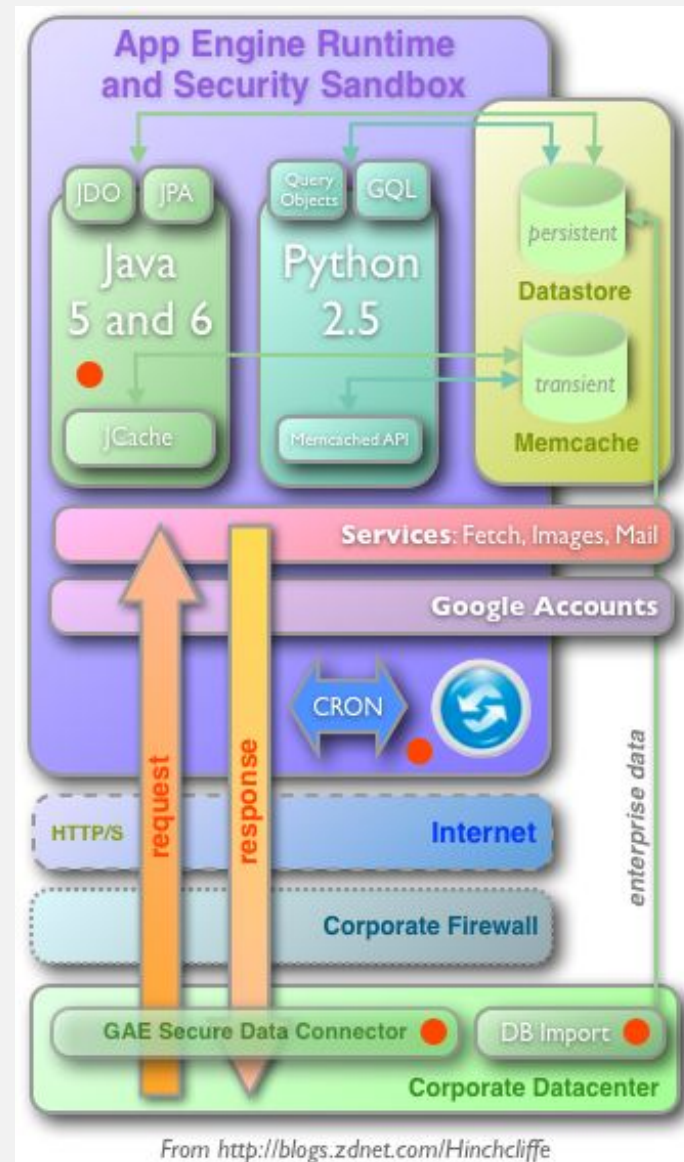




# The Google App Engine - Overview



QA|WARE





# Some GAE Services (1/2)



QA|WARE

## Datastore

- Persistent storage implemented as a key/value database.
- Transactions are atomic. Write operations are strongly consistent. Queries are eventually consistent.
- Definition, querying, and manipulation of data are done through a custom language called GQL (Google Query Language, similar to SQL).
- High-level APIs such as JDO and JPA are available. These are standardized within the Java/JEE ecosystem. The API is implemented using the DataNucleus framework.

## Memcache

- High-performance temporary data storage in memory (in-memory data grid).
- Each entry is stored with a unique key.
- Each entry is limited to 1 MB.
- An expiration time in seconds is specified, after which the entry will be removed from the Memcache.
- Data may also be evicted earlier, depending on Memcache usage.
- The JCache API is available as a high-level API.

# Some GAE Services (2/2)

## URL Fetch

- Access to content on the internet.
- Supported methods: GET, POST, PUT, DELETE, and HEADER.
- Access is allowed on ports within the ranges 80-90, 440-450, and 1024-65535.
- Requests and responses are each limited to 1 MB.

## Users

- Integration with a single sign-on system.
- Supports Google Accounts and OpenID accounts.
- The JAAS API is used as the high-level API.

## XMPP

- Messages can be sent to and received from any XMPP-compatible messaging system.
- Each application has a unique XMPP username.

## All APIs:

- [App Identity](#)
- [Blobstore](#)
- [Google Cloud Storage](#)
- [Capabilities](#)
- [Channel](#)
- [Conversion](#)
- [Images](#)
- [Mail](#)
- [Memcache](#)
- [Multitenancy](#)
- [OAuth](#)
- [Prospective Search](#)
- [Search](#)
- [Task Queues](#)
- [URL Fetch](#)
- [Users](#)
- [XMPP](#)



QA|WARE

# Constraints when running inside of Google App Engine (1 / 2)



QA|WARE

A GAE application runs in a sandbox that restricts the application's behavior. This is done to ensure efficient processing and to protect the infrastructure during auto-scaling.

Not all classes of the standard library may be used.

- No custom threads may be created.
- No access to the runtime environment, such as its class loaders.


# Constraints when running inside of Google App Engine (2 / 2)



QA|WARE

- Communication with other web applications or servers is only possible via URL Fetch, XMPP, or email.
- Requests and responses must not exceed 1 MB in size.
- Webhooks are used as a general architectural mechanism for incoming communication, triggered by events (e.g., warmups), messages, or cron jobs.
- All requests to a GAE application are terminated after 60 seconds.
- Various restrictions apply to data volumes and the number of service calls.

# Problems with most PaaS solutions



- Very opinionated, that's why they work so great in greenfield projects
- Hard to integrate into the brownfield environment of existing enterprises

Enterprises, however, want to profit from the PaaS idea as well.


This led to a engineering specialization called *platform engineering*.



# Cloud native resulted in massive cognitive load 🚀




QA|WARE



Source: Amazon Web Services



# Platform Engineering solves that issue



*“Platform engineering is the discipline of designing and building **toolchains and workflows** that enable **self-service capabilities for software engineering** organizations in the cloud-native era. Platform engineers provide an **integrated** product most often referred to as an **“Internal Developer Platform”** covering the operational necessities of the **entire lifecycle** of an application.”*

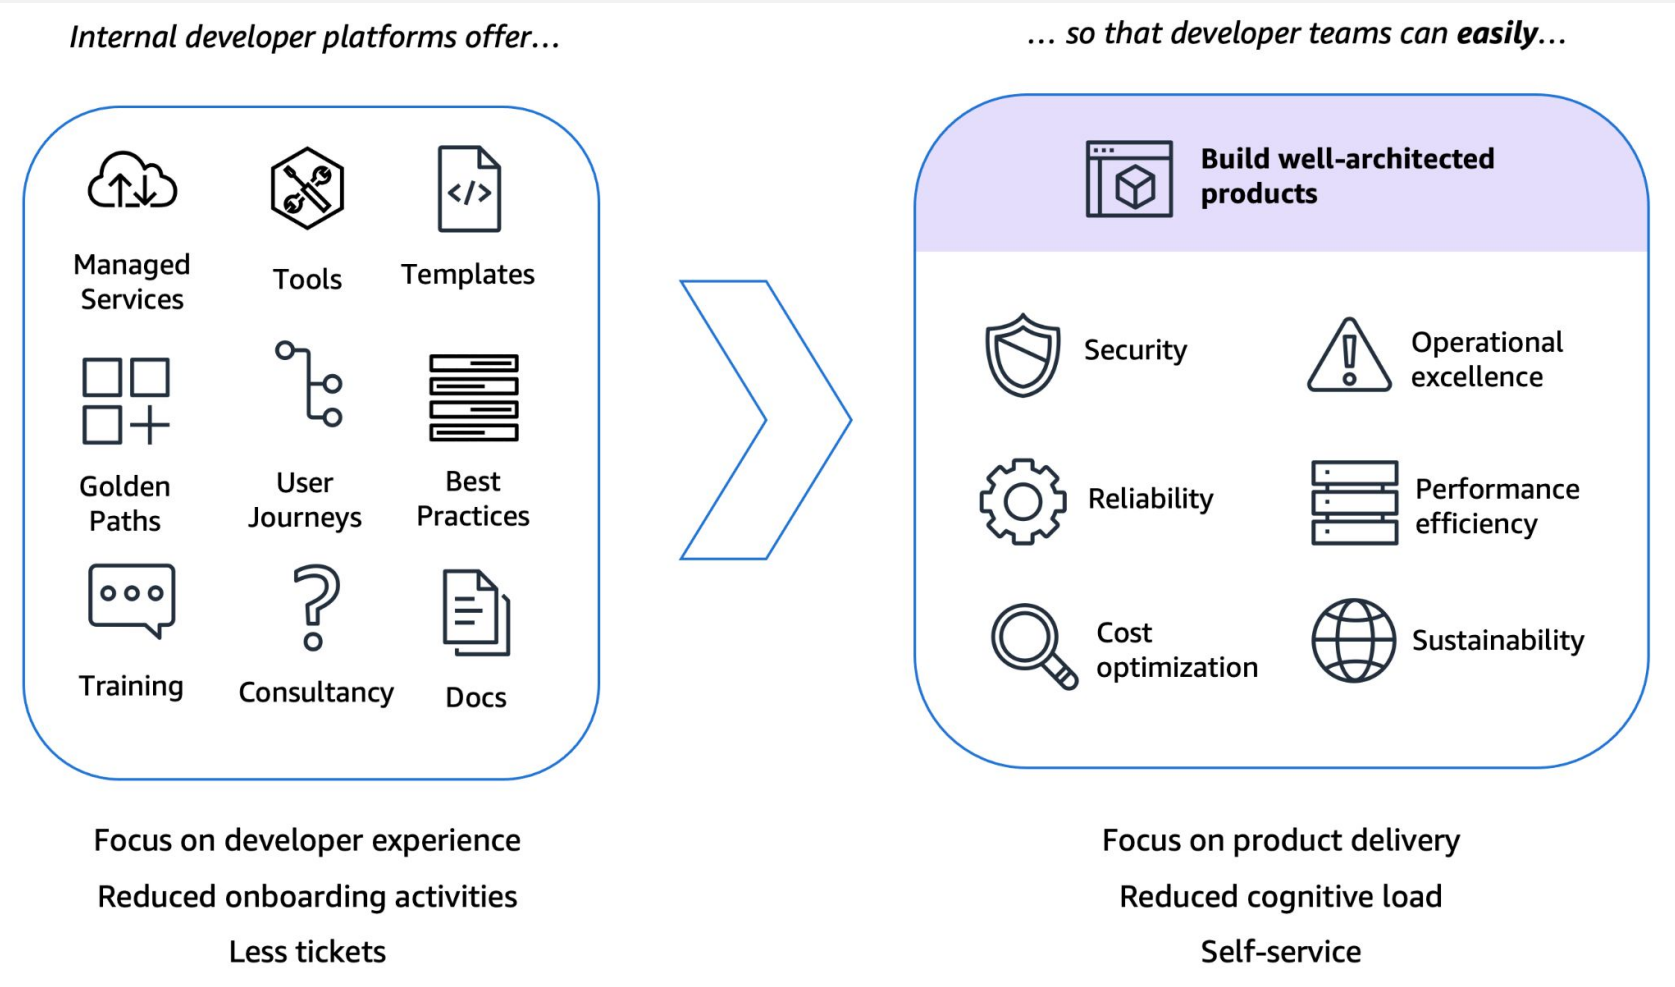
Humanitec

- Specialization of roles, resulting in reduced cognitive load.
- Still DevOps, with a central interface: the platform.
- Fosters reuse and organizational scaling.
- Automated integration means more time for actual software engineering, i.e. creating business value.

# Internal Developer Platforms



QA|WARE



Source: Amazon Web Services