

qaware.de



QA|WARE
SOFTWARE ENGINEERING

Observability

Franz Wimmer

franz.wimmer@qaware.de

What exactly is observability?




How well do I know the internal state of my system when I look at its outputs? (see [Kal])

What does such a modern system output?

- A website?
- A REST API?
- E-mails to the admin?
- A sparsely maintained log file?

What if the system consists of many microservices?



What exactly is observability?



QA|WARE

“Observability” today typically refers to three topics:

- Logs
- Metrics
- Traces

All together, this enables insight into distributed microservices.



The goal of observability: Diagnosability. What do I understand by that?



QA|WARE

*“A **structured approach** to equipping an application with **sensors** in advance so that I can **quickly identify the error** in the event of a **fault** and have the information necessary to repair it.”*

A system can be easily diagnosed if you can easily recognize and correct healthy and unhealthy conditions.



A diagnosable system has a



- short Mean Time to Detect (MTDD)
- short Mean Time to Repair (MTTR)

and thus to a short period of time in which errors remain undetected and even exist.

Diagnosability: A structured approach is necessary to avoid influencing the system too much.



QA|WARE

1. Get an overview:

- a. What are the central components, e.g. login, etc.?
- b. What are the supporting components, e.g. batch and loader jobs, etc.?

2. Identify error limits:

- a. Internal error limits: layers/use cases
- b. External error limits: incoming and outgoing calls

3. Define error classes:

- a. Severity levels: operation still possible, no impact on customer, etc.
- b. Impact: certain functions are not available to customers, etc.

4. Determine the runtime data that is necessary for detection:

- a. Uniformity: data has the same meaning; uniform data formats are defined, etc.
- b. Clearly defined: it is clear what the metric means, e.g. CPU load

5. Define actions:

- a. Alerting: who is notified and when
- b. Create playbooks for errors: How do I get the data, etc.




QA|WARE

Why (Cloud) Observability?


Back then The monolith



QAWARE



Today Microservices



The (Open Source) Stack



QAWARE

Today we use the LGTM stack.:

- Loki (Logs)
- Grafana (Visualization)
- Tempo (Traces)
- Mimir (Metrics)

<https://grafana.com/oss/>

The (Open Source) Stack

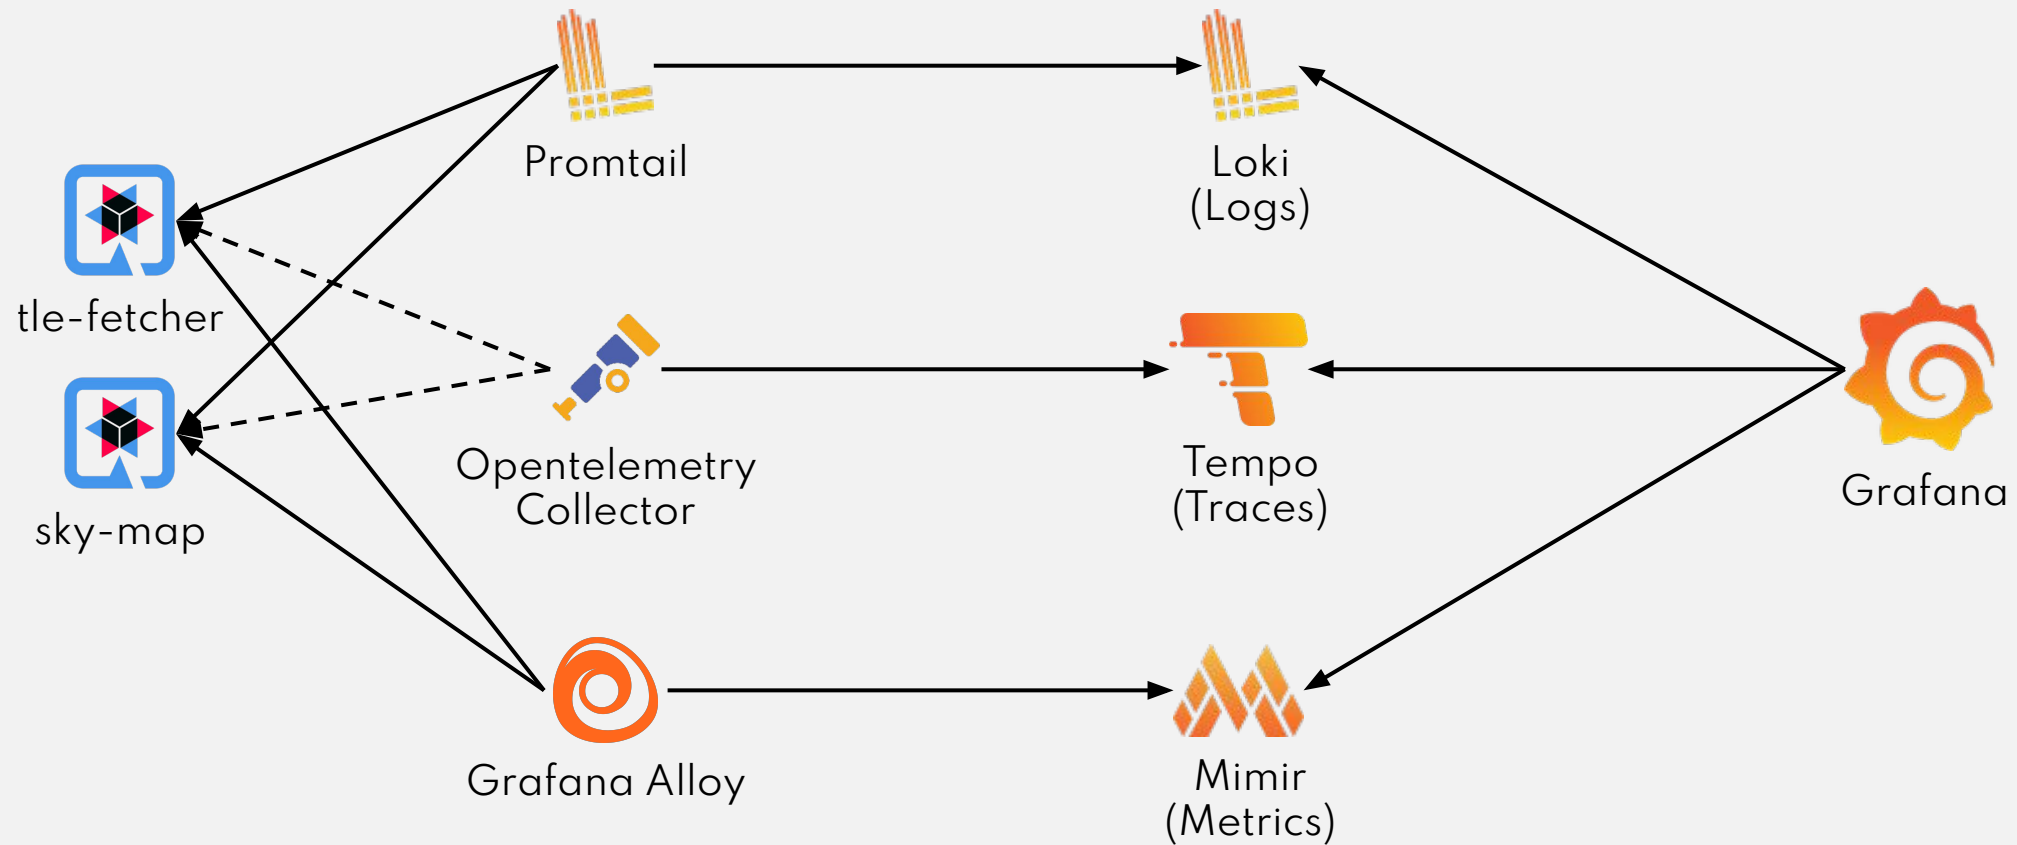


Applications

Data collection

Storage

Visualization





QA|WARE

Logs



Grafana loki

Logs



QA|WARE

- **Logs** are an important part of any IT system.
- In the cloud, we need to look at logs from many different services at the same time.
- Every system has logs, but they are always in a different **format**.
- Example: **Quarkus** webservice:
 - `2023-01-10 20:56:42,122 DEBUG [io.qua.mic.run.bin.ver.VertxHttpServerMetrics]
(vert.x-eventloop-thread-11) requestRouted null HttpRequestMetric [initialPath=/q/metrics,
currentRoutePath=null, templatePath=null, request=io.vertx.core.http.impl.Http1xServerRequest@512b1fd6]`
- Example: **Nginx** webserver:
 - `- 2a02:c207:3005:5132::1 - - [10/Jan/2023:00:00:13 +0100] 0.000 repo.saturn.codefoundry.de "POST
/api/v4/jobs/request HTTP/1.1" 957 204 0 "-" "gitlab-runner 15.6.1 (15-6-stable; go1.18.8; linux/amd64)"`

Logging. Please be structured. Please be well-considered. /1



QA|WARE

- Define a log format and ensure that all services use the same format.
- Define a diagnostic context = information that helps in the event of an error, e.g.
 - Traceld
 - UserId
 - SessionId
- Use structured logging (e.g. JSON)
 - Simplifies handling in the analysis
- Use asynchronous logging (provided that the recipients support it) to prevent blockages
 - If all of this is sent over TCP etc. and then filtered and sorted later, the order doesn't matter 😎

Logging. Please be structured. Please be well-considered. /2



QAWARE

- Don't log too much, but log every
 - exception,
 - every error,
 - every meaningful piece of information ...
 - ... but not multiple times.
- Use log levels. To do this, define which category should have which level, e.g.:
 - WARN: Errors that affect a single request but not the stability of the service.
 - ERROR: Errors that affect the stability of the service.

Loki



QA|WARE

Loki is the log aggregator in the Grafana ecosystem.

Passive log storage – logs have to be obtained differently.

Promtail is installed locally if possible and pushes the logs to **Loki**.


Since: 2018

License: AGPL 3.0

Architecture



QAWARE



Promtail configuration



QAWARE

```
clients:
  - url: "https://loki.qaware.de/loki/api/v1/push"
    basic_auth:
      username: loki
      password: *****

scrape_configs:
  - job_name: journal
    journal:
      json: true
      max_age: 12h
      labels:
        host: saturn
        job: systemd-journal
    relabel_configs:
      - source_labels: ['__journal__systemd_unit']
        target_label: 'unit'
```

Caution: don't use labels too dynamically!

<https://grafana.com/blog/2020/04/21/how-labels-in-loki-can-make-log-queries-faster-and-easier/>

LogQL



QAWARE

Loki has its own query language.

```
count_over_time({foo="bar"}[1m]) > 10
```

```
{host="$host", job="systemd-journal"}  
  | json  
  | line_format "{{ .unit }}: {{ .MESSAGE }}"  
  |= "$search"
```



QA|WARE

Metrics



Metrics: Basics



QA|WARE

Metrics are measurements that reflect the current state of the system.

Examples:

- **CPU** load
- Size of the **JVM Heap**
- Number of **calls to an interface**

The metrics are provided by the system to be monitored.

Metrics can also have metadata.

Our metrics storage: Grafana Mimir



QAWARE

- Mimir (2022) stores **metrics**
- It contains a lot of code from the older tool Prometheus (2012)
- Technical foundation:
 - Time series database for numerical time series
 - Collectors for metrics and text-based format for metrics
 - Query language (PromQL) for time series
 - Alerting based on time series

Metrics: Basics



QAWARE

Grafana Alloy pulls metrics from all configured targets every **15 seconds**.

The transport is usually done via HTTP:

GET /metrics

No aggregation is done here yet. The consumer is responsible for that.

```
# HELP process_open_fds Number of open file descriptors.
# TYPE process_open_fds gauge
process_open_fds 8
# HELP process_start_time_seconds Start time of the process since unix epoch in seconds.
# TYPE process_start_time_seconds gauge
process_start_time_seconds 1.65566242485e+09
# HELP process_virtual_memory_bytes Virtual memory size in bytes.
# TYPE process_virtual_memory_bytes gauge
process_virtual_memory_bytes 1.782685696e+09

# HELP http_server_requests_seconds
# TYPE http_server_requests_seconds summary
http_server_requests_seconds_count{method="GET",status="200",uri="/tle",} 1.0
http_server_requests_seconds_sum{method="GET",status="200",uri="/tle",} 8.183356641
# HELP http_server_requests_seconds_max
# TYPE http_server_requests_seconds_max gauge
http_server_requests_seconds_max{method="GET",status="200",uri="/tle",} 8.183356641
```

Metric Types



QA|WARE

Metrics have different types to visualize data:

- **Counter**
A number that can either be incremented or reset.
- **Gauge**
A metric that can change in either direction.
- **Histogram**
Values in “buckets” - e.g. the duration of requests.
- **Summary**
Similar to histogram, summarizes values.

PromQL



QAWARE

PromQL is a powerful query language for filtering/aggregating the desired metrics.

The entire time series:

```
http_requests_total
```

Filter by label:

```
http_requests_total{job="apiserver", handler="/api/comments"}
```


Calculate rate:

```
rate(http_requests_total[5m])[30m:1m]
```

Mimir: Architecture (Write)



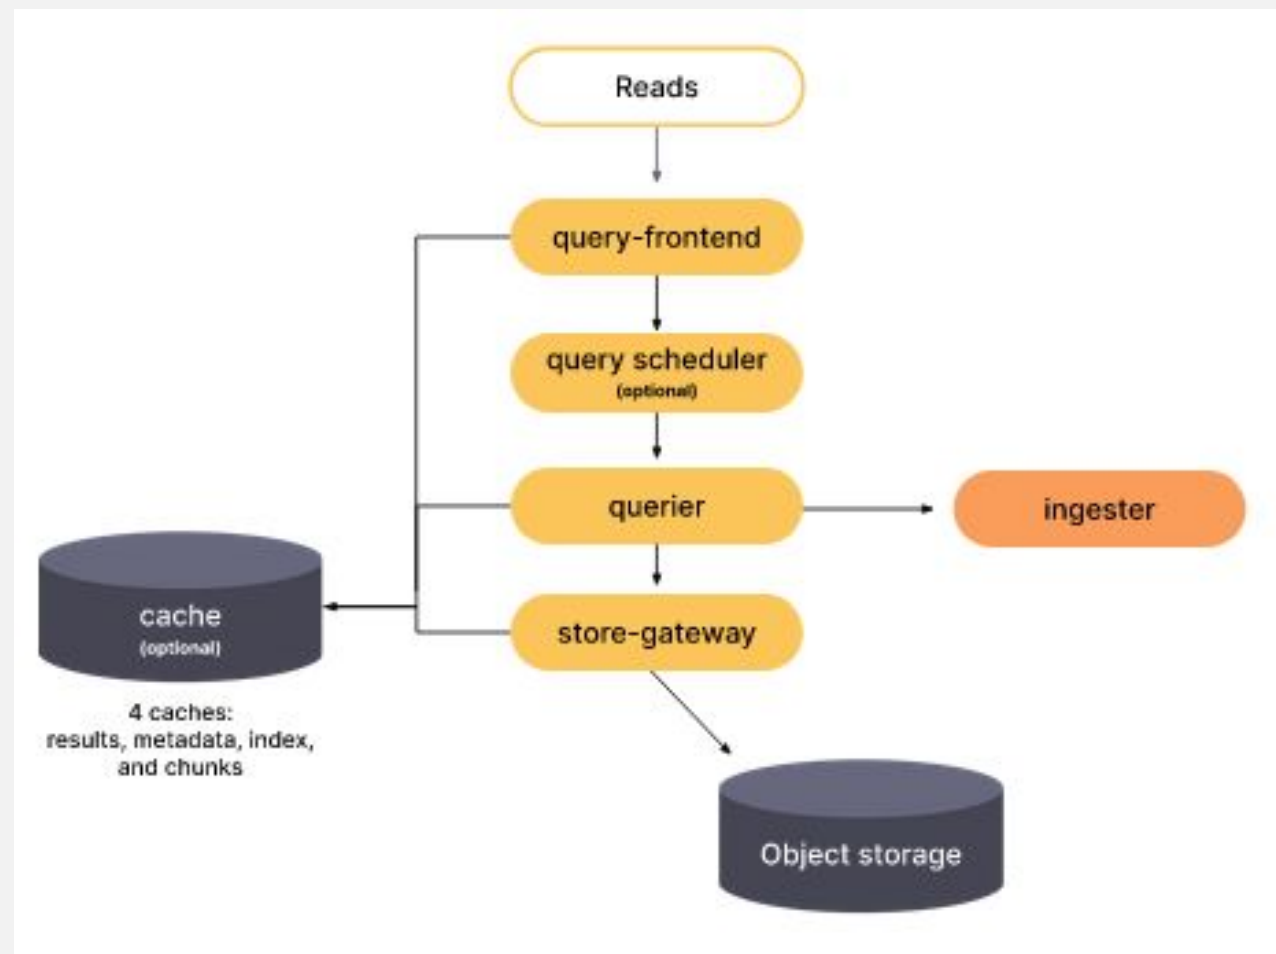
QAWARE



Mimir: Architecture (Read)



QAWARE



Configuration (application)

Example: Quarkus & Micrometer



QAWARE

Gradle:

```
implementation("io.quarkus:quarkus-micrometer-registry-prometheus")
```

Maven:

```
<dependency>  
  <groupId>io.quarkus</groupId>  
  <artifactId>quarkus-micrometer-registry-prometheus</artifactId>  
</dependency>
```

Grafana Alloy



QAWARE

- Scrapes metrics...
 - ...and forwards them to Mimir.
-
- Easy to use (binary + config file)
 - Good integration into Kubernetes, e.g. using the service discovery
 - Configuration in flow mode

Configuration in flow mode (1)



QAWARE

```
// Discover Kubernetes pods to collect metrics from.
discovery.kubernetes "pods" {
    role = "pod"
}

// Scrape metrics from Kubernetes pods and send to a prometheus.remote_write
// component.
prometheus.scrape "default" {
    targets = discovery.kubernetes.pods.targets
    forward_to = [prometheus.remote_write.default.receiver]
}
```


Configuration in flow mode (2)



QAWARE

```
// Get an API key from disk.
local.file "apikey" {
    filename = "/var/data/my-api-key.txt"
    is_secret = true
}

// Collect and send metrics to a Prometheus remote_write endpoint.
prometheus.remote_write "default" {
    endpoint {
        url = "http://localhost:9009/api/prom/push"

        basic_auth {
            username = "MY_USERNAME"
            password = local.file.apikey.content
        }
    }
}
```



QA|WARE

Traces



Distributed Tracing



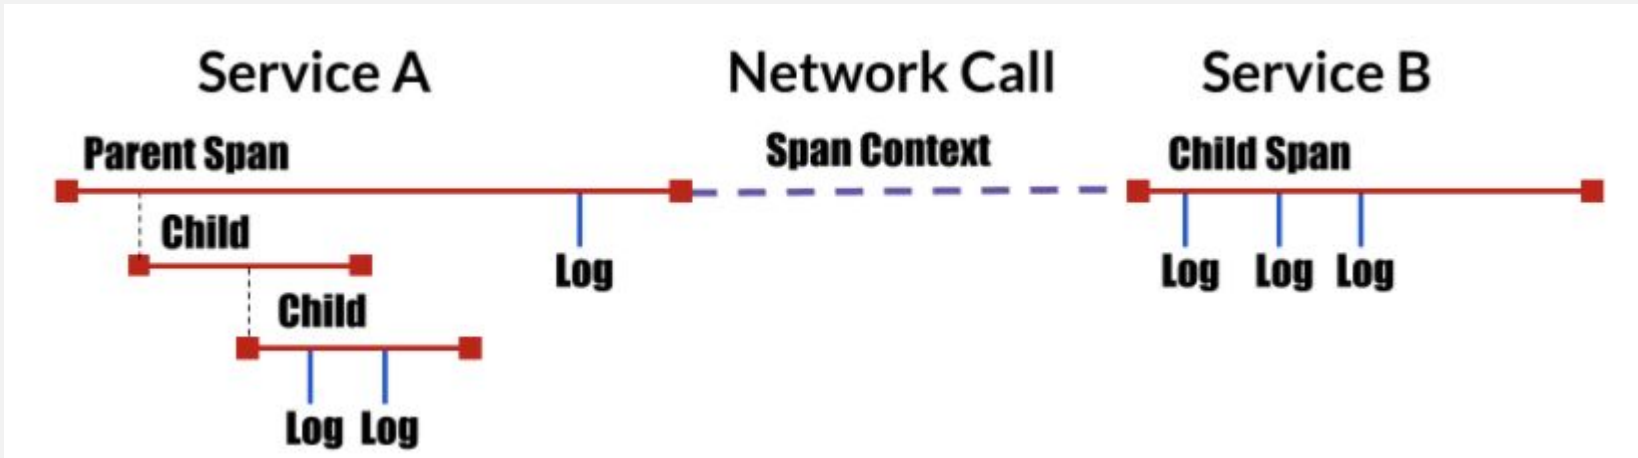
QAWARE

- Distributed Tracing: Technology for tracing calls and processes in distributed software systems.
- Today's considerations go back to the paper: Dapper, a Large-Scale Distributed Systems Tracing Infrastructure
- Idea: Each service forwards information from the caller to the next service and back again. This results in a directed graph.
 - Each service must be instrumented. Otherwise there will be a gap.
 - The smallest unit is called a span. A span has a duration and descriptive information. A trace is a set of 1..n spans.
- Many implementations exist for different programming languages and technologies.
 - <https://opentelemetry.io/docs/languages/>
 - <https://opentelemetry.io/docs/zero-code/>

Distributed Tracing



QAWARE




Traces: Foundation

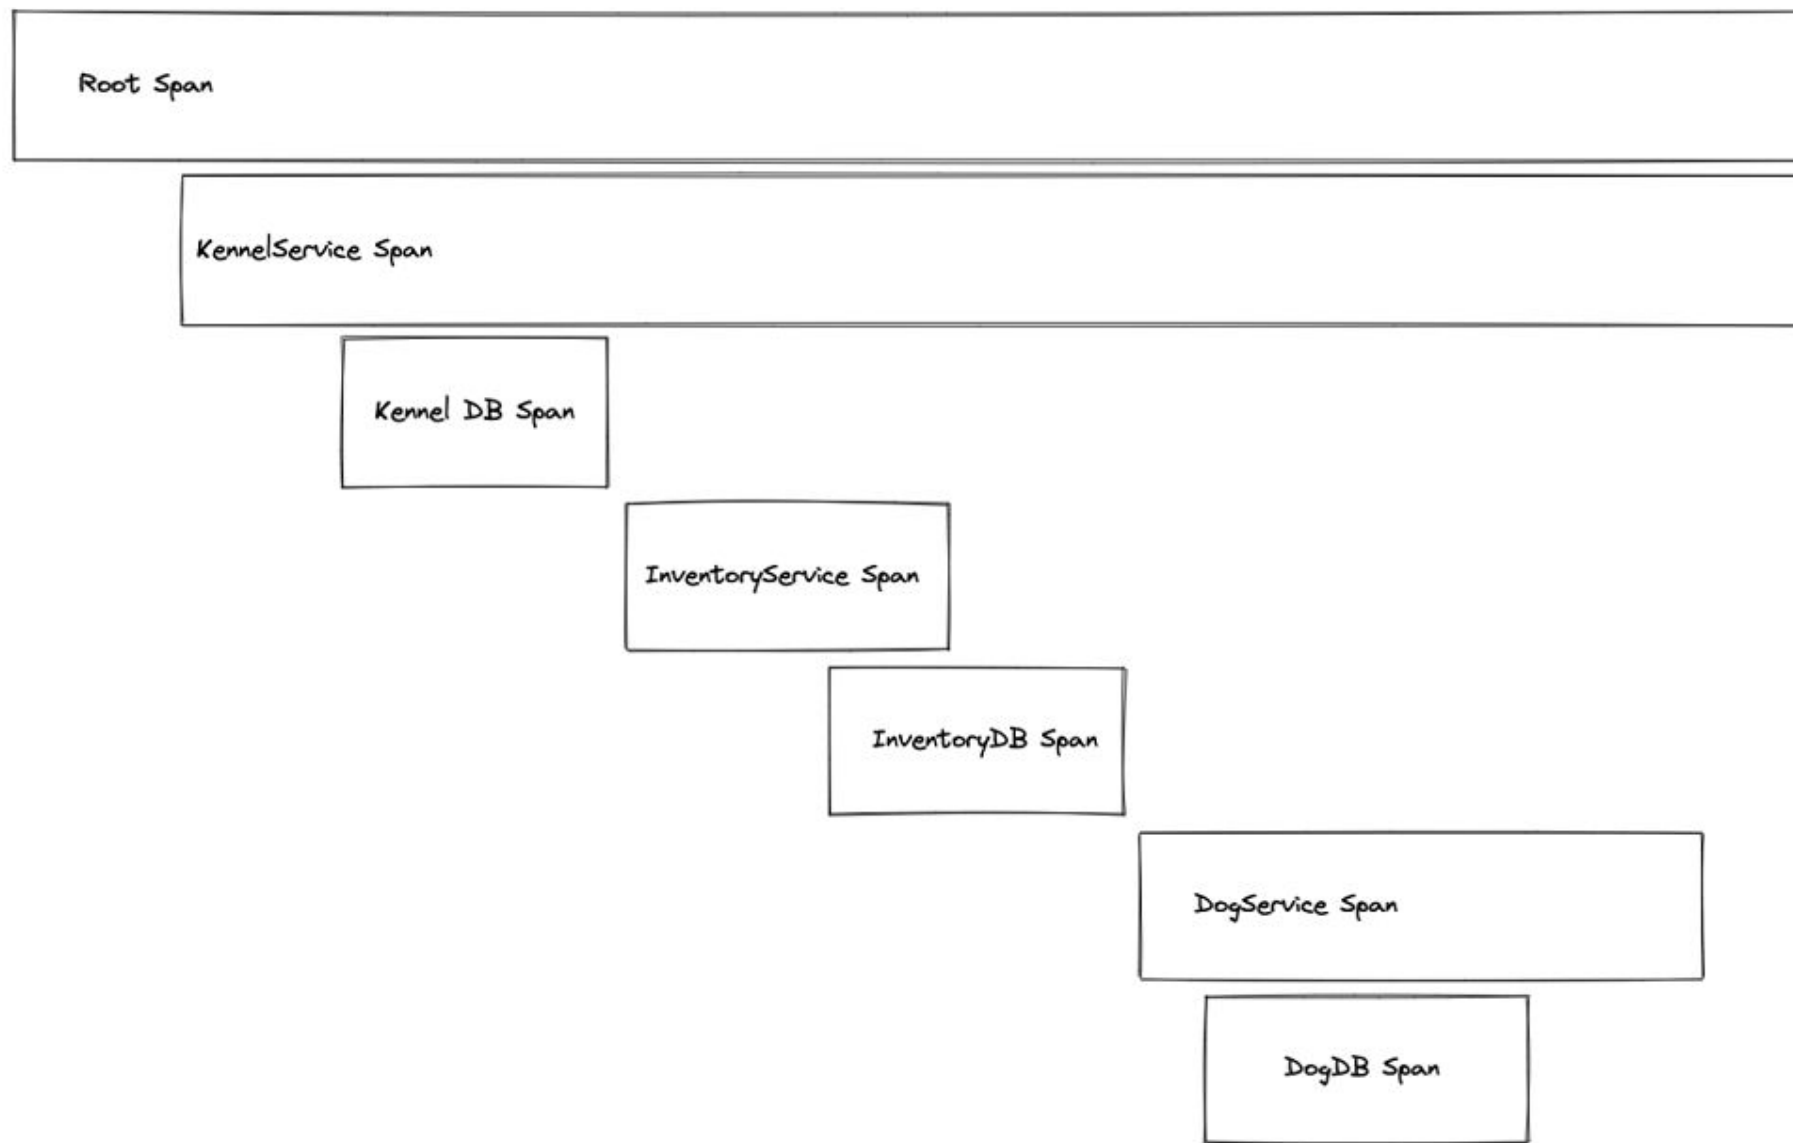



QAWARE

Traces visualize which paths a request has taken through the microservices.....

... and where an error may have occurred!





Grafana Tempo



QAWARE

Traces in the Grafana ecosystem!

SDKs for various platforms

Compatible with various tracing agents:

- Jaeger
- Zipkin
- Opentelemetry

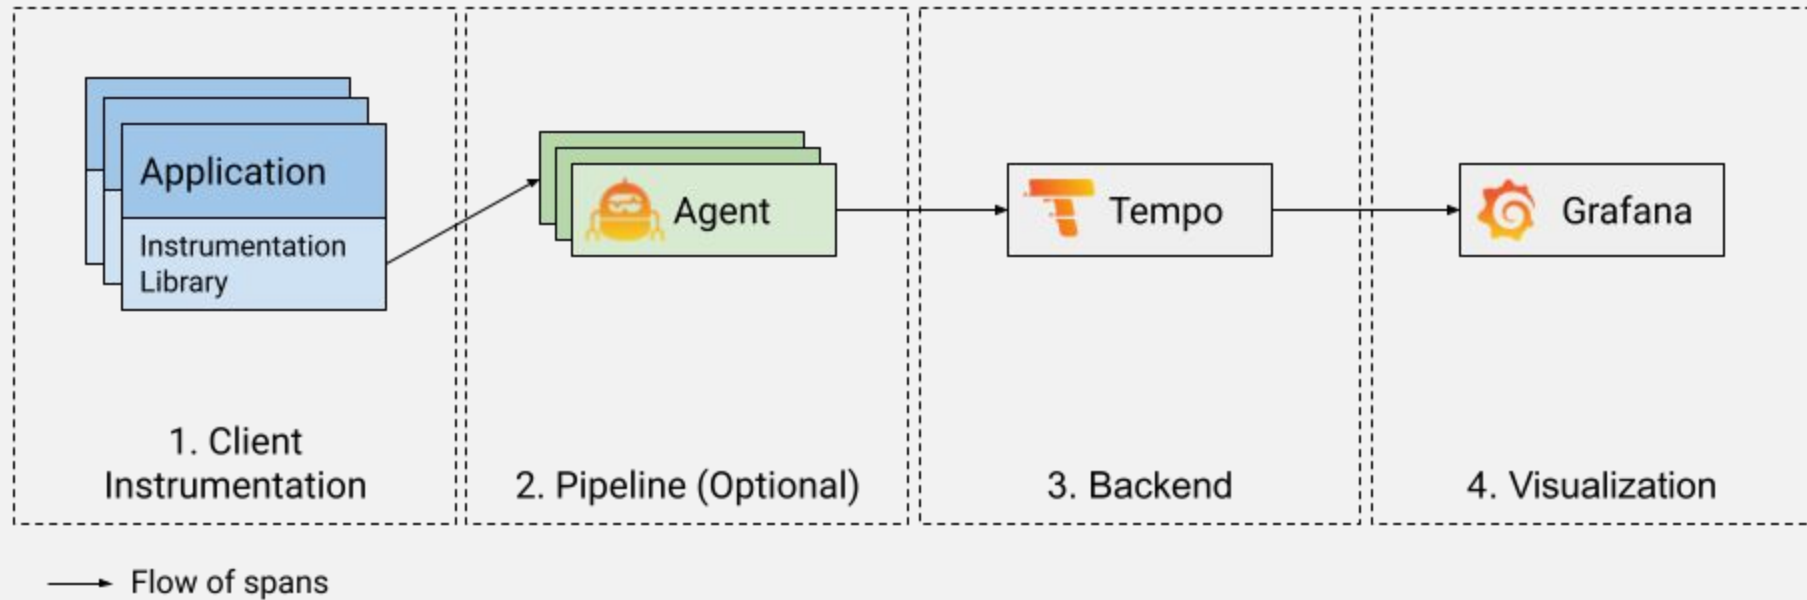
Since: 2020

License: AGPL 3.0

Architecture



QAWARE



Combining metrics and traces? Exemplars!

Supported in Micrometer 1.9, for example




Image:
<https://grafana.com/blog/2021/03/31/intro-to-exemplars-which-enable-grafana-tempo-distributed-tracing-at-massive-scale/>



QAWARE

Visualization



Grafana



QAWARE

Visualization for many different data sources:

- Mimir
- Loki
- Tempo
- InfluxDB
- ...

Since: 2014

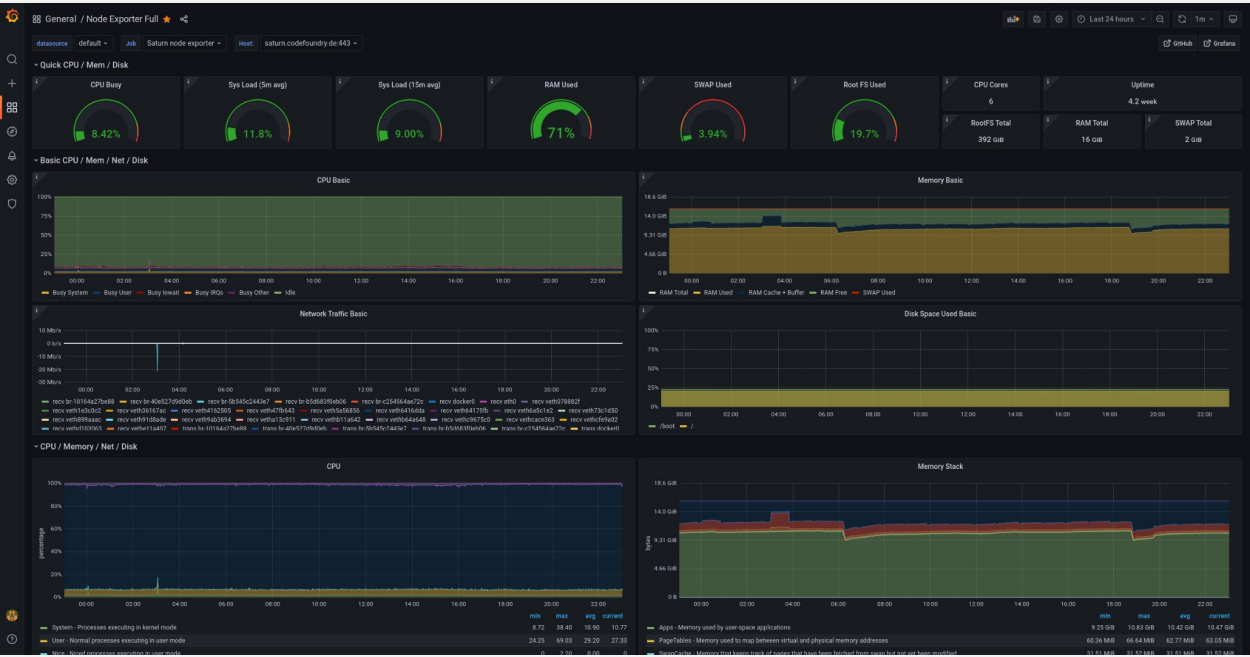
License: AGPL 3.0

Dashboards



QA|WARE

Data of all types is visualized in dashboards.



Grafana - Dashboards



QAWARE

Don't worry: you don't have to laboriously click together dashboards yourself.

<https://grafana.com/grafana/dashboards/>



QA|WARE

Showcase



QAWARE

More from the Grafana universe

Load tests: k6



QA|WARE

Built with Go – the tests are written in JavaScript.

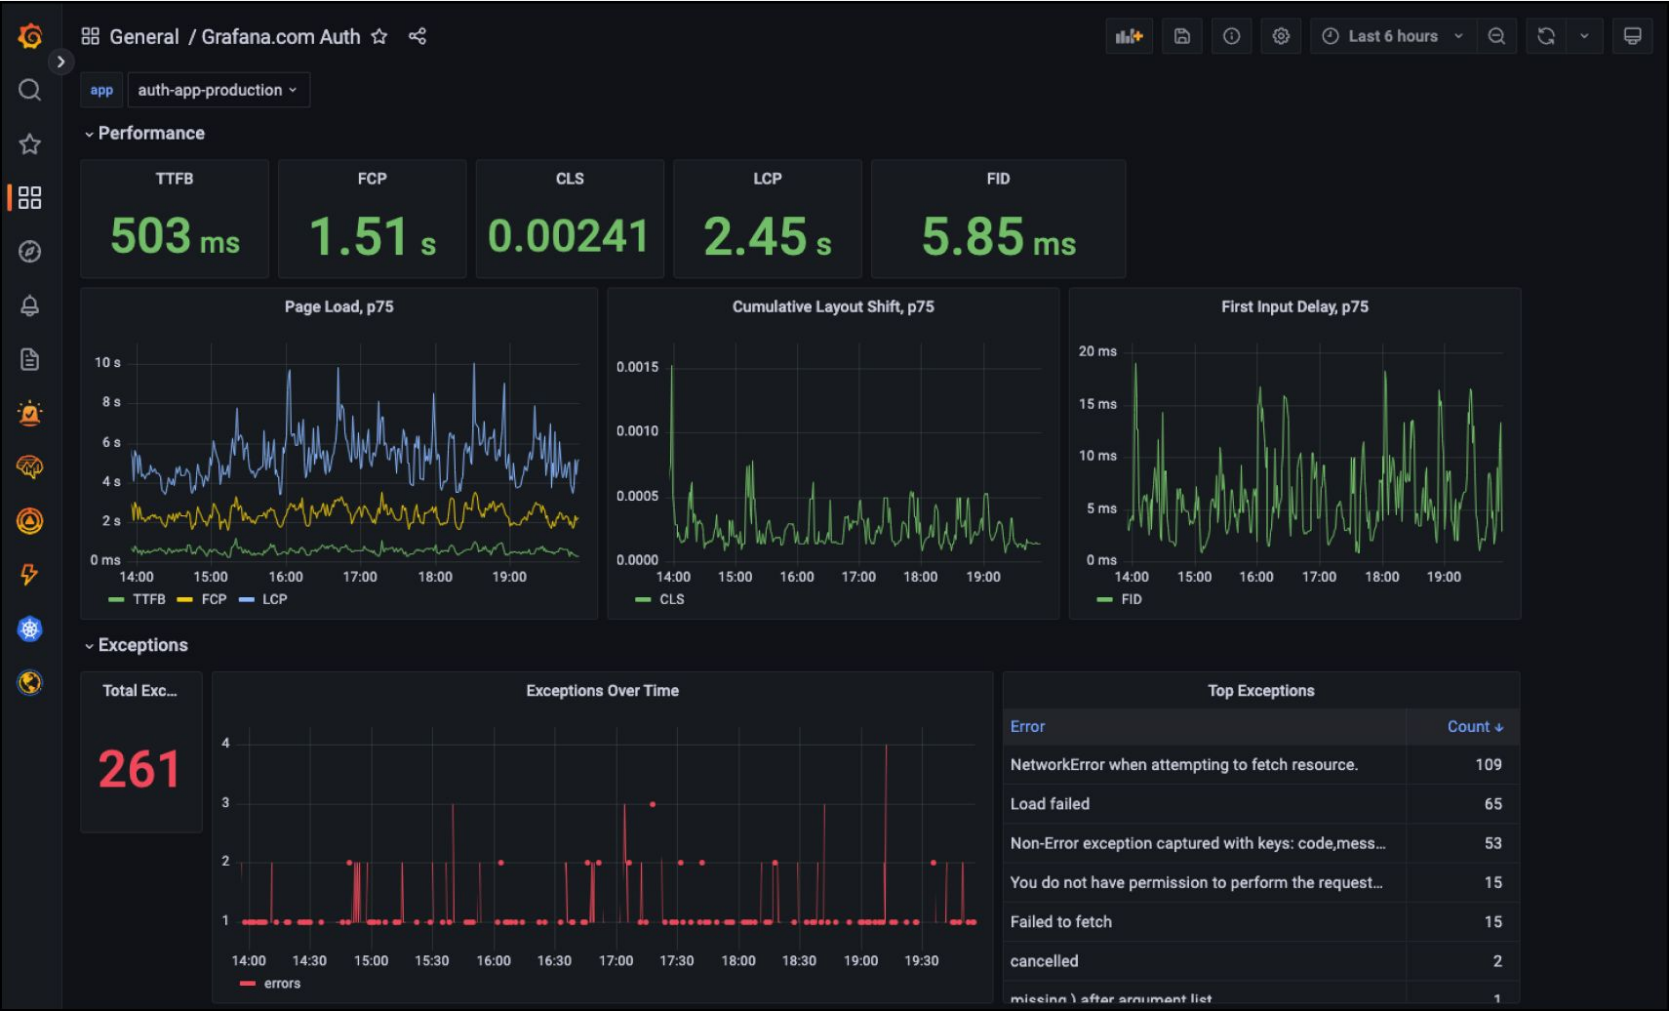
Test results end up in an InfluxDB – visualization is done with Grafana.



Frontend observability: Faro



QAWARE



<https://grafana.com/oss/faro/>

Further reading



QAWARE

Blog: <https://blog.qaware.de/posts/cloud-observability-grafana-spring-boot/>

Showcase: <https://github.com/zalintyre/cloud-observability-grafana-spring-boot>

Grafana @ Heise Mastering Kubernetes:

<https://de.slideshare.net/QAware/cloud-observability-mit-loki-prometheus-tempo-und-grafana>

Grafana: <https://grafana.com/>

Mimir: <https://grafana.com/docs/mimir/latest/>

Loki: <https://grafana.com/oss/loki/>

Tempo: <https://grafana.com/oss/tempo/>

Sources



QAWARE

[Kal] R.E. Kalman,
On the general theory of control systems,
IFAC Proceedings Volumes,
Volume 1, Issue 1,
1960,
Pages 491-502,
ISSN 1474-6670,
[https://doi.org/10.1016/S1474-6670\(17\)70094-8](https://doi.org/10.1016/S1474-6670(17)70094-8).
(<https://www.sciencedirect.com/science/article/pii/S1474667017700948>)

qaware.de



QA|WARE
SOFTWARE ENGINEERING

QAware GmbH

Aschauer Straße 32
81549 München
Tel. +49 89 232315-0
info@qaware.de



twitter.com/qaware



linkedin.com/company/qaware-gmbh



xing.com/companies/qawaregmbh



slideshare.net/qaware



github.com/qaware